Azure's Data Factory & Blob Storage

Lab 06

by

Diane Howard, Nishava Inc.

Deep Azure @McKesson

Overview

- What is Azure's Data Factory, How to use Data Factory, Costs
- Create Storage Account
- Create Active Directory for Client ID, Client Secret, Tenant ID
- Data Factory Demo in Python:
    1. Create Resource Group
    2. Create Data Factory
    3. Create Linked Storage for Blob Storage and Blob Sink
    4. Initialize Blob Storage for input/output data
    5. Create Pipeline
    6. Monitor your Pipeline
    7. Run your Pipeline

Objective of Demo

- Create a Data Factory in Azure within Python code that will ingest data from an Excel spreadsheet (or Blob), perform an identify transformation (identical copy) and transfer data to another Blob data store as a sink.

- This is accomplished via a workflow (job) initialized within the Data Factory.

Data Factory

- There are 2 versions of Data Factory within Azure: V1 and V2 (Preview)

- Defined as data-driven jobs (*aka workflows*) that have pipelines to move & transform data

1. **Ingest** data from disparate data stores (e.g., AZ blob/file/tables, SQL, Cosmos, Amazon Redshift, Informix, PostgresSQL, NoSQL, Amazon S3, FTP, HDFS, …)

https://docs.microsoft.com/en-us/azure/data-factory/concepts-datasets-linked-services

2. **Transform** or **process** the data by using compute services such as the following:
   - Azure HDInsight Hadoop
   - Spark
   - Azure Data Lake Analytics
   - Azure Machine Learning

3. **Publish output data** to data stores (e.g., AZ Blob, AZ Cosmos DB, SQL Server, Data Warehouse, Oracle, AZ Table storage, AZ Filesystem)

Example of Data Factory Usage

- Central place to manage processing of web log analytics, click stream analysis, social sentiment, sensor data analysis, geo-location analysis, etc.

TRANSFORMATIONS

- A gaming company collects logs from games in the cloud.
  - analyze logs to gain insights into customer preferences, demographics, usage behavior

Data Factory V1 vs V2

V1

- ✓ Create data pipelines to move and transform data
- ✓ Run pipelines on a specified schedule (hourly, daily, weekly, etc.)
- ✓ Visualizations to display the lineage and dependencies between your data pipelines
- ✓ Monitor data pipelines
  - pinpoint issues and setup monitoring alerts.

V2 (Preview)

Primary:

1. **Control flow**:
   Branching, looping & conditional processing.
2. **Deploy and run SQL Server Integration Services (SSIS) packages** in Azure.

- ✓ Support for virtual network (VNET) environments.
- ✓ Scale out with on-demand processing power.
- ✓ Support on-demand Spark cluster.
- ✓ Flexible scheduling to support incremental data loads.
- ✓ Triggers for executing data pipelines.

Available APIs

V1 (Nov 2016)

Batch processing of time series data.

- AZ Portal
- Copy Wizard
- Visual Studio
- Azure PowerShell
- Azure Resource Manager template
- REST API
- .NET API

V2 (Sept 2017)

'General-purpose hybrid data integration service'

- AZ Portal (note: limited capabilities)
- Azure PowerShell
- Languages:
  .NET & Python
- REST API

Data Factory V2 Costs

The pricing for Data Factory usage is calculated based on the following factors:

1. Number of activities run.
   Orchestration of activities
2. Volume of data moved.
   Data Movement
3. SQL Server Integration Services (SSIS) compute hours.
4. Whether a pipeline is active or not.

Inactive Pipelines are charged!

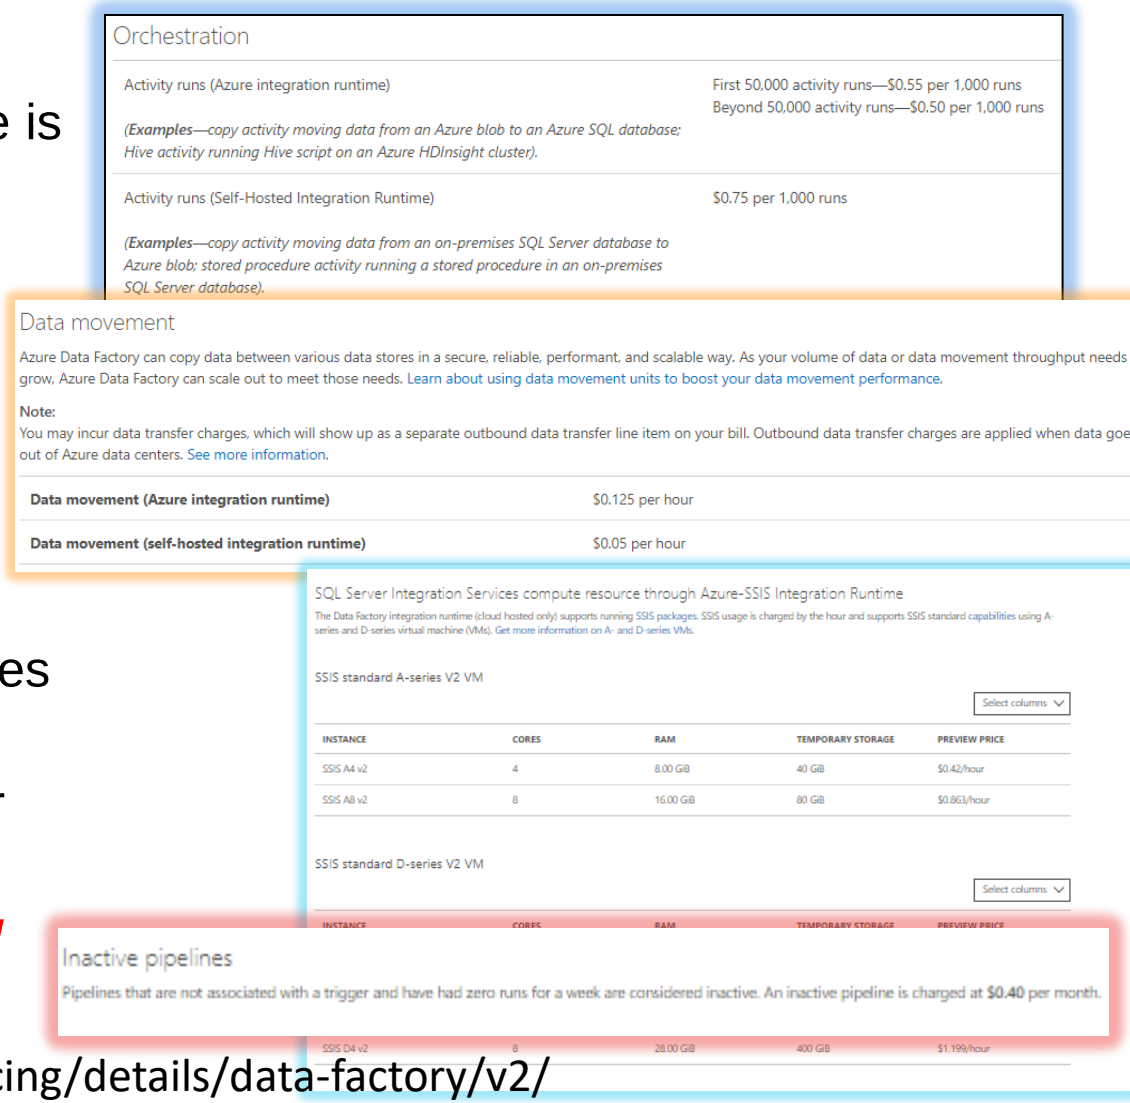

Orchestration

Activity runs (Azure integration runtime) *(**Examples**—copy activity moving data from an Azure blob to an Azure SQL database; Hive activity running Hive script on an Azure HDInsight cluster).*	First 50,000 activity runs—$0.55 per 1,000 runs Beyond 50,000 activity runs—$0.50 per 1,000 runs
Activity runs (Self-Hosted Integration Runtime) *(**Examples**—copy activity moving data from an on-premises SQL Server database to Azure blob; stored procedure activity running a stored procedure in an on-premises SQL Server database).*	$0.75 per 1,000 runs

Data movement

Azure Data Factory can copy data between various data stores in a secure, reliable, performant, and scalable way. As your volume of data or data movement throughput needs grow, Azure Data Factory can scale out to meet those needs. Learn about using data movement units to boost your data movement performance.

Note:
You may incur data transfer charges, which will show up as a separate outbound data transfer line item on your bill. Outbound data transfer charges are applied when data goe out of Azure data centers. See more information.

Data movement (Azure integration runtime)	$0.125 per hour
Data movement (self-hosted integration runtime)	$0.05 per hour

SQL Server Integration Services compute resource through Azure-SSIS Integration Runtime

The Data Factory integration runtime (cloud hosted only) supports running SSIS packages. SSIS usage is charged by the hour and supports SSIS standard capabilities using A-series and D-series virtual machine (VMs). Get more information on A- and D-series VMs.

SSIS standard A-series V2 VM

Select columns

INSTANCE	CORES	RAM	TEMPORARY STORAGE	PREVIEW PRICE
SSIS A4 v2	4	8.00 GiB	40 GiB	$0.42/hour
SSIS A8 v2	8	16.00 GiB	80 GiB	$0.863/hour

SSIS standard D-series V2 VM

Select columns

INSTANCE	CORES	RAM	TEMPORARY STORAGE	PREVIEW PRICE
SSIS D4 v2	8	28.00 GiB	400 GiB	$1.199/hour

Inactive pipelines

Pipelines that are not associated with a trigger and have had zero runs for a week are considered inactive. An inactive pipeline is charged at **$0.40** per month.

https://azure.microsoft.com/en-us/pricing/details/data-factory/v2/

Prerequisites to Create a Data Factory V2

Python 2.7, 3.3, 3.4, 3.5 or 3.6

- Install Python SDK for Azure packages
    - Azure Management Resources
    - Data Factory
- Portal: Obtain your Subscription ID
- Portal: Create an Azure Storage Account & Blob container
- Portal: Create a data file and upload data file to Blob container
- Portal: Create an app in Active Directory (Client ID, Client Secret).
- Python: Create a Data Factory, Linked Service, Pipeline

My Development Environment

- Windows 10
- Python 3.6.2
- Anaconda

```
c:\Users\dhoward>python --version
Python 3.6.2 :: Anaconda, Inc.
```

ANACONDA.

Windows

Windows 10

Home

Python 3.6.2

Release Date: 2017-07-17

Install Python Package Azure Management Resources

Python SDK for Data Factory (supports Python 2.7, 3.3, 3.4, 3.5 and 3.6)

1. Open Command Prompt as Administrator
2. Install Python package for Azure Management Resources

   pip install azure-mgmt-resource
3. Install Python Package for Azure Data Factory

   pip install azure-mgmt-datafactory

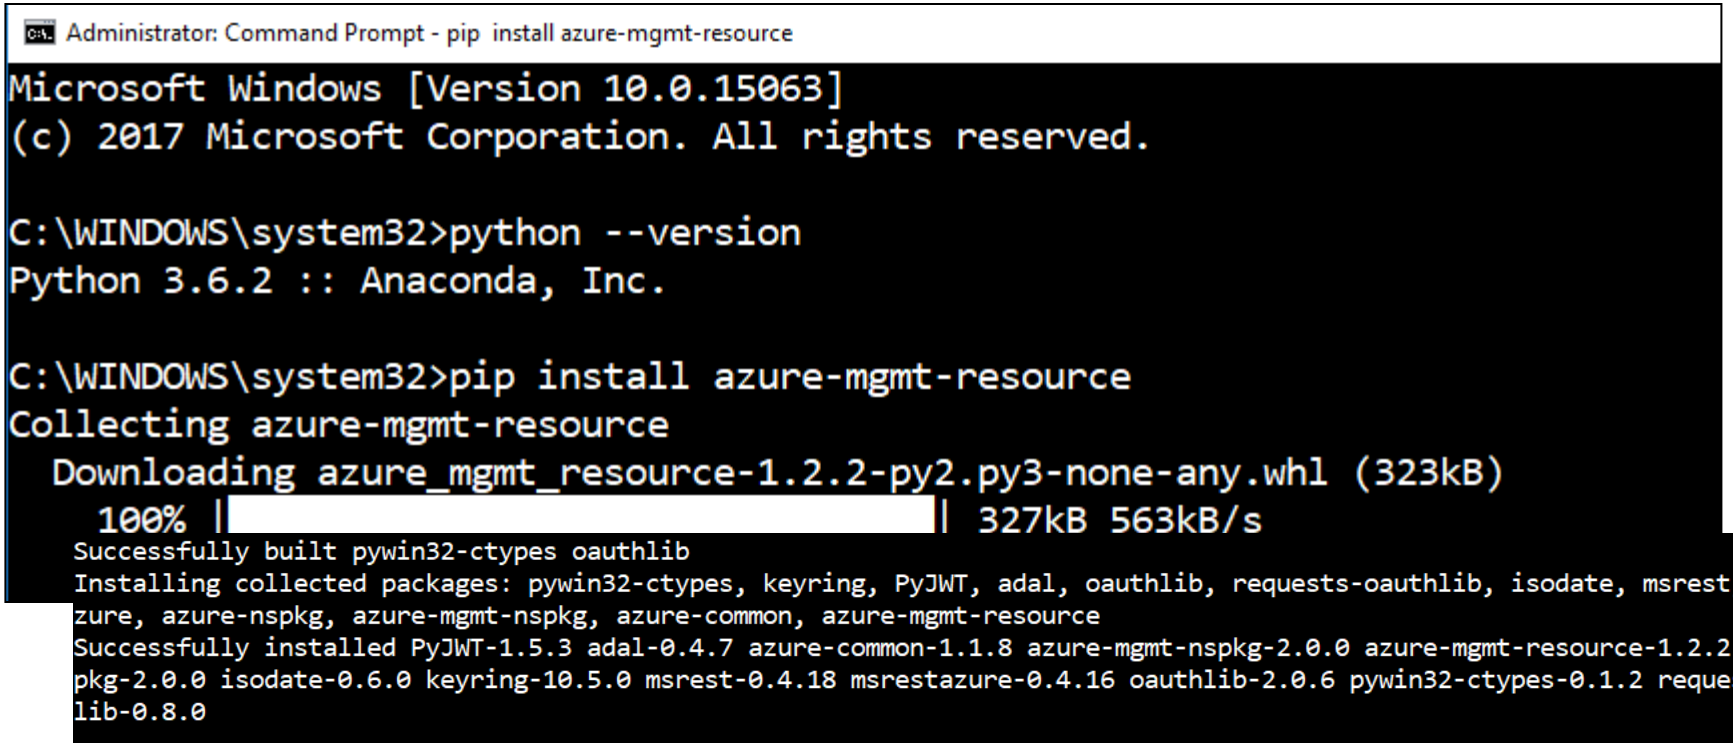

Install Python Package Data Factory

- Open Command Prompt as Administrator
- Install Python Package for Azure Data Factory

```
C:\WINDOWS\system32>pip install azure-mgmt-datafactory
Collecting azure-mgmt-datafactory
  Downloading azure_mgmt_datafactory-0.2.1-py2.py3-none-any.whl (249kB)
    100% |████████████████████████████████| 256kB 1.6MB/s
Requirement already satisfied: azure-mgmt-nspkg>=2.0.0 in c:\users\dhoward\anaconda3_64bit\lib\si
te-packages (from azure-mgmt-datafactory)
Requirement already satisfied: msrestazure~=0.4.11 in c:\users\dhoward\anaconda3_64bit\lib\site-p
ackages (from azure-mgmt-datafactory)
Requirement already satisfied: azure-common~=1.1 in c:\users\dhoward\anaconda3_64bit\lib\site-pac
kages (from azure-mgmt-datafactory)
Requirement already satisfied: azure-nspkg>=2.0.0 in c:\users\dhoward\anaconda3_64bit\lib\site-pa
ckages (fro
Requirement already satisfied: oauthlib>=0.6.2 in c:\users\dhoward\anaconda3_64bit\lib\site-packages (from requests-oauthl
ib>=0.5.0->msrest~=0.4.17->msrestazure~=0.4.11->azure-mgmt-datafactory)
Requirement already satisfied: pycparser in c:\users\dhoward\anaconda3_64bit\lib\site-packages (from cffi>=1.7->cryptograp
hy>=1.1.0->adal~=0.4.0->msrestazure~=0.4.11->azure-mgmt-datafactory)
Installing collected packages: azure-mgmt-datafactory
Successfully installed azure-mgmt-datafactory-0.2.1

C:\WINDOWS\system32>
```

Check your Roles in Your Subscription

- You should be a Contributor Role. If not someone within McKesson Admin should set users up as Contributors.

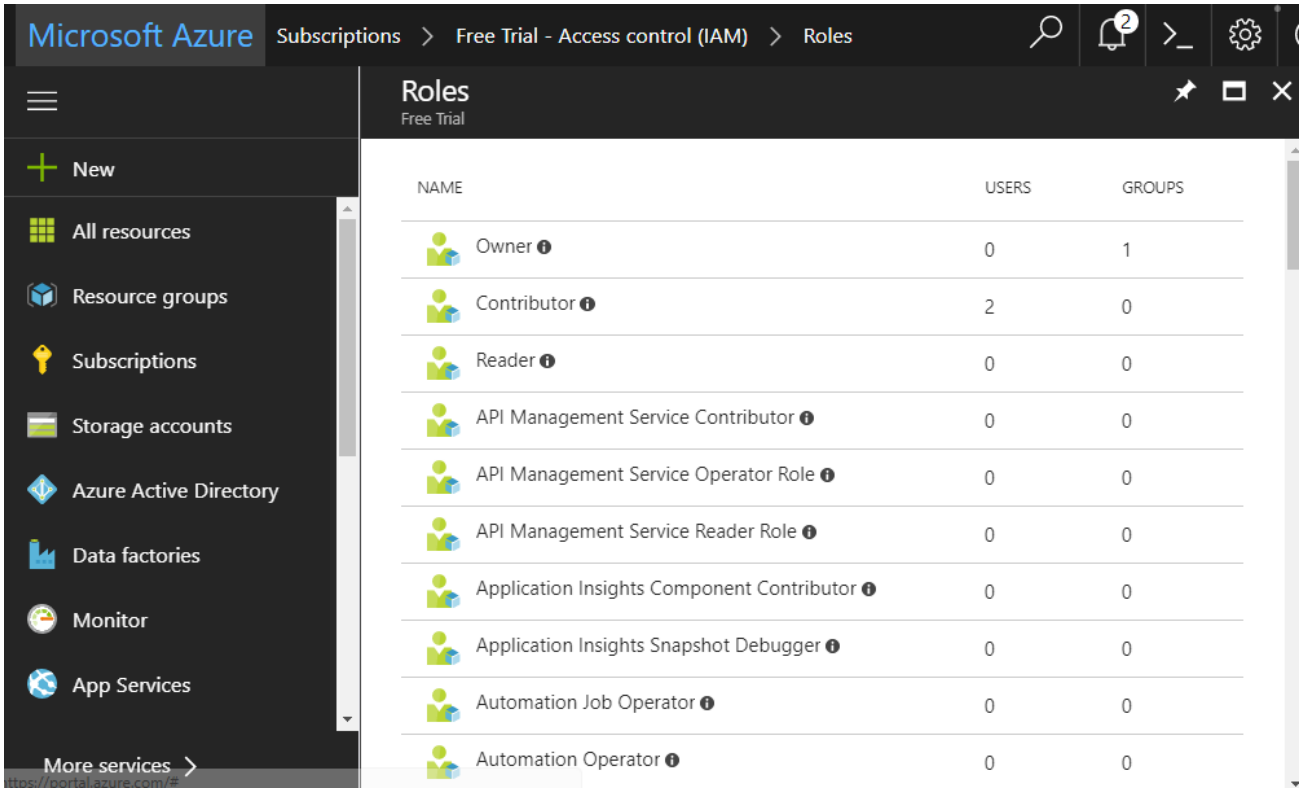

Obtain your Subscription ID

- Uniquely identifies your subscription to use AZ services.

- In the left navigation panel, click Subscriptions. Copy your subscription ID.

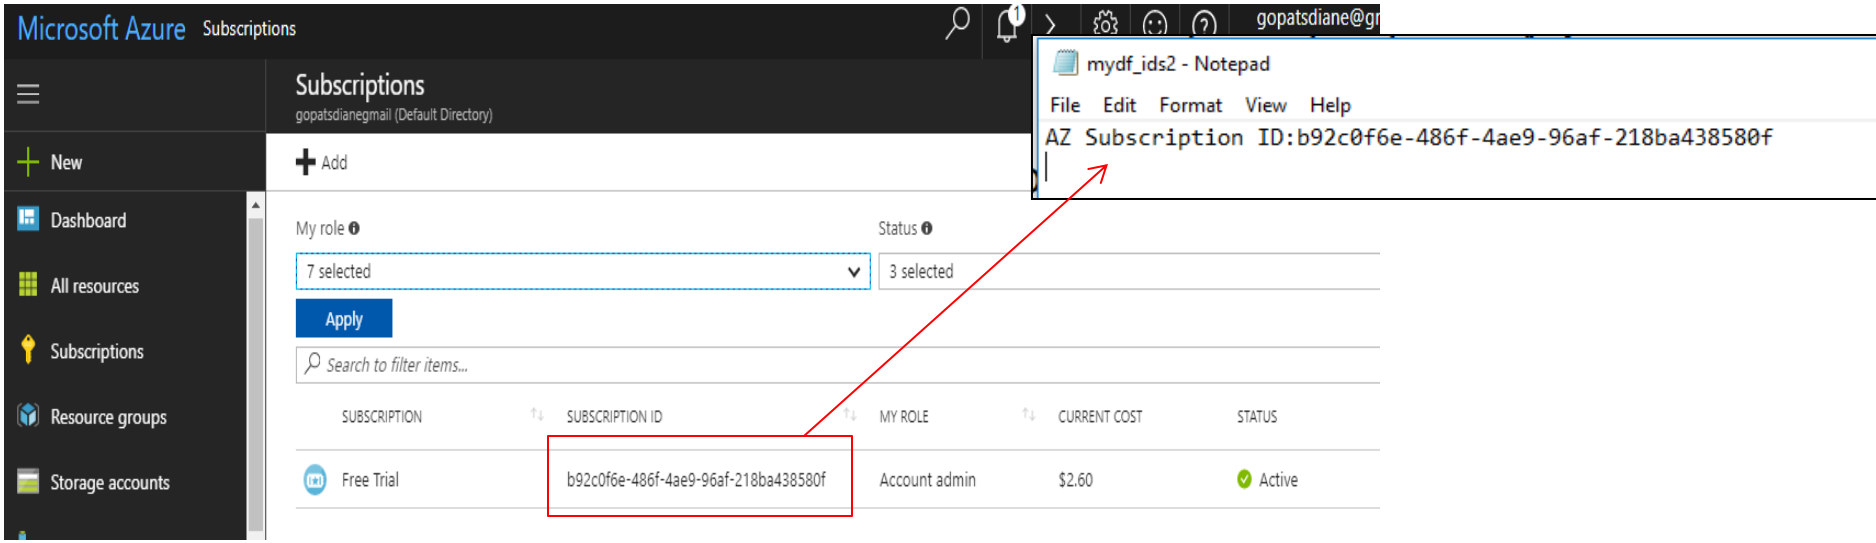

Create a Blob Storage Account

Select:

1. Name: Must be unique within Azure and lowercase.
2. Deployment model: Resource Mgr
3. Account Kind: Blob or Gen Purpose
4. Use default for Replication.

(Other options changes price.)

5. Create new resource group or use an existing one.
6. Location: select any location close to where you are located.

(I use East US).

7. Virtual Networks (Preview): Disabled

Optional: Pin to Dashboard

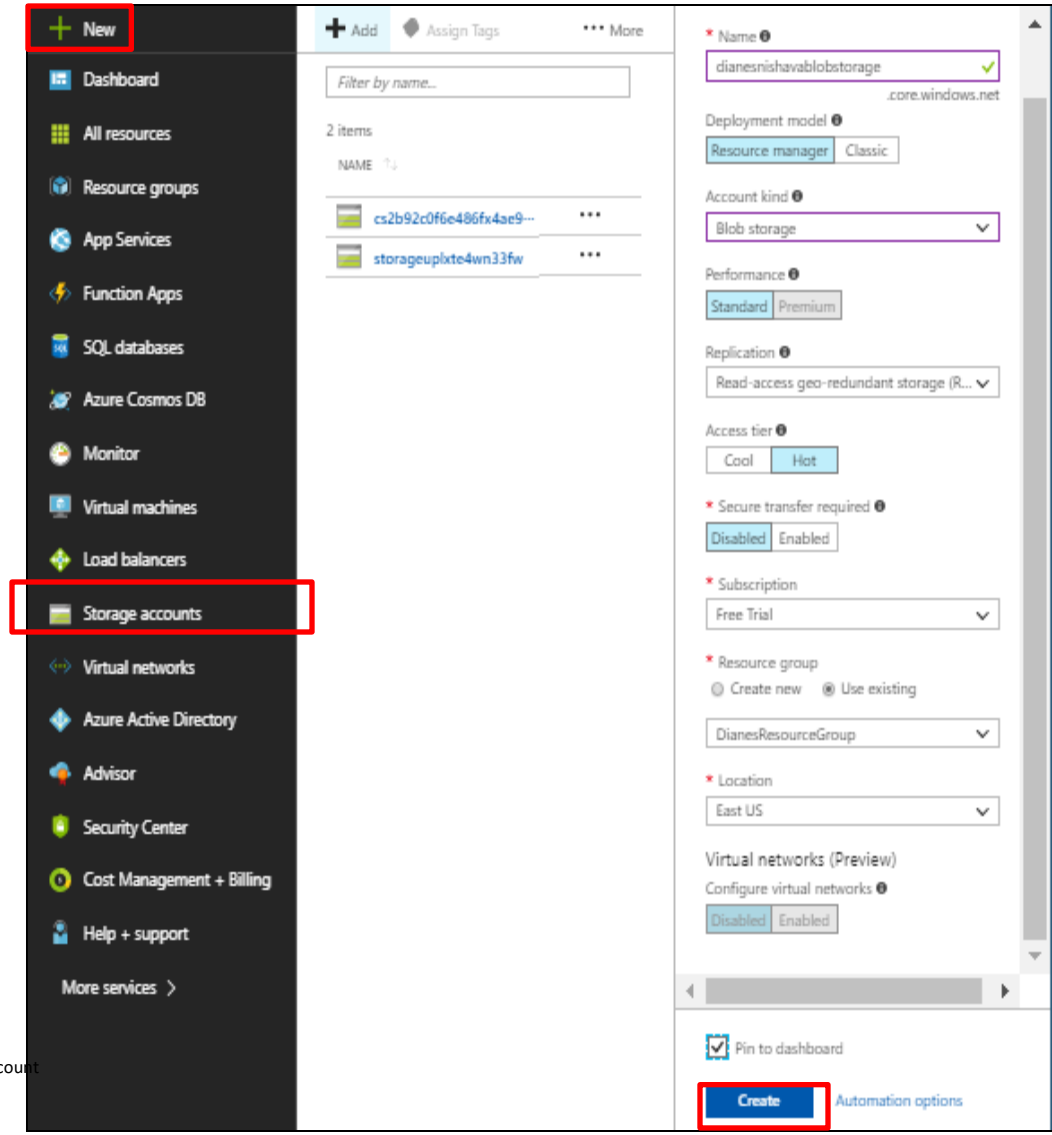

https://docs.microsoft.com/en-us/azure/storage/common/storage-create-storage-account#create-a-storage-account

Wait for your Deployment

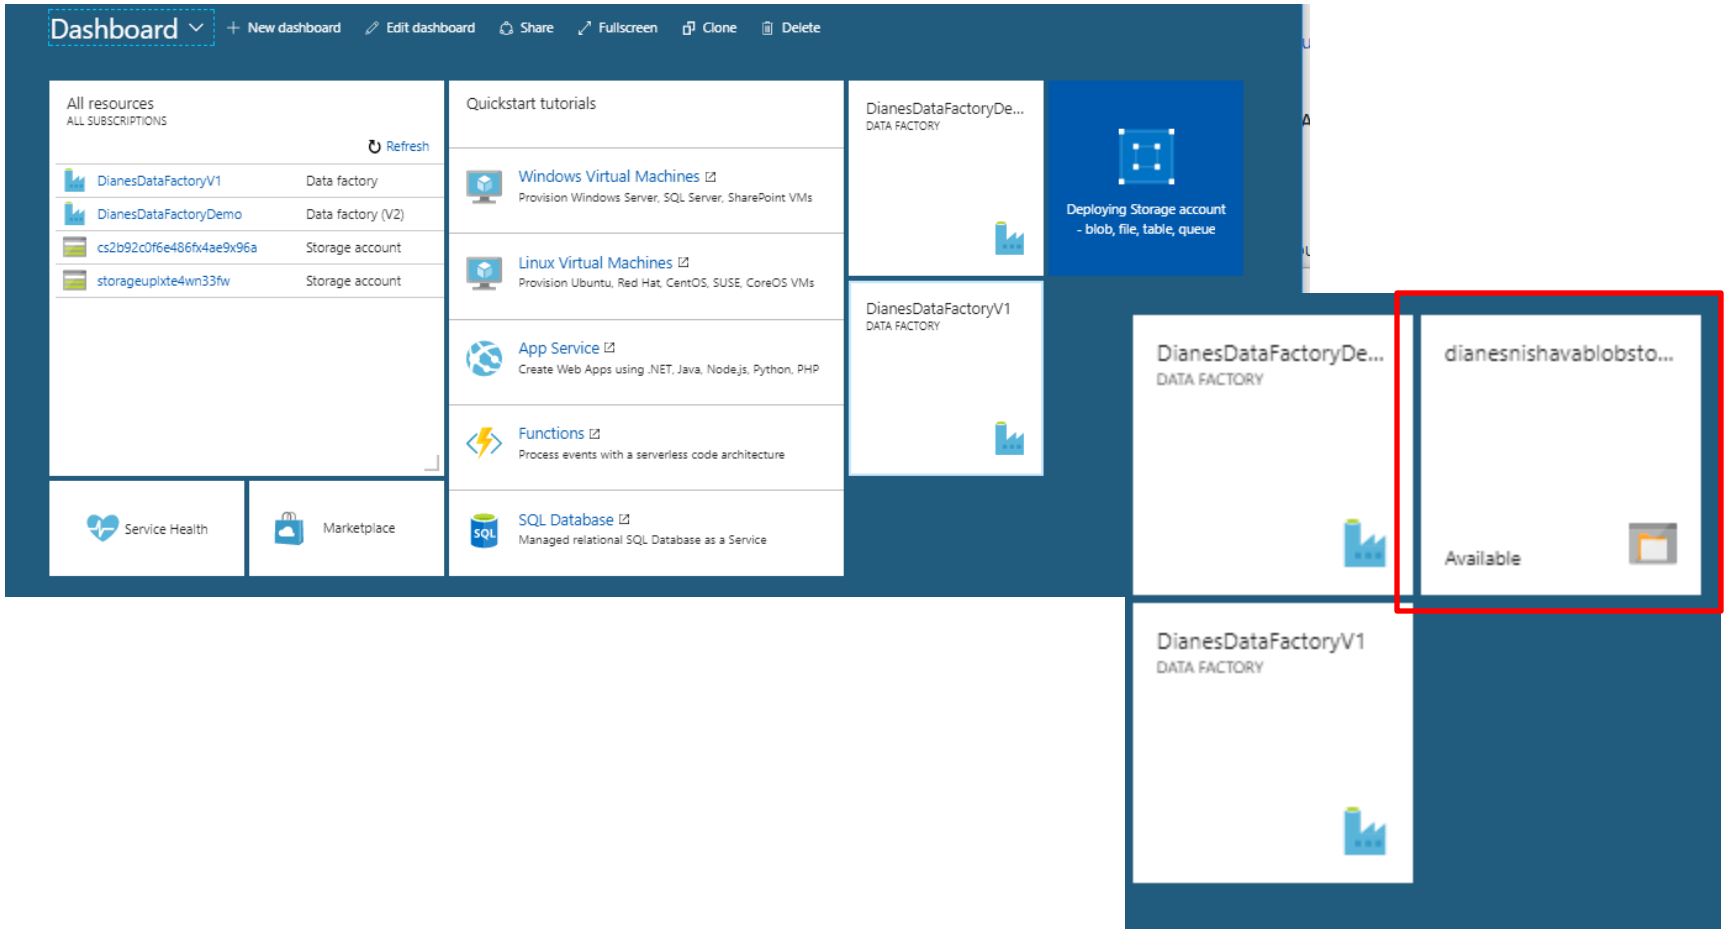

Create a Container for your Blob Storage

- Stores unlimited number of blobs.
- Container name must be lowercase.
- Blob: A file of any type and size.
- **Azure Storage** offers three types of blobs: block blobs, page blobs & append blobs.

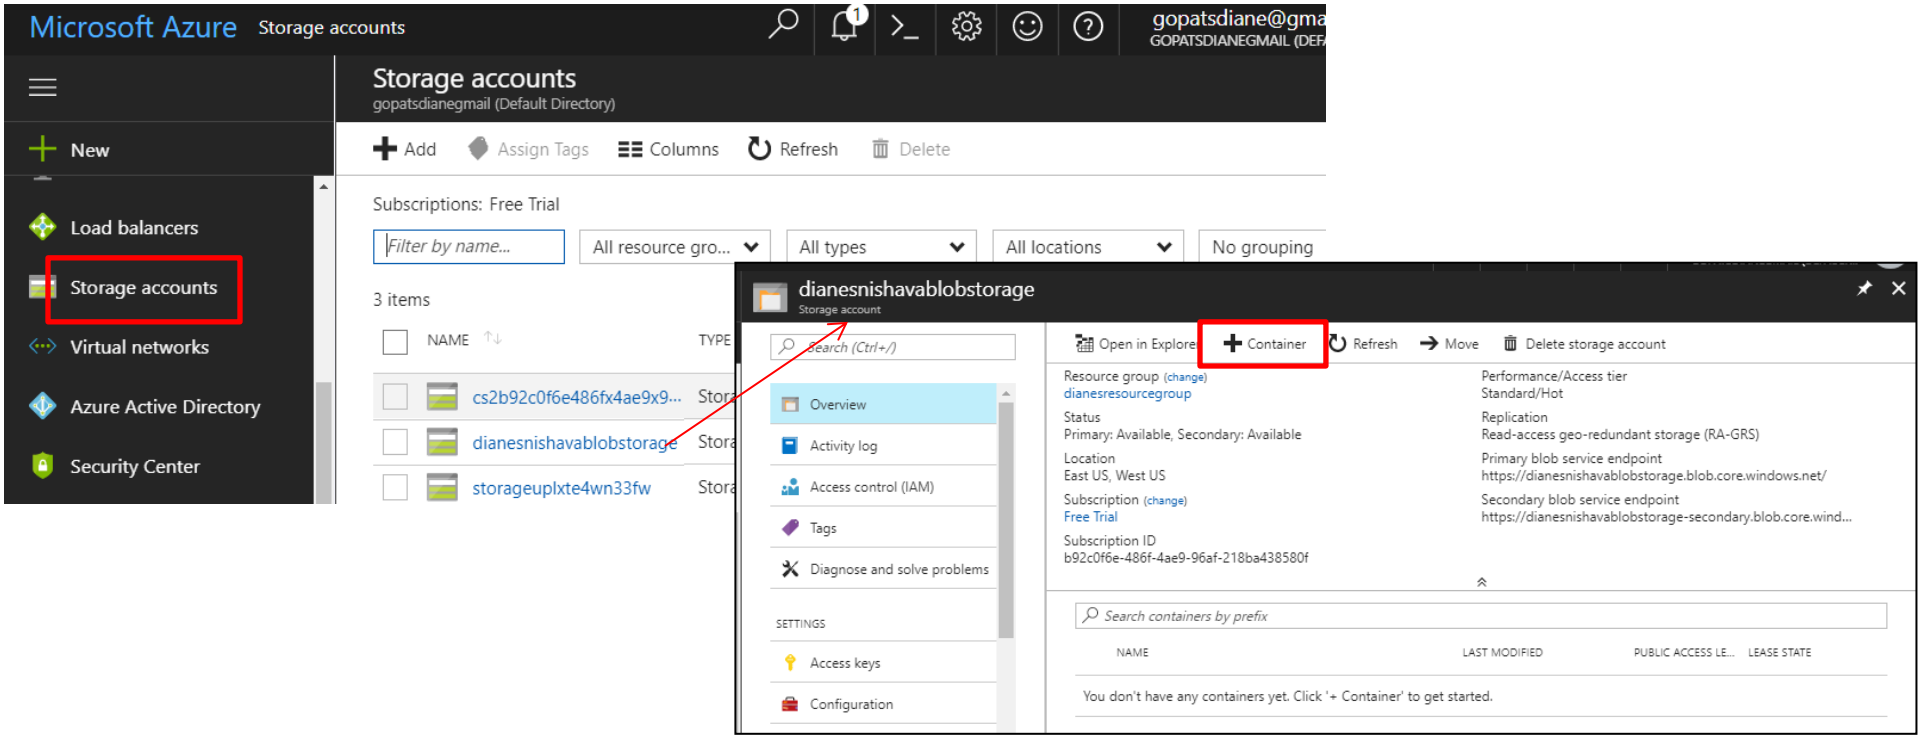

Create Storage Container

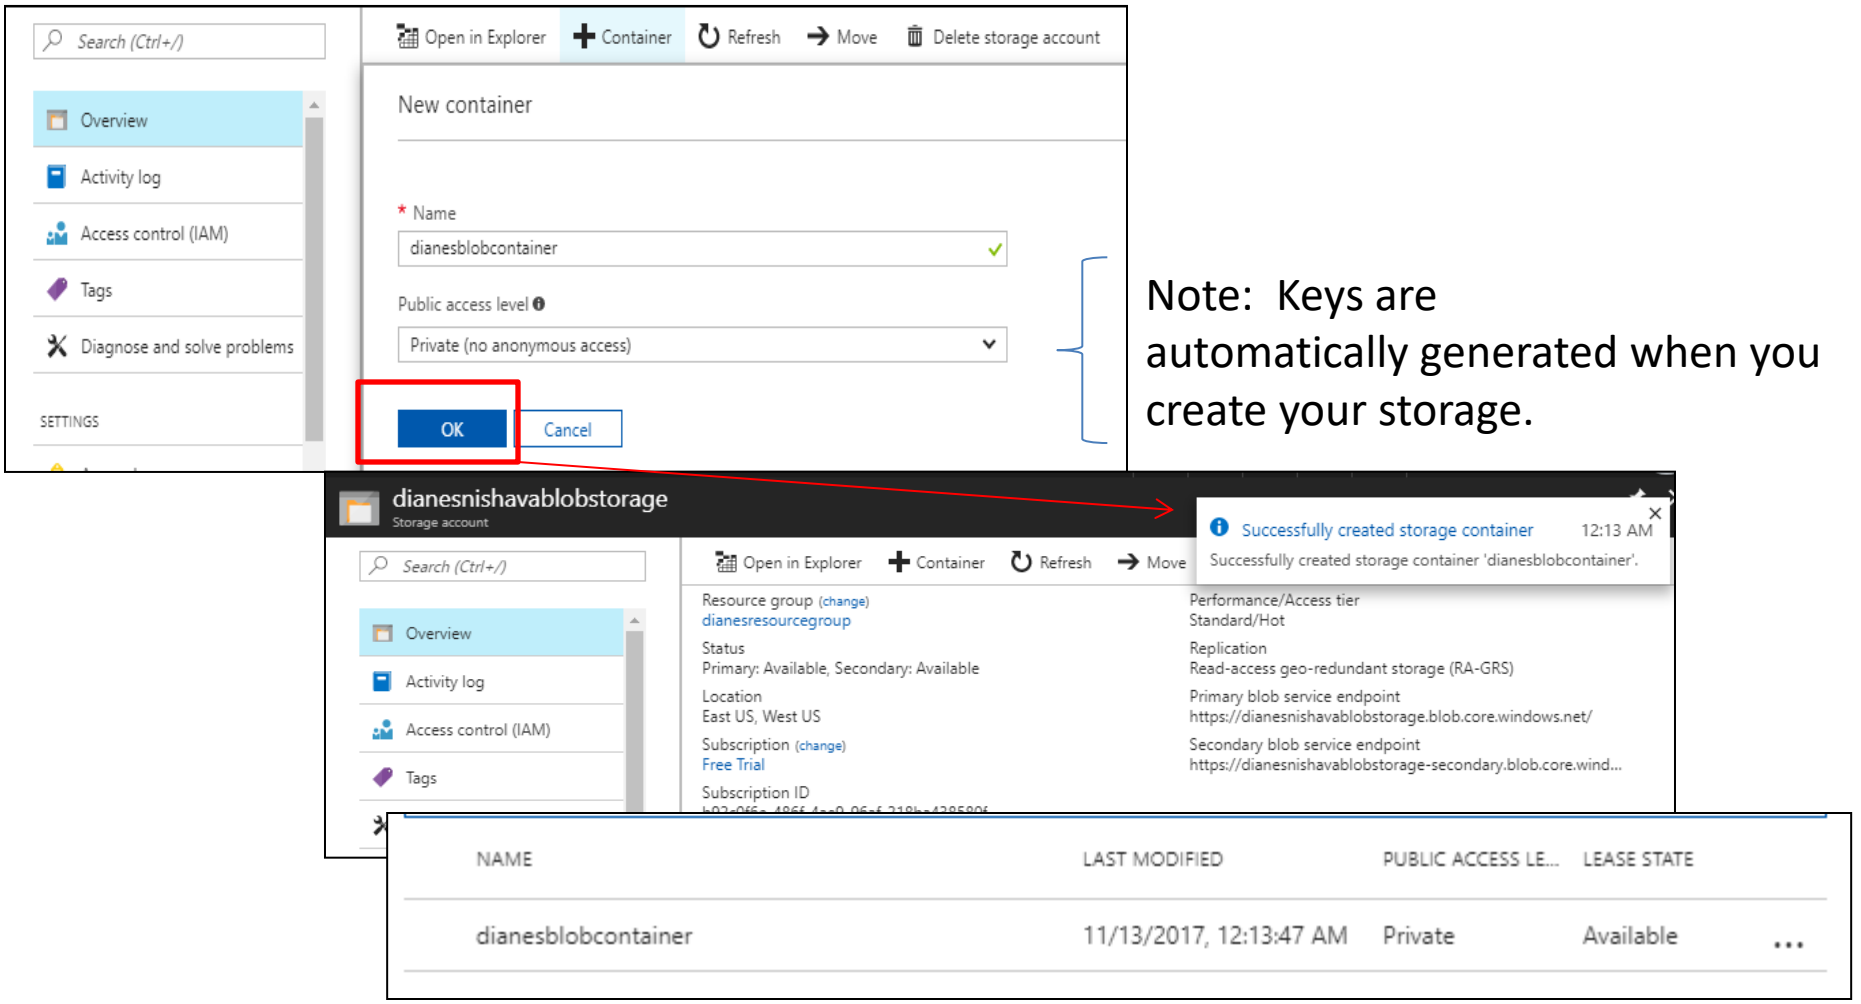

Storage Access Keys

- Two 512-bit storage access keys are created automatically when you create general-purpose or blob storage.
- Used for authentication to access your storage.
- Never share your storage access keys!

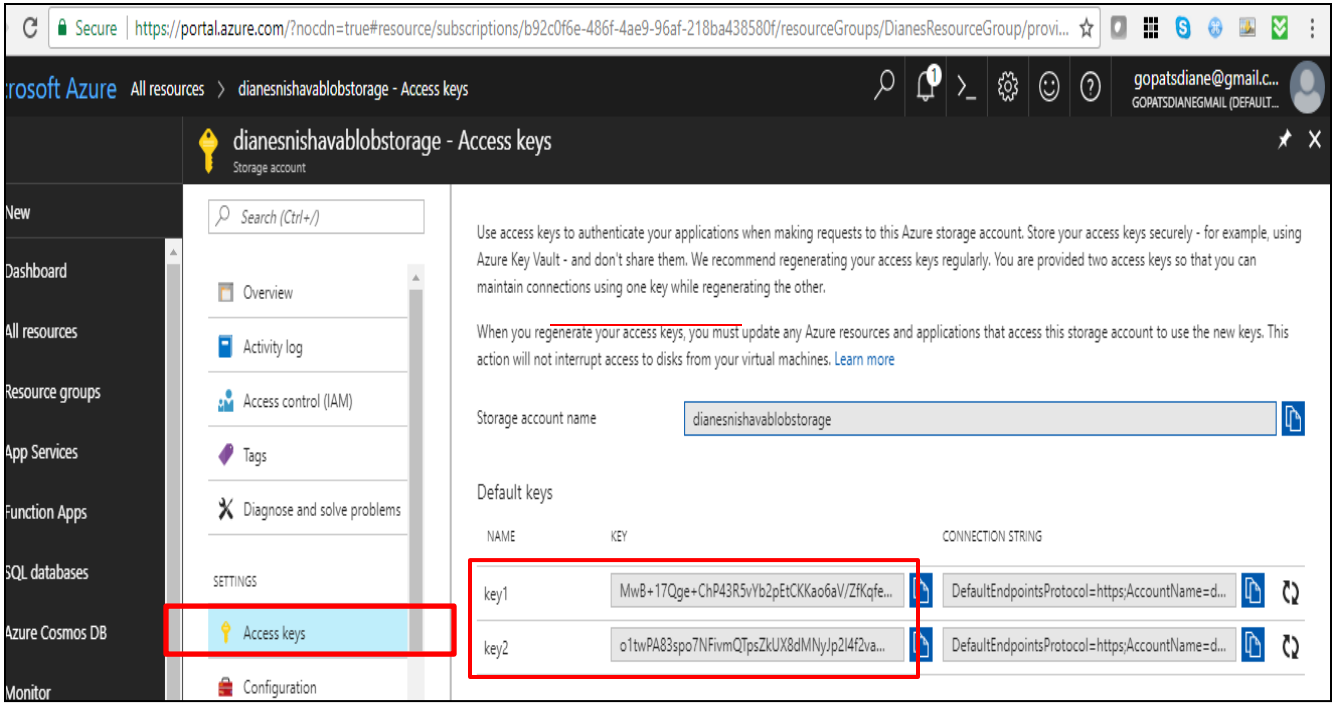

Save Subscription ID, Storage Info and Key in Notepad

- We will need these later in our Python Code.

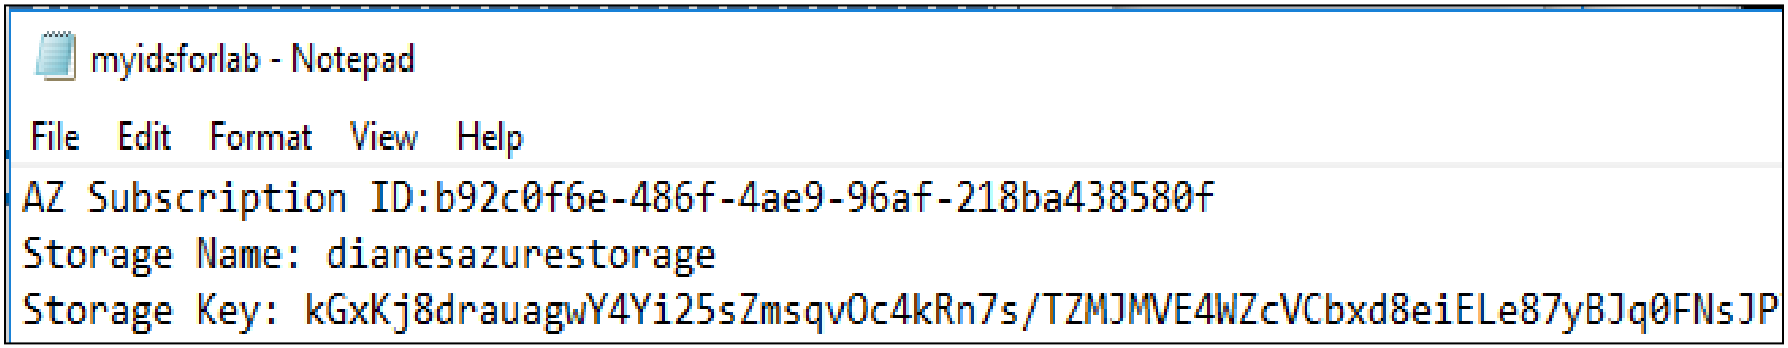

Upload Data File to new Blob Container

employee.txt

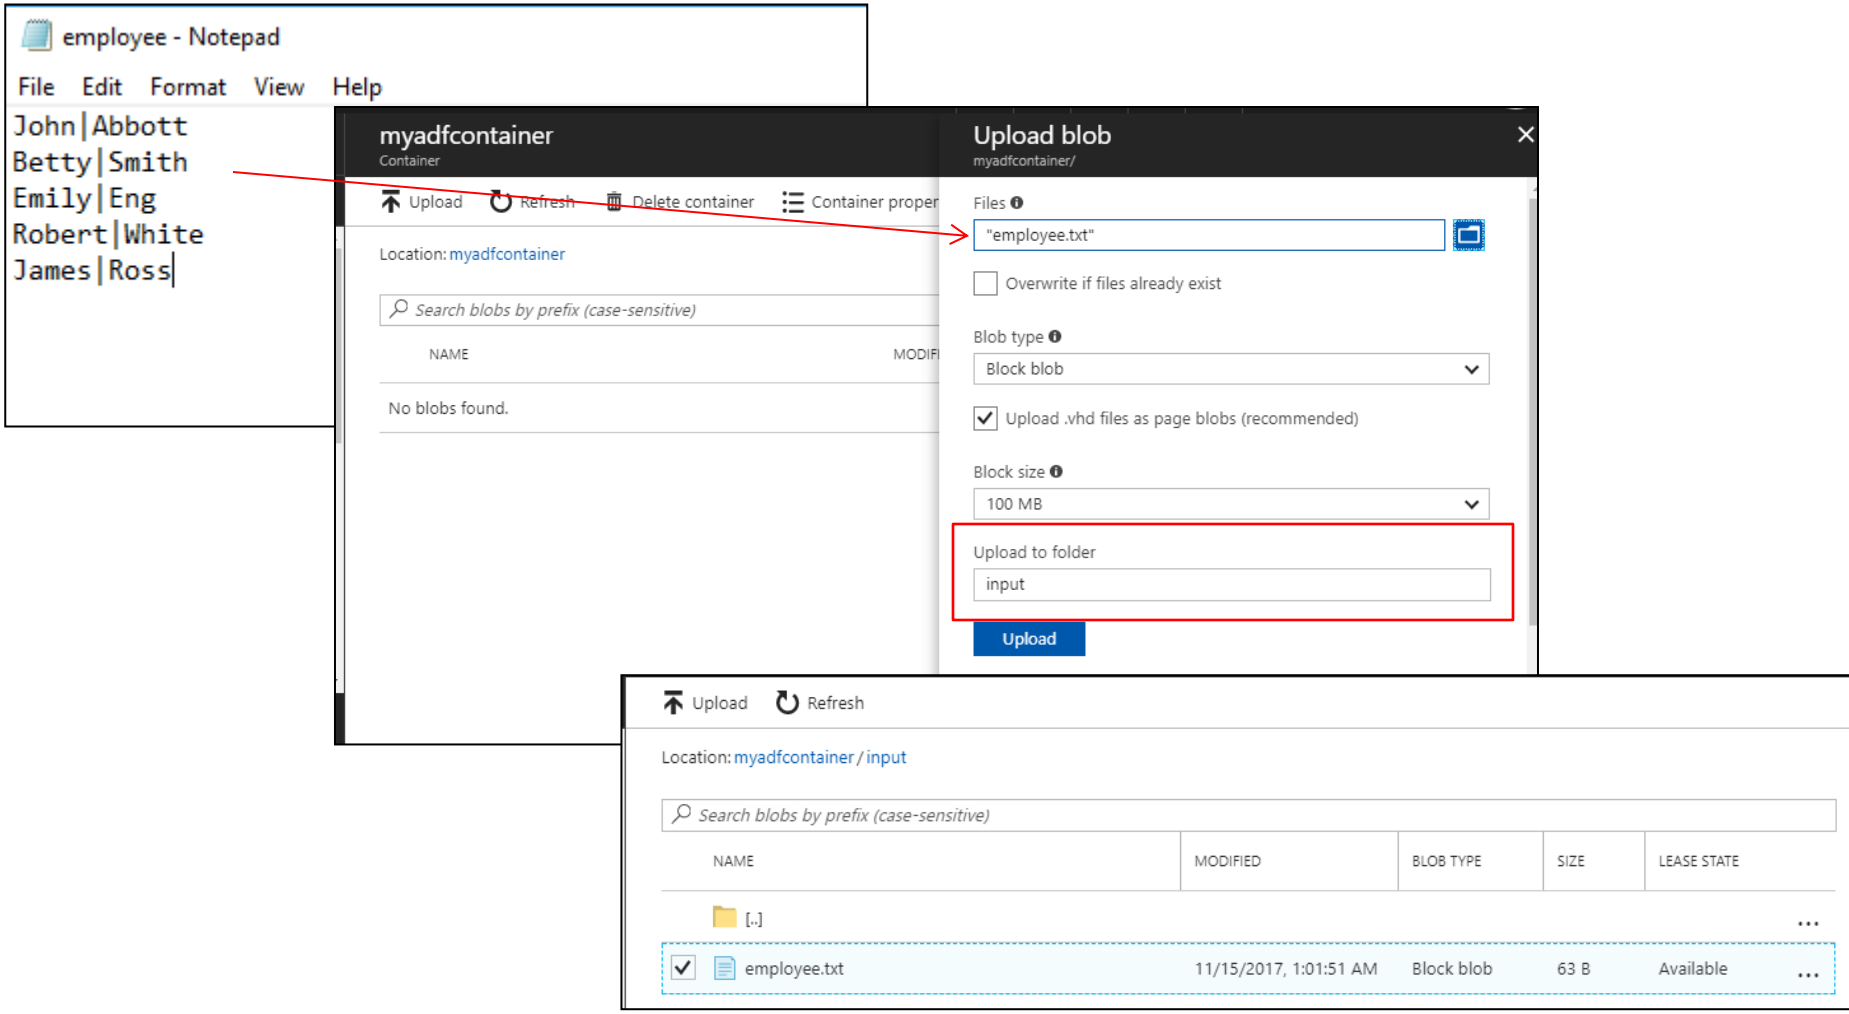

Save Container Name and folder and data file name in Notepad

- We will need these later in our Python Code.

Prerequisites to Create a Data Factory V2

Python 2.7, 3.3, 3.4, 3.5 or 3.6

- ✓ Install Python SDK for Azure packages
  - Azure Management Resources
  - Data Factory
- ✓ Portal: Obtain your Subscription ID
- ✓ Create an Azure Storage Account & Blob container
- ✓ Portal: Create a data file and upload data file to Blob container
- Portal: Create an app in Active Directory (Client ID, Client Secret).
- Python: Create a Data Factory, Linked Service, Pipeline

Check Azure Active Directory Permissions

- Active Directory Permissions allows resources to be created by apps in AZ.
- Go to Active Directory to check if you can register an App.
- Default is set to Yes!

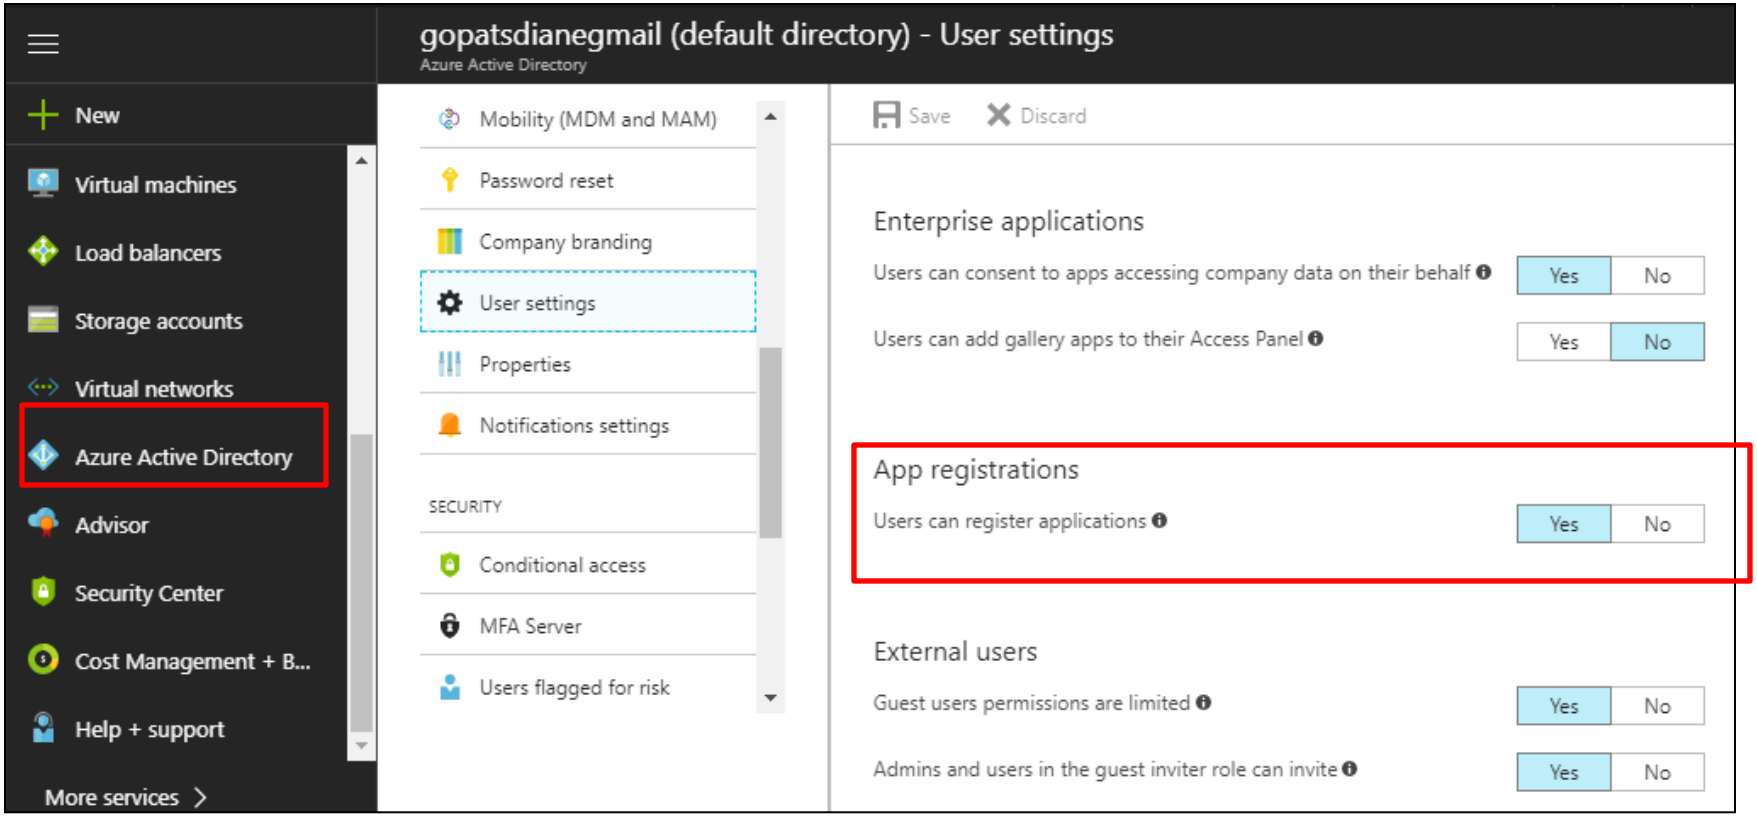

https://docs.microsoft.com/en-us/azure/azure-resource-manager/resource-group-create-service-principal-portal

Create an Azure Active Directory application

- Set up an Azure Active Directory (AD) application and assign permissions.
- Used by an application that will need to access or modify resources.

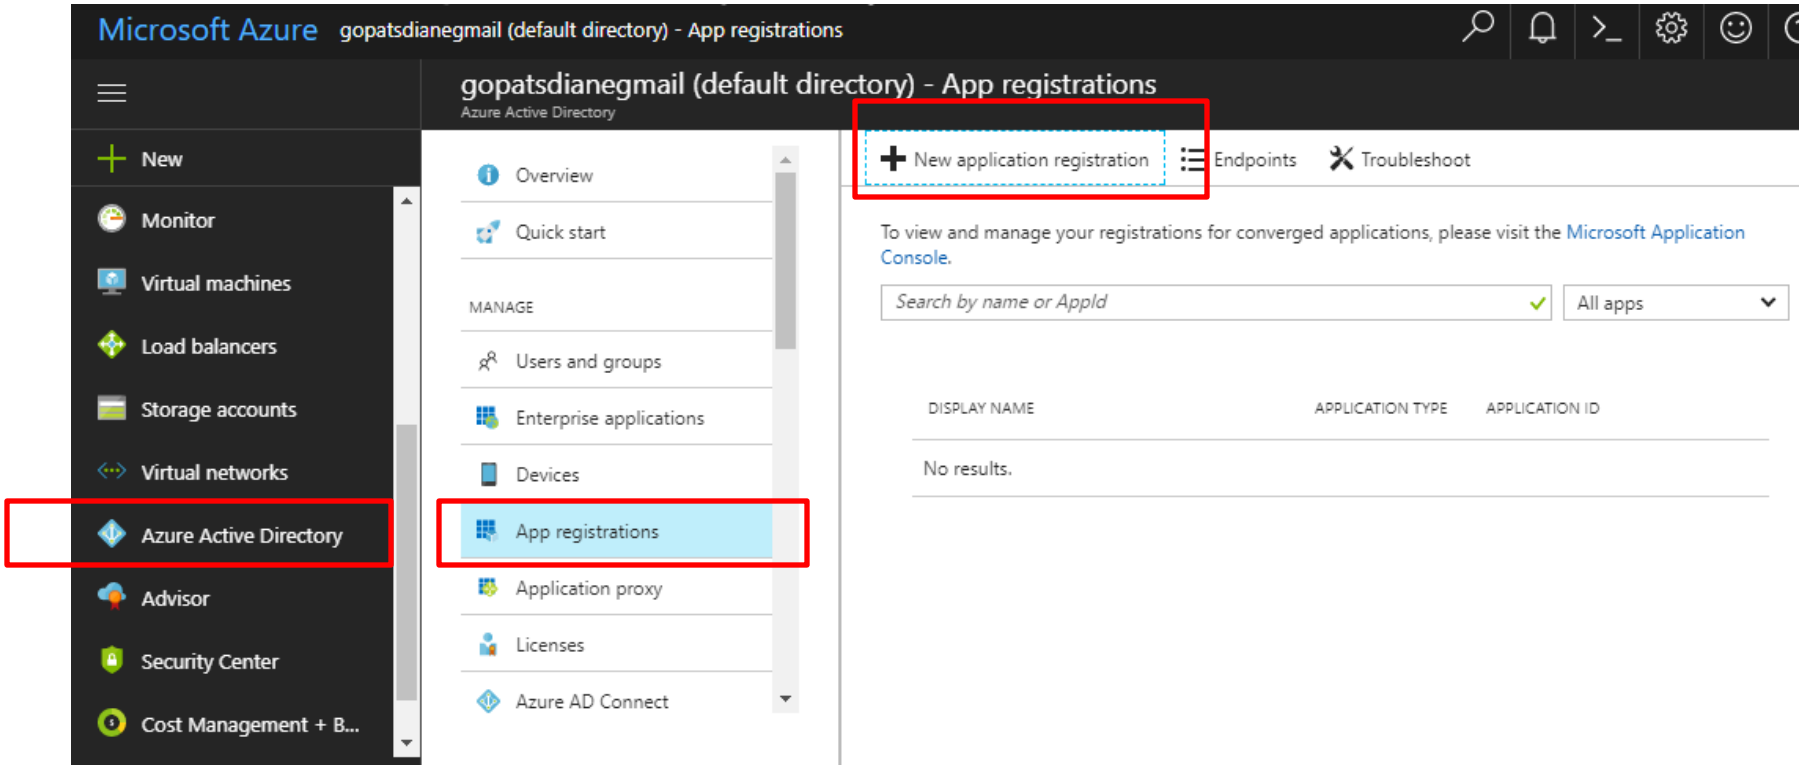

Create an Azure Active Directory application

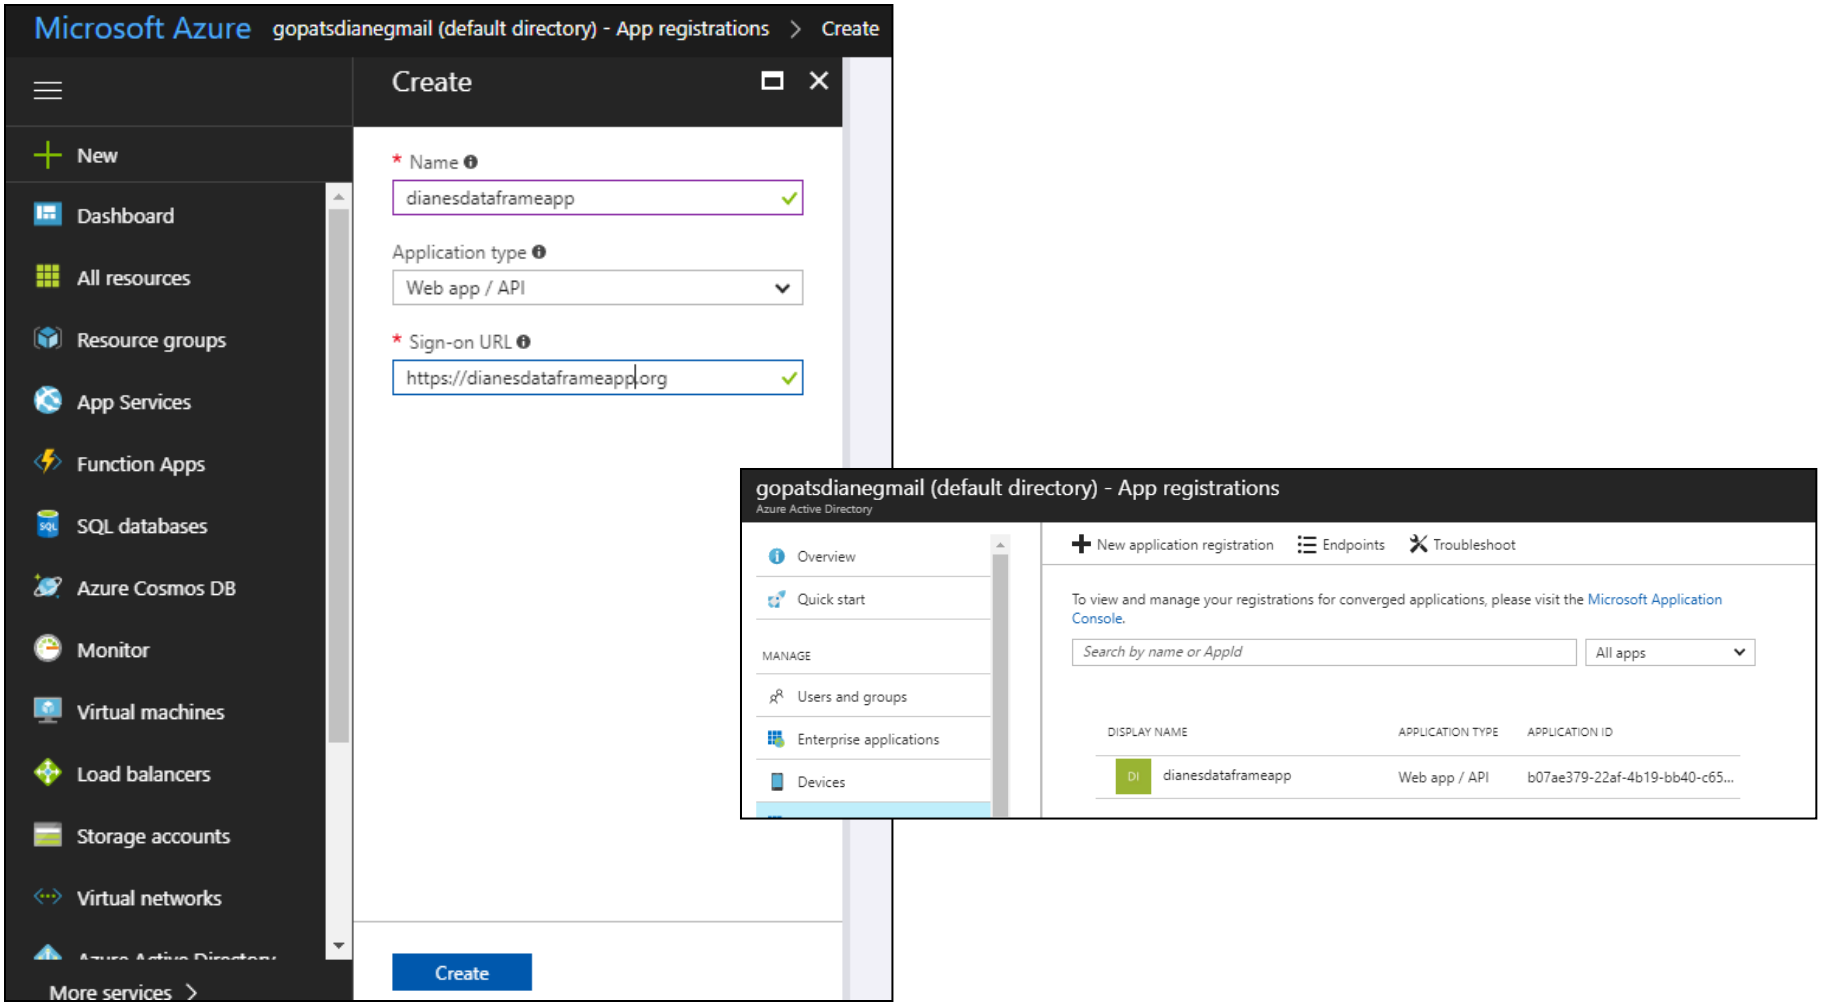

Assign a Role to your new App

- Under New -> More Services -> Subscriptions
- Select your Subscription -> Access Control (IAM) -> Add

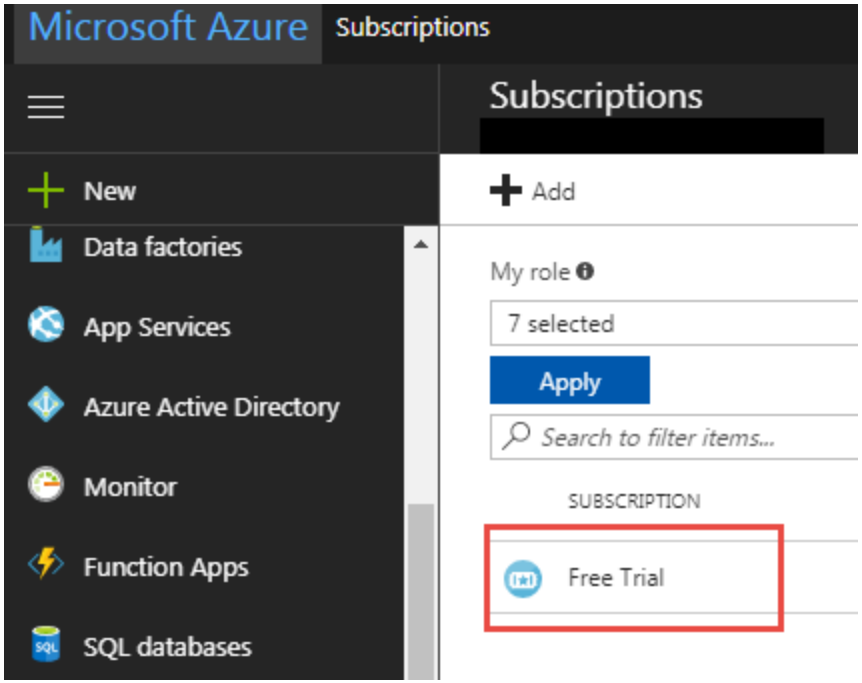

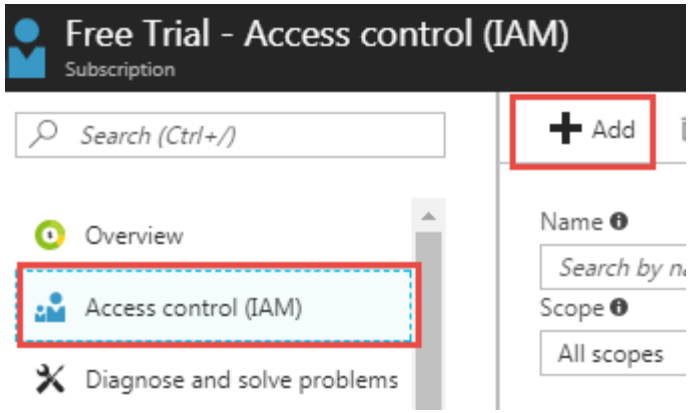

Assign a Role

- Under Role select Contributor. Note:
- You need to have this permission allowed if you are using the McKesson subscription.

- Search for your data frame app, select it and then select Save

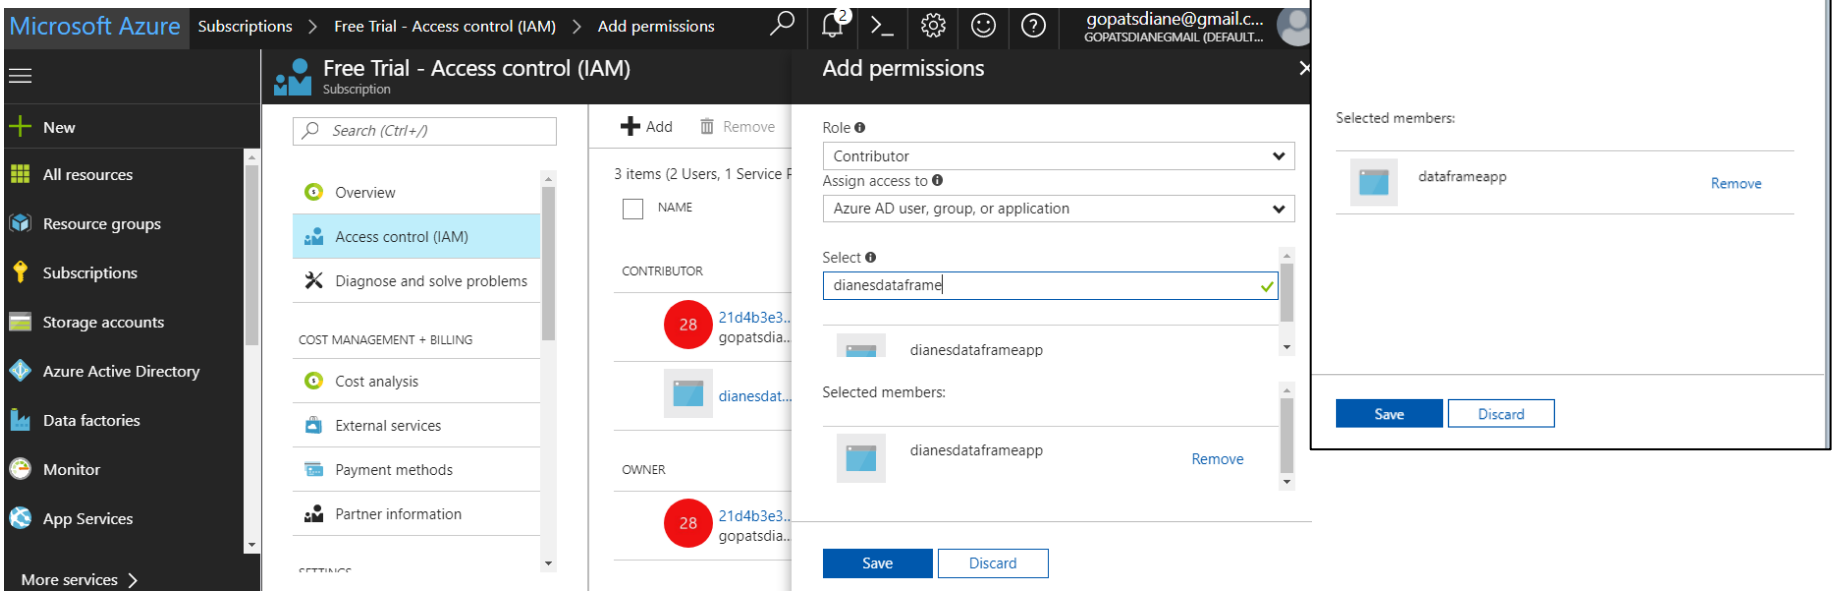

Obtain the Application ID

- Need the Application ID for your Python and .NET code
- Save it into Notepad as you will need to add it to your Python or .NET code.

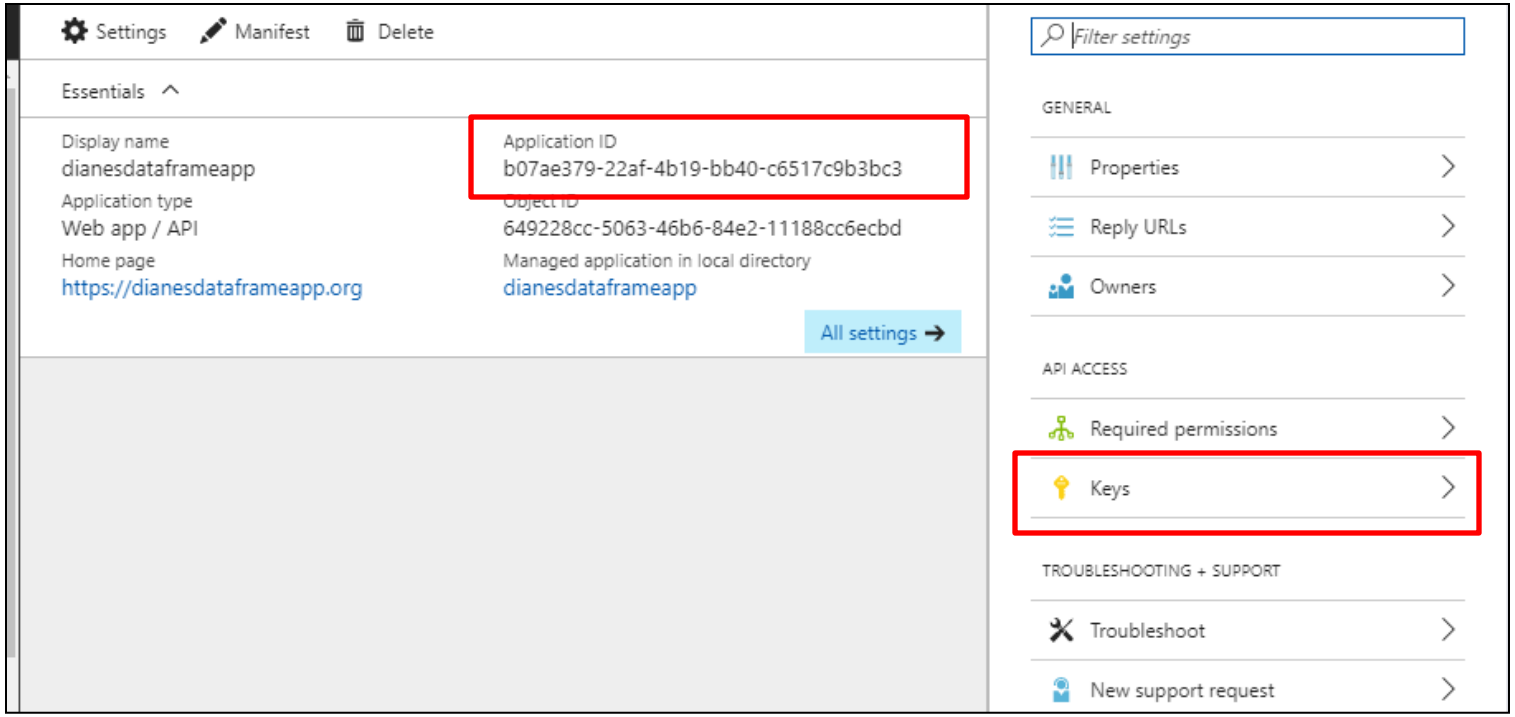

Create Authentication (Secret Key)

- Set up your Secret Key

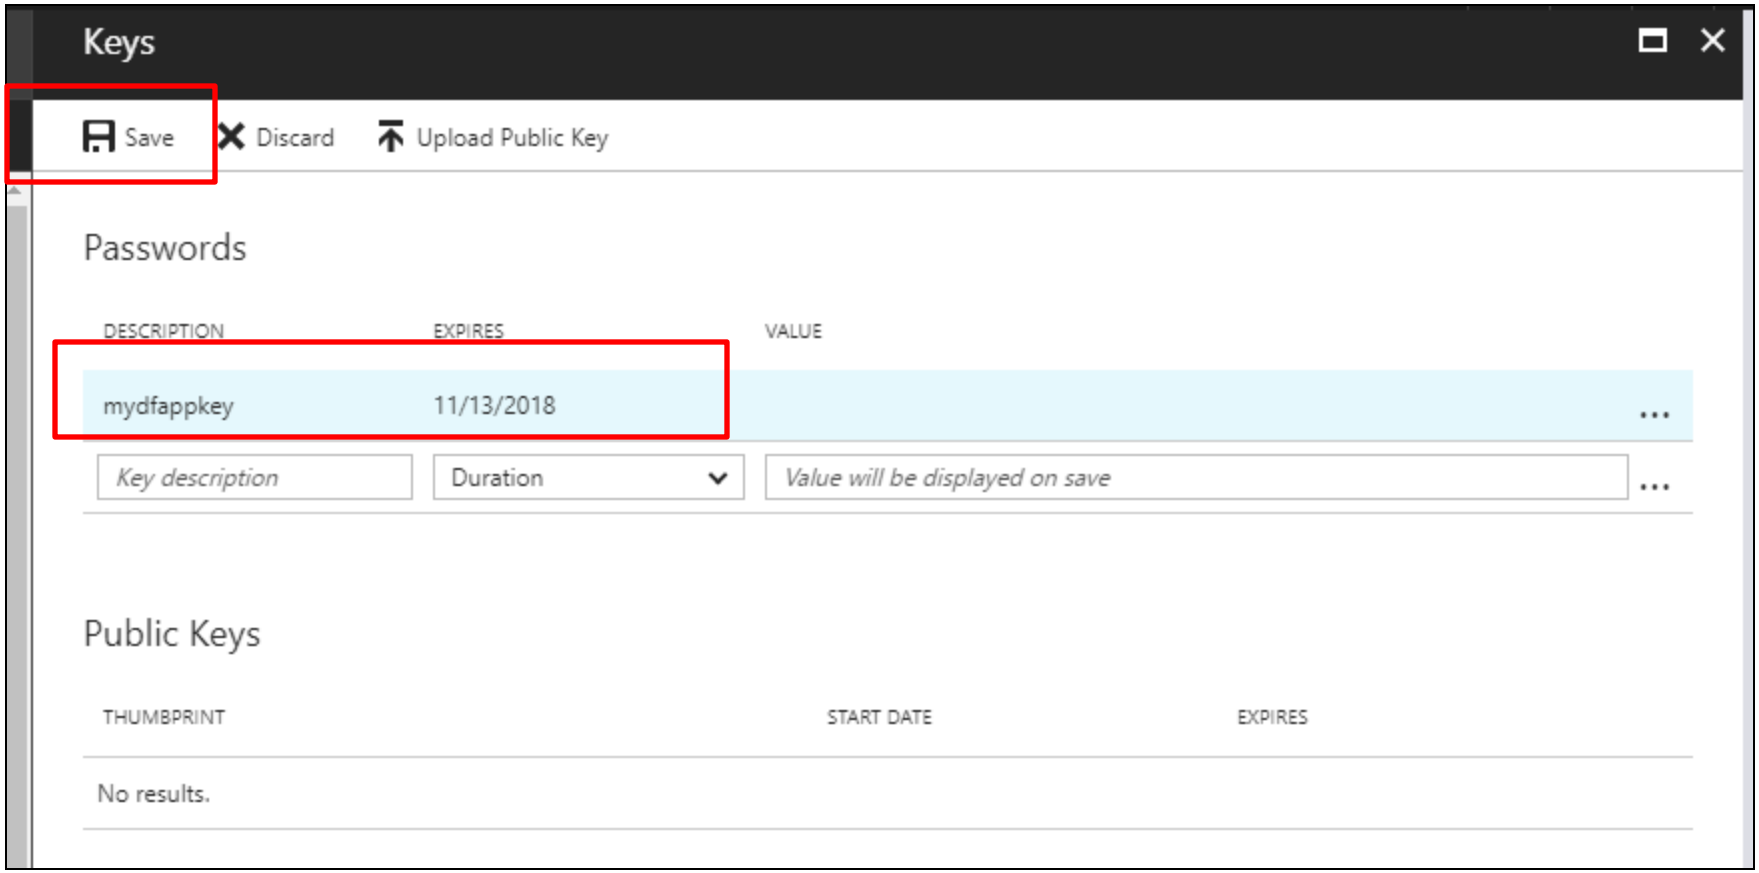

Copy your Secret Key Value

- Appears one time only!
- This is also known as your *Tenant ID*.
- Save your Secret Key in Notepad!

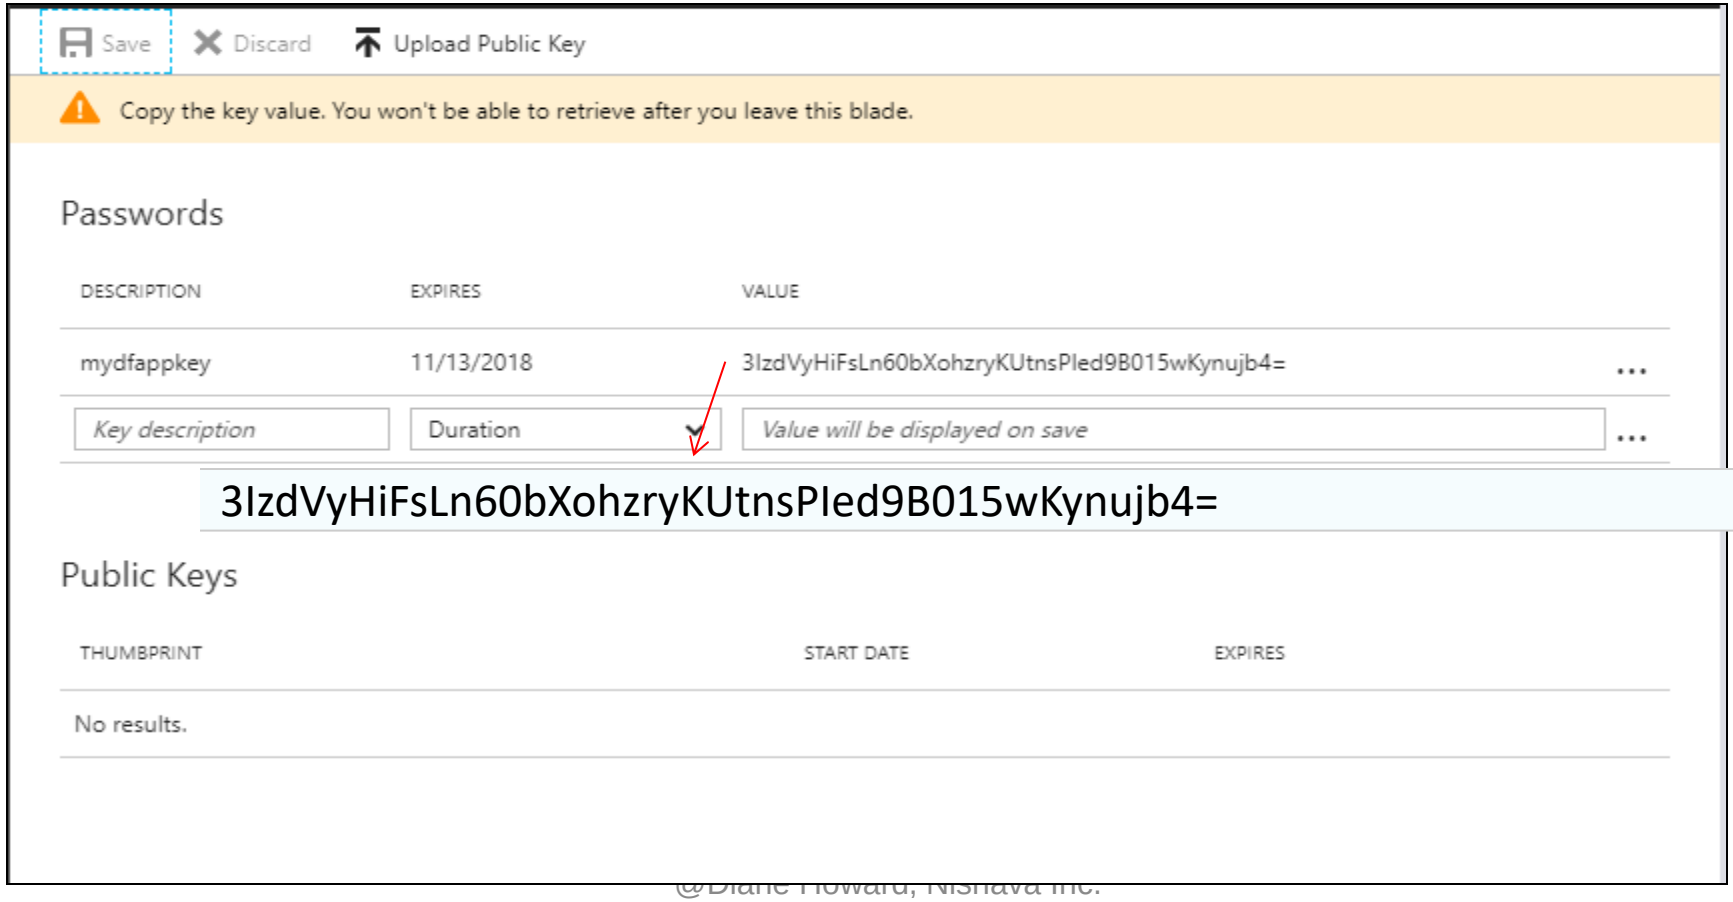

Get your Tenant ID

- The Directory ID = your Tenant ID.
- Copy it to Notepad

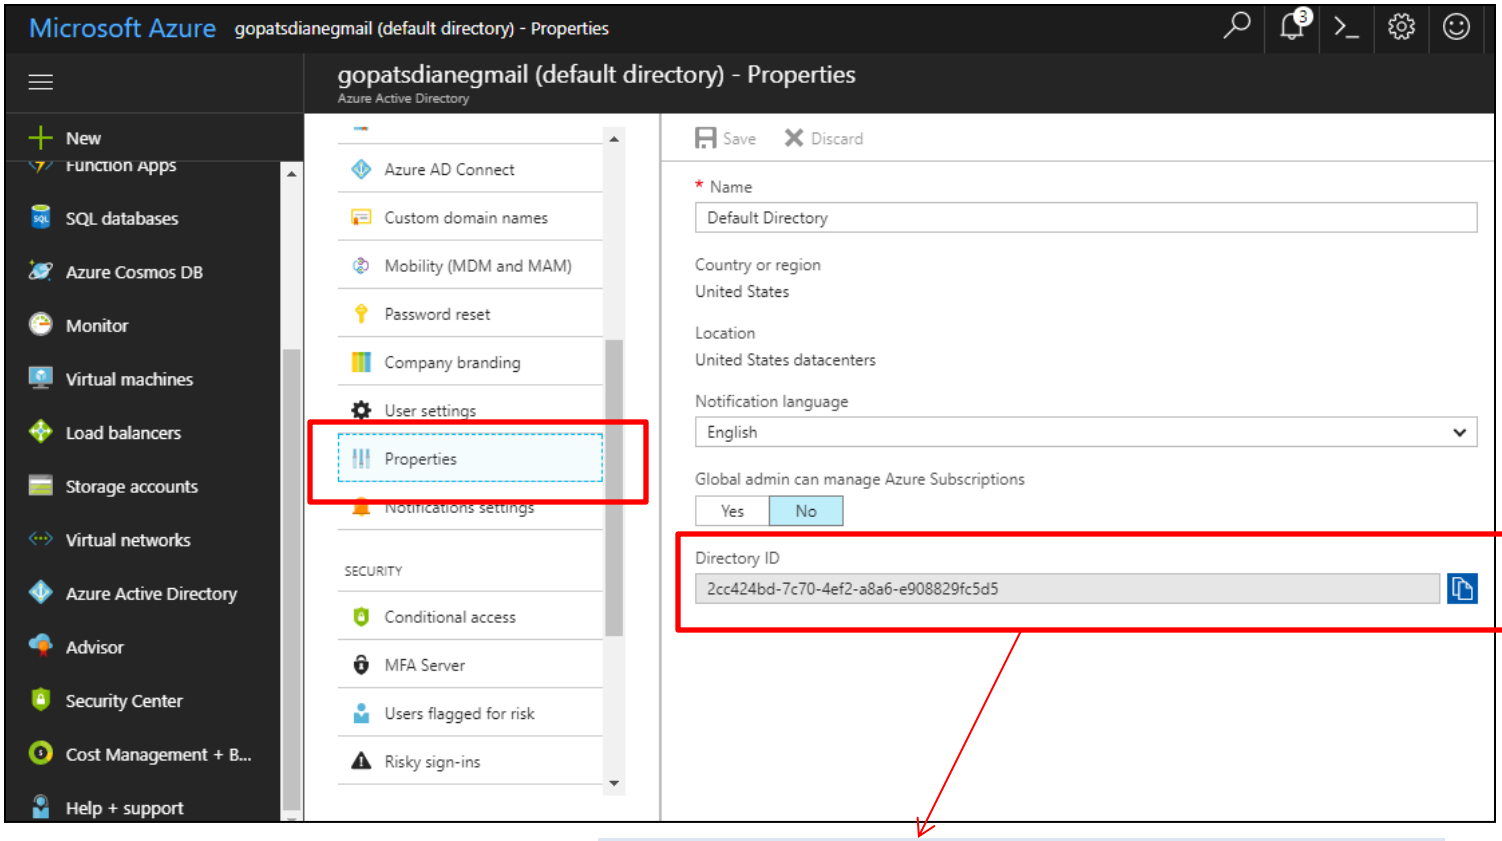

2cc424bd-7c70-4ef2-a8a6-e908829fc5d5

Save in Notepad

- Application ID: b07ae379-22af-4b19-bb40-c6517c9b3bc3
- Client Secret Key: 3IzdVyHiFsLn60bXohzryKUtnsPIed9B015wKynujb4=
- Tenant ID: 2cc424bd-7c70-4ef2-a8a6-e908829fc5d5
- We will need these values later in our Python or .NET code.

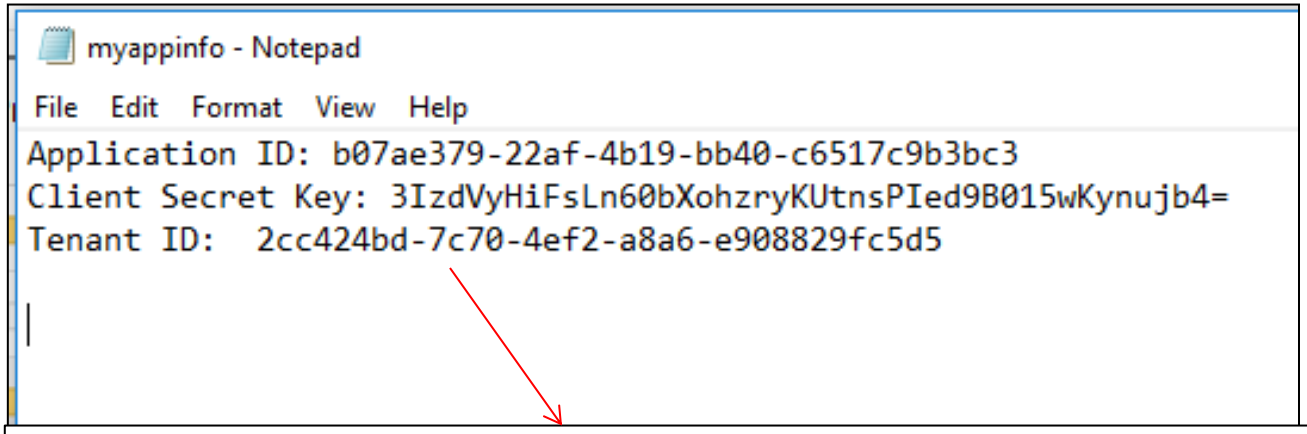

Python Code

```
# Specify your Active Directory client ID, client secret, and tenant ID
    credentials = ServicePrincipalCredentials(client_id='<Active Directory application/client ID>',
secret='<client secret>', tenant='<Active Directory tenant ID>')
    resource_client = ResourceManagementClient(credentials, subscription_id)
    adf_client = DataFactoryManagementClient(credentials, subscription_id)
```

Prerequisites to Create a Data Factory V2

Python 2.7, 3.3, 3.4, 3.5 or 3.6

- ✓ Install Python SDK for Azure packages
  - Azure Management Resources
  - Data Factory
- ✓ Portal: Obtain your Subscription ID
- ✓ Create an Azure Storage Account & Blob container
- ✓ Portal: Create a data file and upload data file to Blob container
- ✓ Portal: Create an app in Active Directory (App ID, Client Secret Key, Directory ID/Tenant ID).
- Python: Create a Data Factory, Linked Service, Pipeline

Python Code

- Add Azure imports

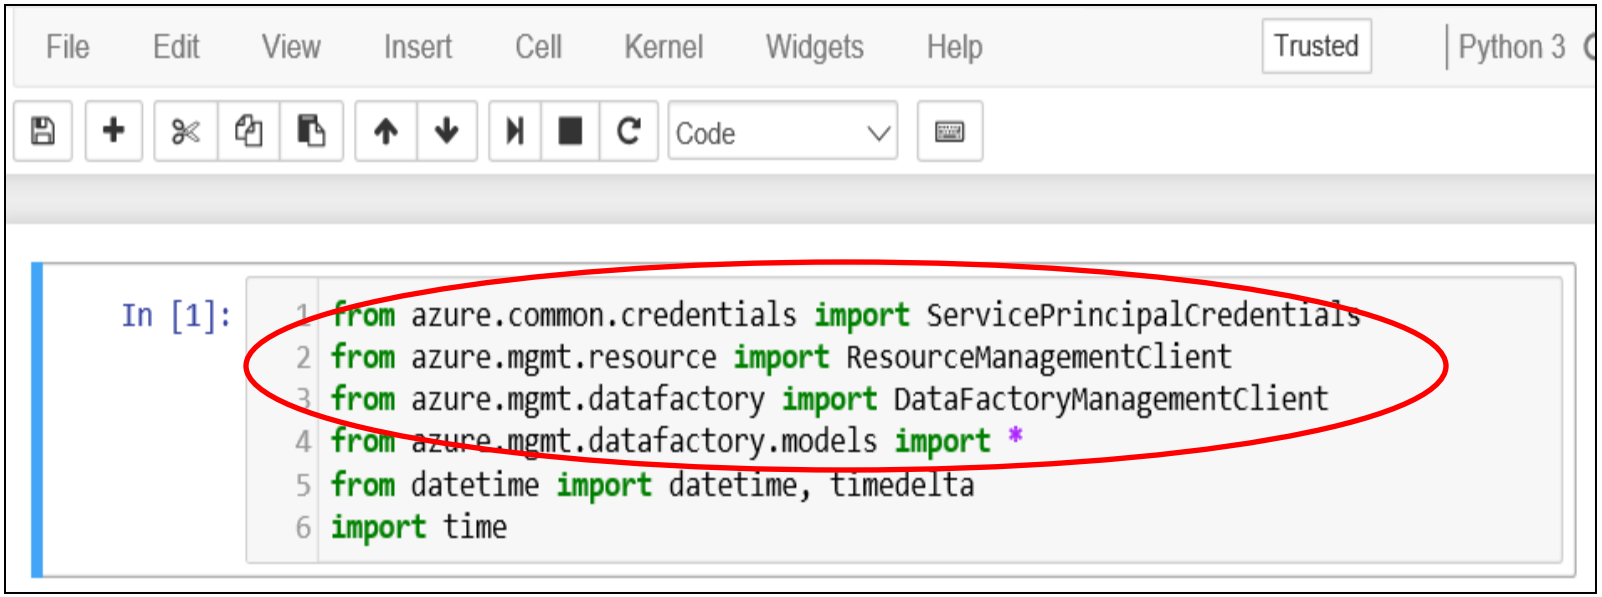

Print Azure Resources & Status

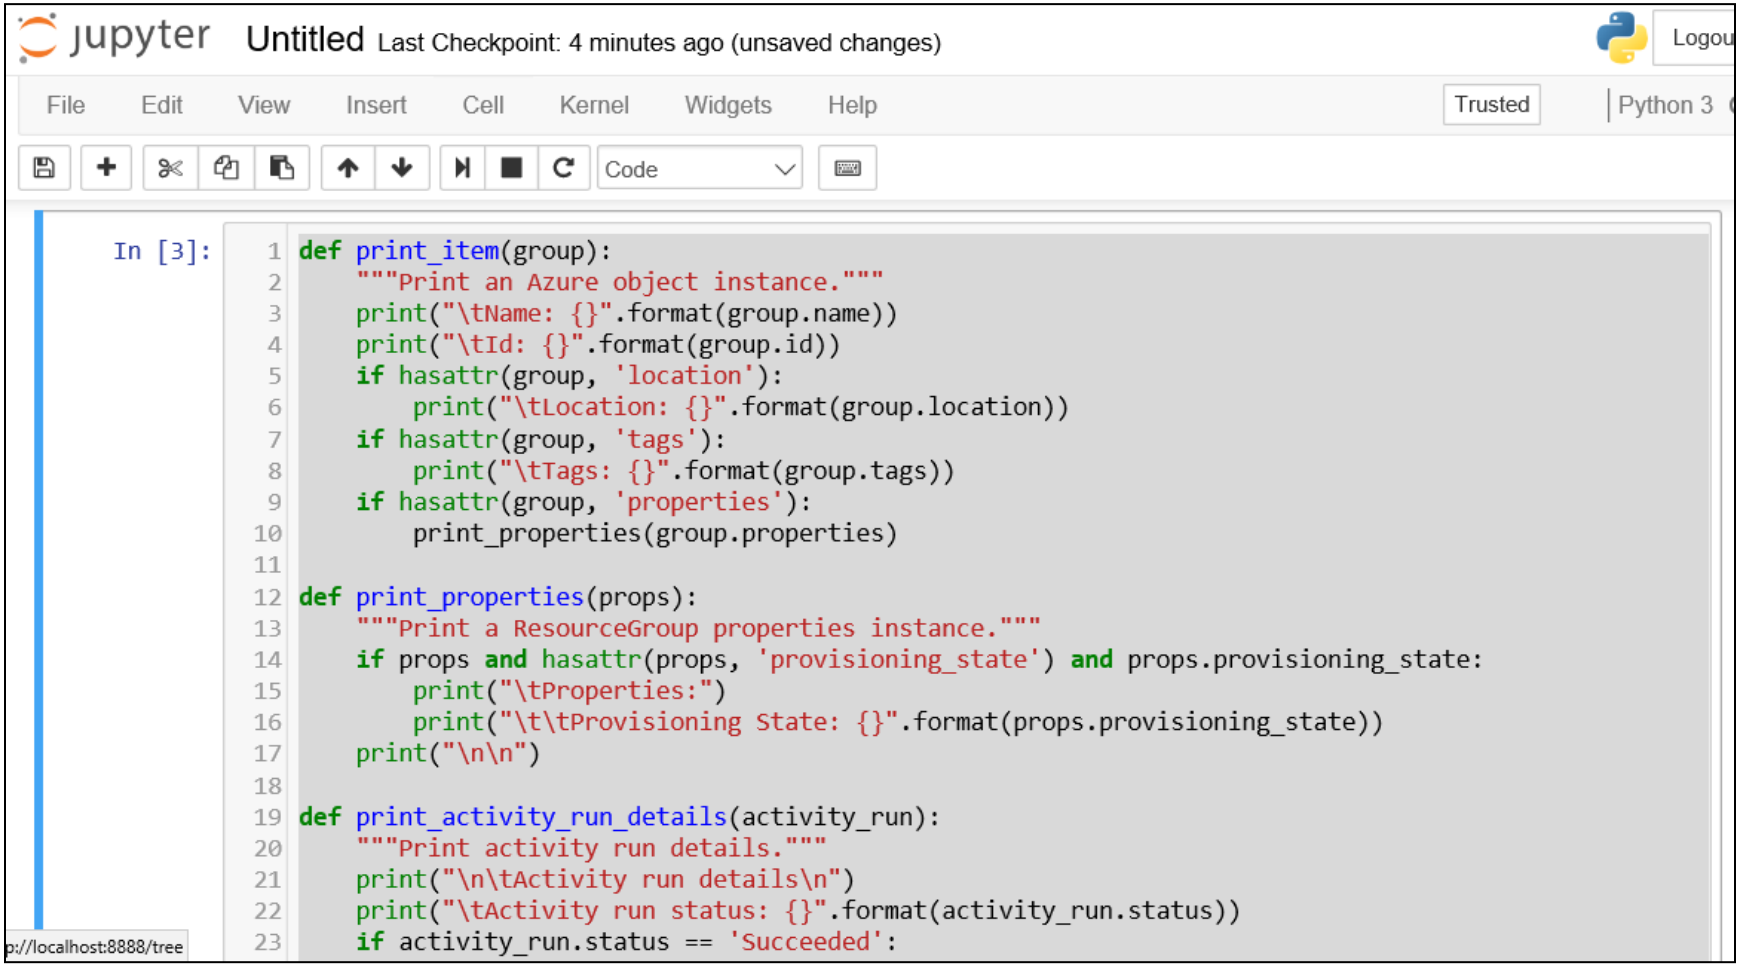

Print AZ Resources and Status - continued

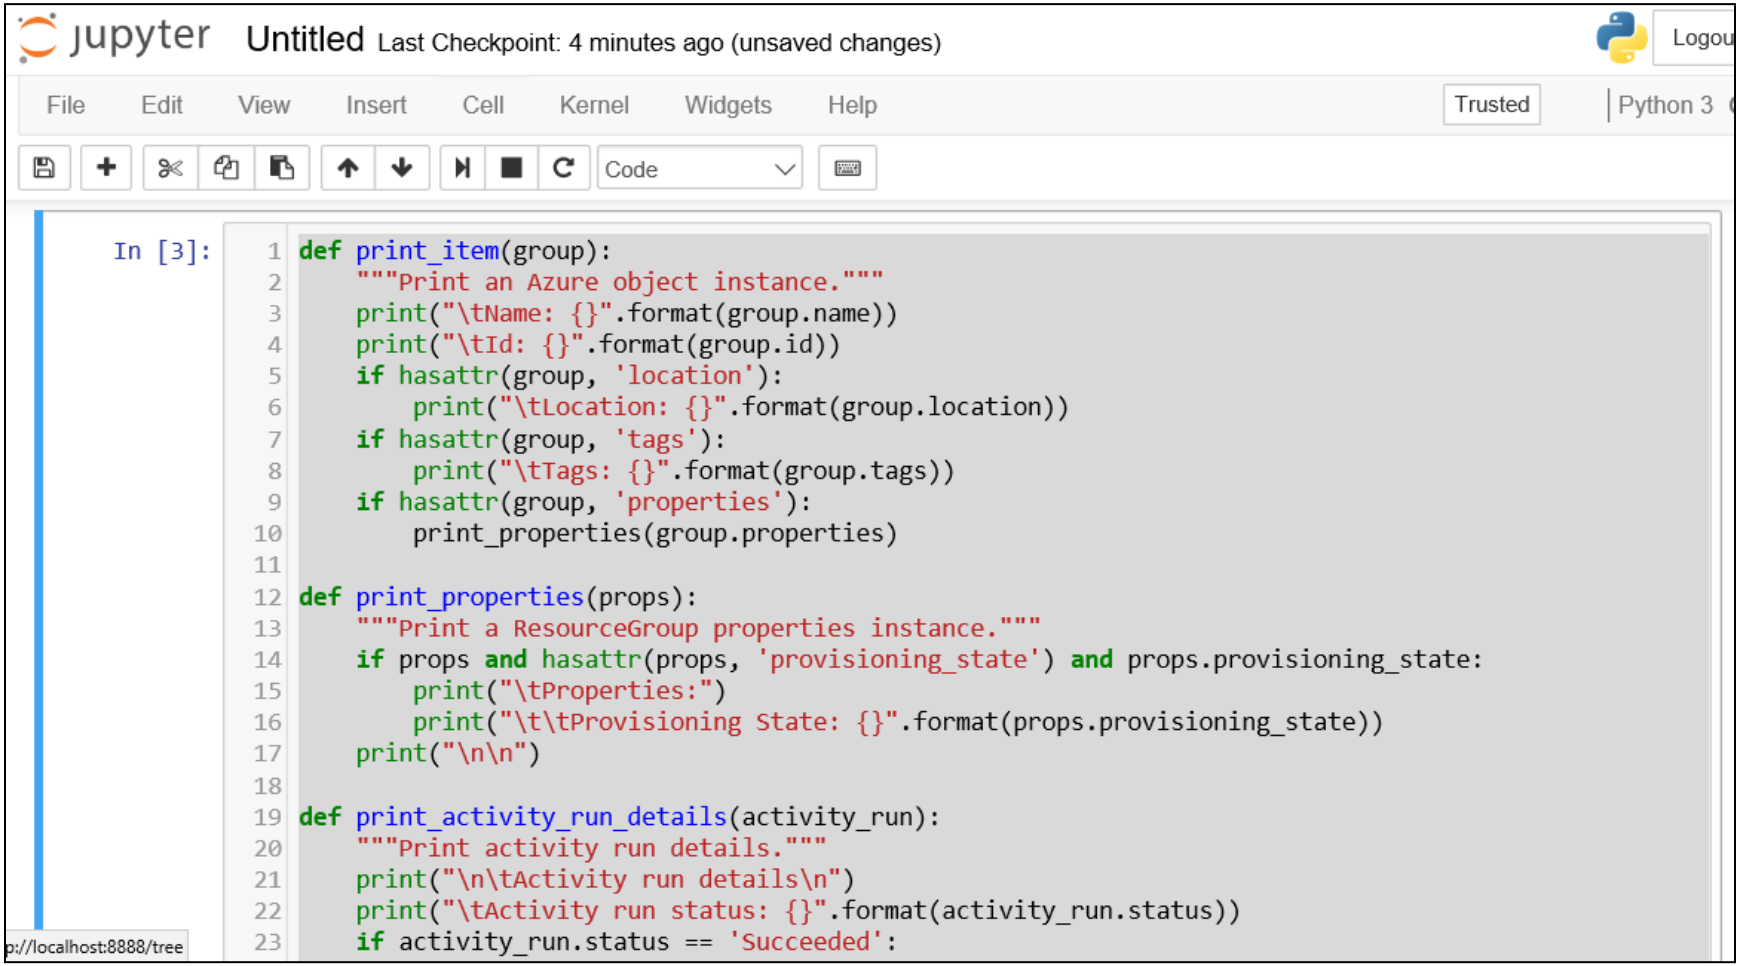

Main: Initialize Variables

- Initialize Resource Group Name, Data Factory Name, Subscription ID and Active Directory credentials

```
48 def main():
49
50 # Azure subscription ID
51     subscription_id = 'b92c0f6e-486f-4ae9-96af-218ba438580f'
52
53 # This program creates this resource group. If it's an existing resource group, comment out the code that creates the resour
54     rg_name = 'DianesRG'
55
56 # The data factory name. It must be globally unique.
57     df_name = 'DianesDF'
58
59 # Specify your Active Directory client ID, client secret, and tenant ID
60     credentials = ServicePrincipalCredentials(client_id='b07ae379-22af-4b19-bb40-c6517c9b3bc3', secret='3IzdVyHiFsLn60bXohzr
61     resource_client = ResourceManagementClient(credentials, subscription_id)
62     adf_client = DataFactoryManagementClient(credentials, subscription_id)
63
64     rg_params = {'location':'eastus'}
65     df_params = {'location':'eastus'}
66
```

Main: Create AZ Resources

- Create Resource Group, Data Factory, and Storage Linked Service

```
67 # Create the resource group
68 # Comment out if the resource group already exits
69     resource_client.resource_groups.create_or_update(rg_name, rg_params)
70
71 # Create a data factory
72     df_resource = Factory(location='eastus')
73     df = adf_client.factories.create_or_update(rg_name, df_name, df_resource)
74     print_item(df)
75     while df.provisioning_state != 'Succeeded':
76         df = adf_client.factories.get(rg_name, df_name)
77         time.sleep(1)
78
79 # Create an Azure Storage linked service
80     ls_name = 'storageLinkedService'
81
```

Main: Create AZ Resources

- Define Input Blob Data Source, Output Blob Sink, Copy Job (Activity)

```
# Create an Azure blob dataset (input)
    ds_name = 'ds_in'
    ds_ls = LinkedServiceReference(ls_name)
    blob_path= 'adfv2tutorial/input'
    blob_filename = 'input.txt'
    ds_azure_blob= AzureBlobDataset(ds_ls, folder_path=blob_path, file_name = blob_filename)
    ds = adf_client.datasets.create_or_update(rg_name, df_name, ds_name, ds_azure_blob)
    print_item(ds)

# Create an Azure blob dataset (output)
    dsOut_name = 'ds_out'
    output_blobpath = 'adfv2tutorial/output'
    dsOut_azure_blob = AzureBlobDataset(ds_ls, folder_path=output_blobpath)
    dsOut = adf_client.datasets.create_or_update(rg_name, df_name, dsOut_name, dsOut_azure_blob)
    print_item(dsOut)

# Create a copy activity
    act_name =  'copyBlobtoBlob'
    blob_source = BlobSource()
    blob_sink = BlobSink()
    dsin_ref = DatasetReference(ds_name)
    dsOut_ref = DatasetReference(dsOut_name)
    copy_activity = CopyActivity(act_name,inputs=[dsin_ref], outputs=[dsOut_ref], source=blob_source, sink=blob_sink)
```

Main: Create AZ Resources

- Create the Pipeline for the Copy Activity, Pipeline Run, Set time delay for the Run, Monitor the Pipeline

```
# Create a pipeline with the copy activity
    p_name =  'copyPipeline'
    params_for_pipeline = {}
    p_obj = PipelineResource(activities=[copy_activity], parameters=params_for_pipeline)
    p = adf_client.pipelines.create_or_update(rg_name, df_name, p_name, p_obj)
    print_item(p)

# Create a pipeline run.
    run_response = adf_client.pipelines.create_run(rg_name, df_name, p_name,
        {
        }
    )

# Monitor the pipeline run
    time.sleep(30)
    pipeline_run = adf_client.pipeline_runs.get(rg_name, df_name, run_response.run_id)
    print("\n\tPipeline run status: {}".format(pipeline_run.status))
    activity_runs_paged = list(adf_client.activity_runs.list_by_pipeline_run(rg_name, df_name, pipeline_run.run_id, datetime
    print_activity_run_details(activity_runs_paged[0])
```

Python Output

```
Name: DianesDF2
        Id: /subscriptions/b92c0f6e-486f-4ae9-96af-
218ba438580f/resourceGroups/dianesrg2/providers/Microsoft.DataFactory/factories/DianesDF2
        Location: eastus
        Tags: {}
        Name: storageLinkedService
        Id: /subscriptions/b92c0f6e-486f-4ae9-96af-
218ba438580f/resourceGroups/DianesRG2/providers/Microsoft.DataFactory/factories/DianesDF2/linkedservices/storageLinkedService

        Name: ds_in
        Id: /subscriptions/b92c0f6e-486f-4ae9-96af-
218ba438580f/resourceGroups/DianesRG2/providers/Microsoft.DataFactory/factories/DianesDF2/datasets/ds_in

        Name: ds_out
        Id: /subscriptions/b92c0f6e-486f-4ae9-96af-
218ba438580f/resourceGroups/DianesRG2/providers/Microsoft.DataFactory/factories/DianesDF2/datasets/ds_out

        Name: copyPipeline
        Id: /subscriptions/b92c0f6e-486f-4ae9-96af-
218ba438580f/resourceGroups/DianesRG2/providers/Microsoft.DataFactory/factories/DianesDF2/pipelines/copyPipeline
*** after run response and before pipeline_run run_response.run_id = 13411ae4-c962-11e7-839d-e006e630d1f8

Datetime with no tzinfo will be considered UTC.
Datetime with no tzinfo will be considered UTC.

        Pipeline run status: Succeeded

        Activity run details

        Activity run status: Succeeded
        Number of bytes read: 18
```

Visual Studio output from ADF run

- Within the portal in SQL databases select your database that you just created.
- Select Tools editor to query your table

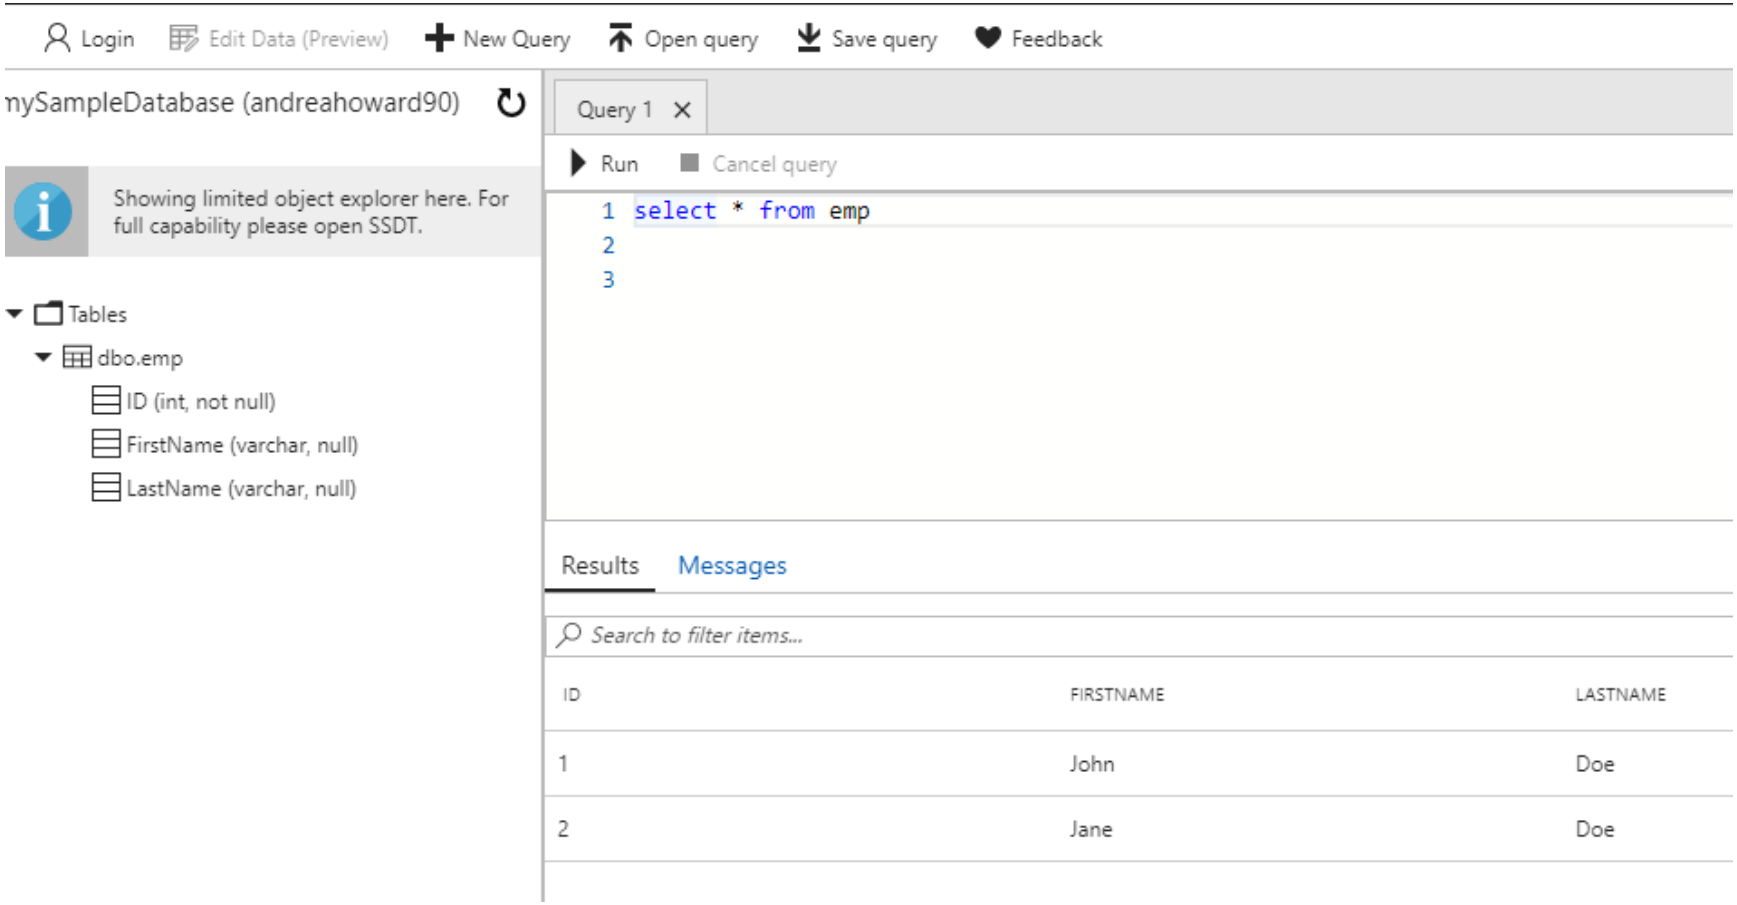

Successful copy of our Workflow

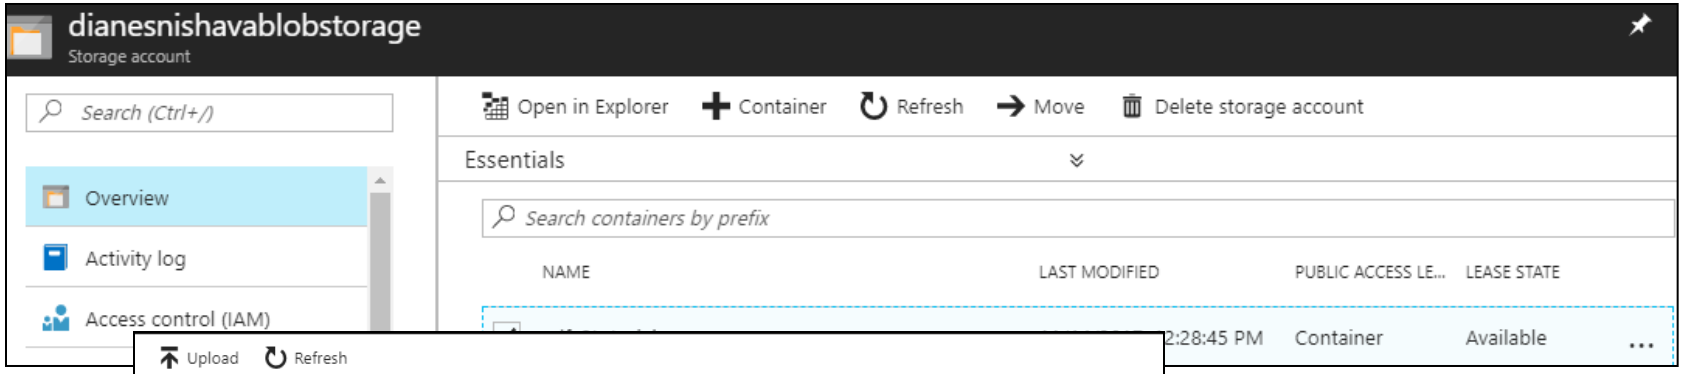

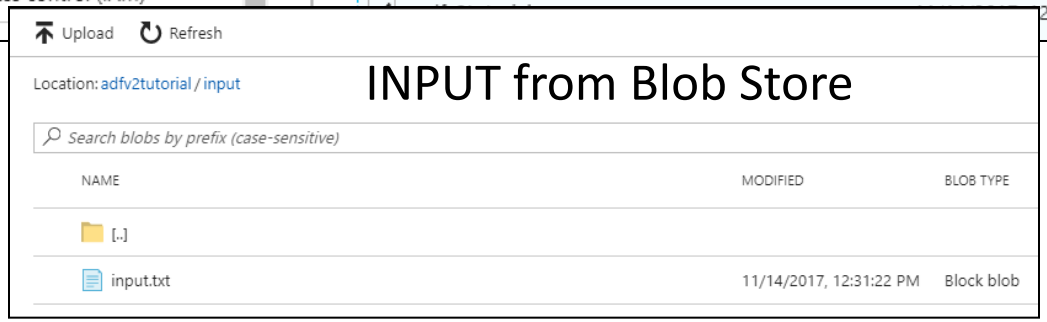

INPUT from Blob Store

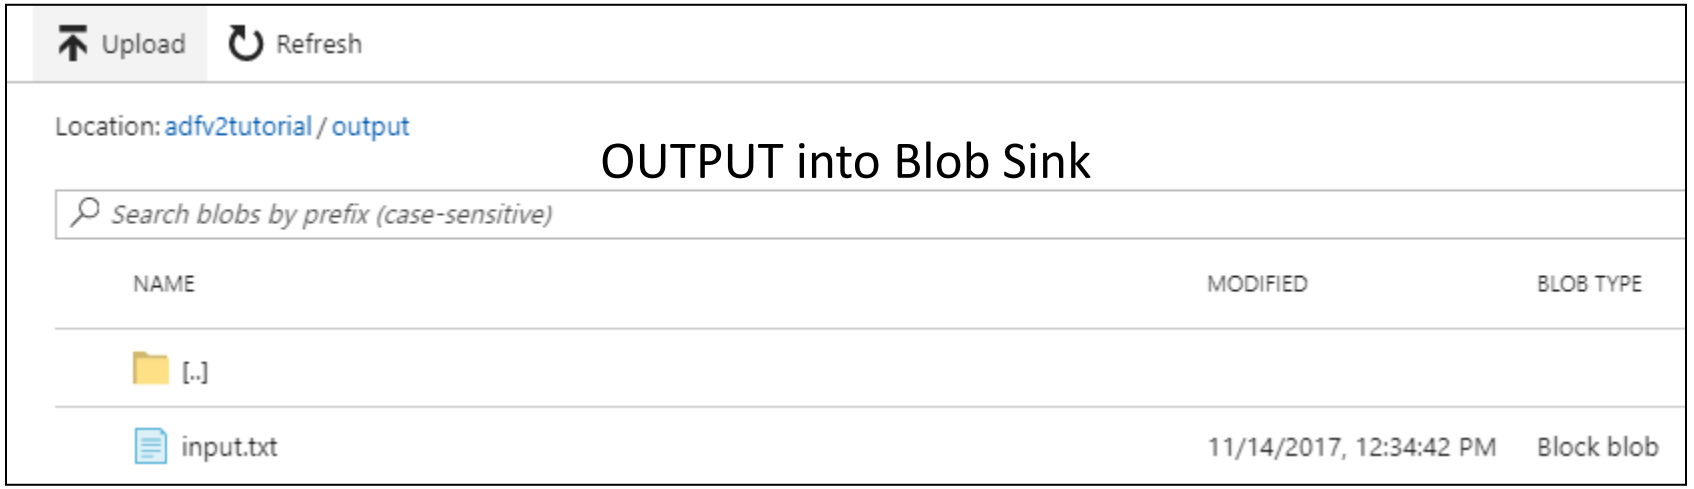

OUTPUT into Blob Sink

Log Activity from Azure Dashboard

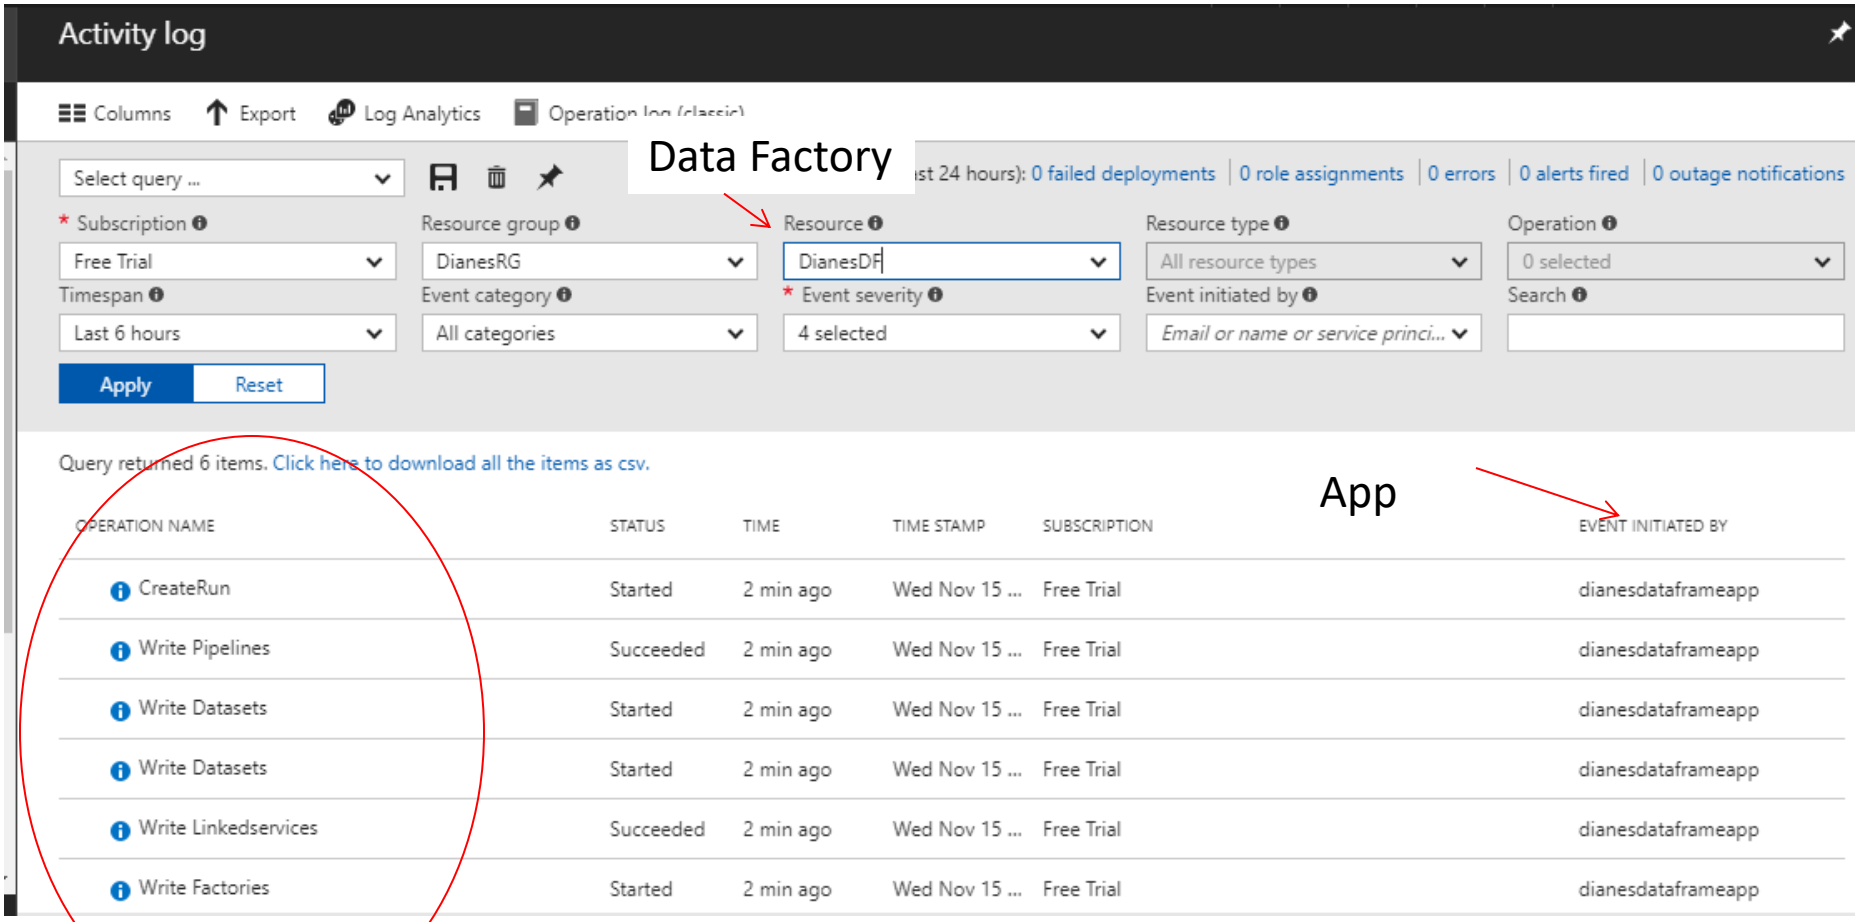

Examination of Activities in Data Factory

- In your Portal you can view the workflow activities (Pipeline runs, Activity Runs, Trigger Runs.

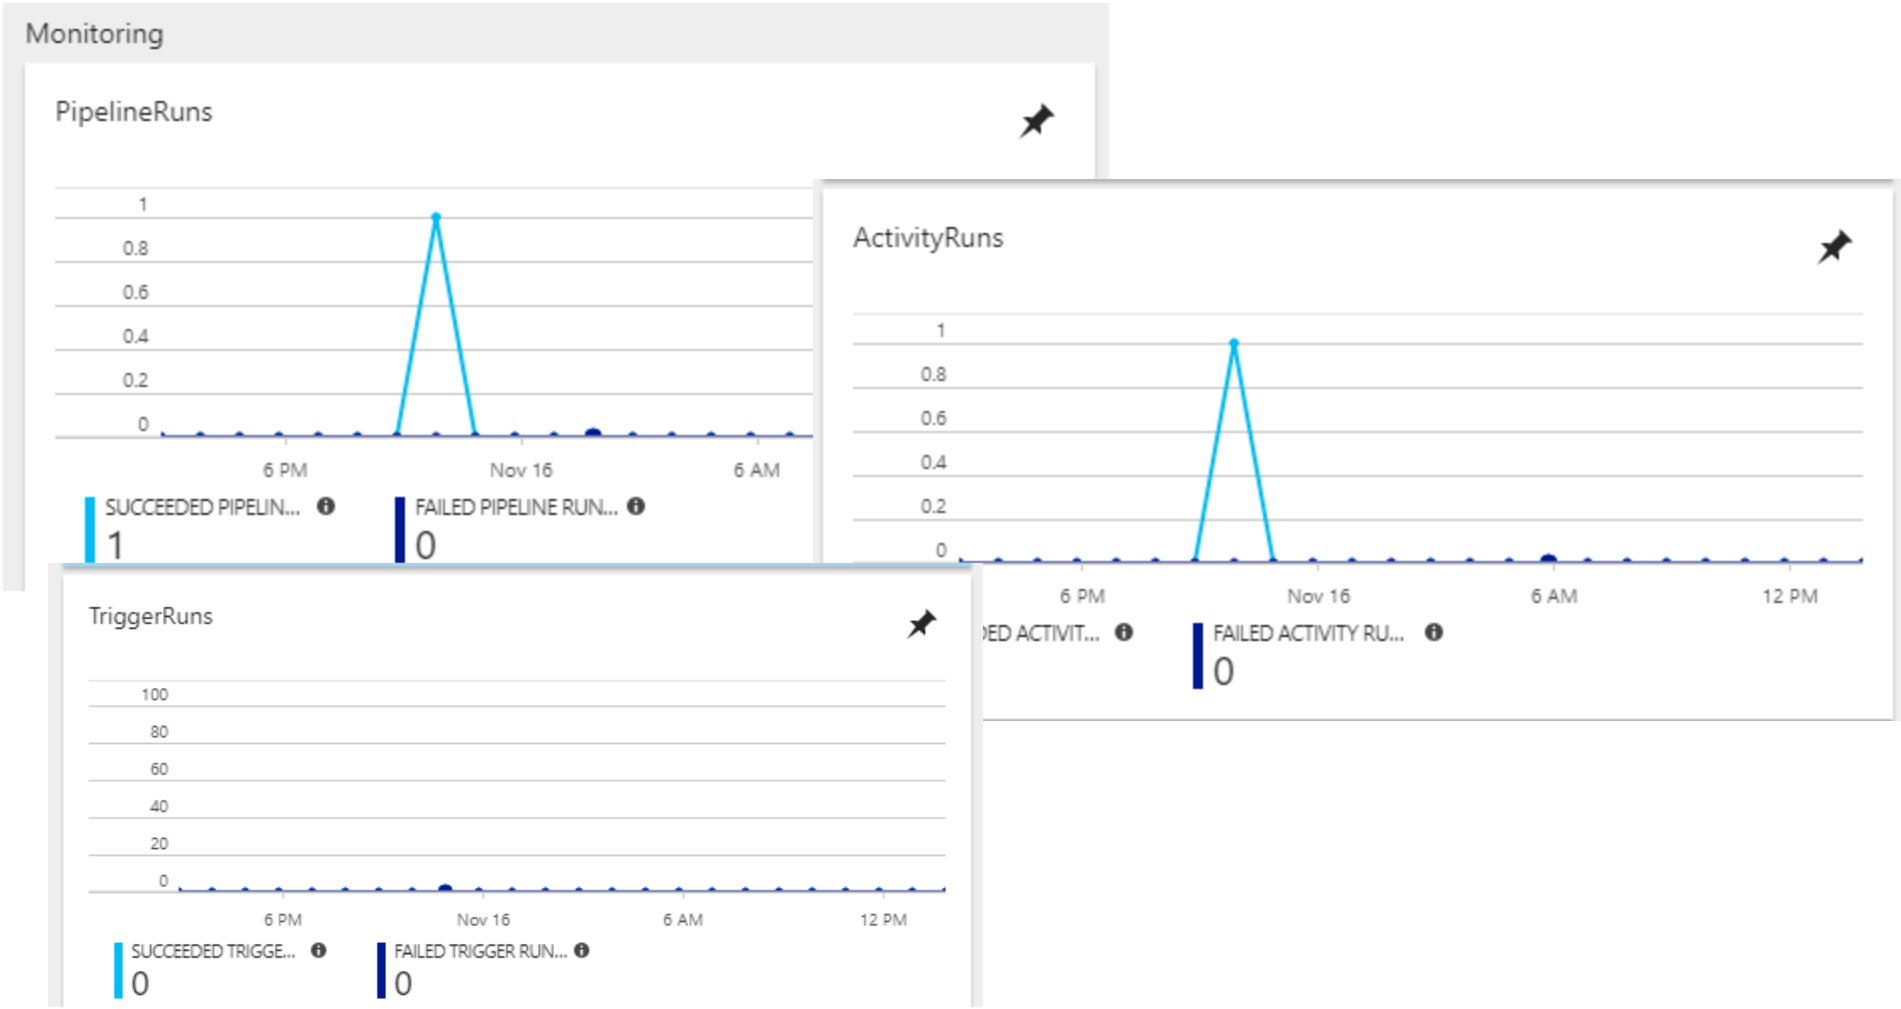

Summary

- We created a simple Data Factory in Python for a small data file which was uploaded to a Blob Store and watched the progress of the job: *Copy the data file to our Blob Sink*.
- AZ Resources Needed:
  - Subscription Info (ID)
  - Active Directory (Client ID, Client Secret Key) Storage Container (Name, ID)
  - Resource Group (Name)
- We did not perform any manipulation and analytics analysis of our data which is the heart of using Data Factory.

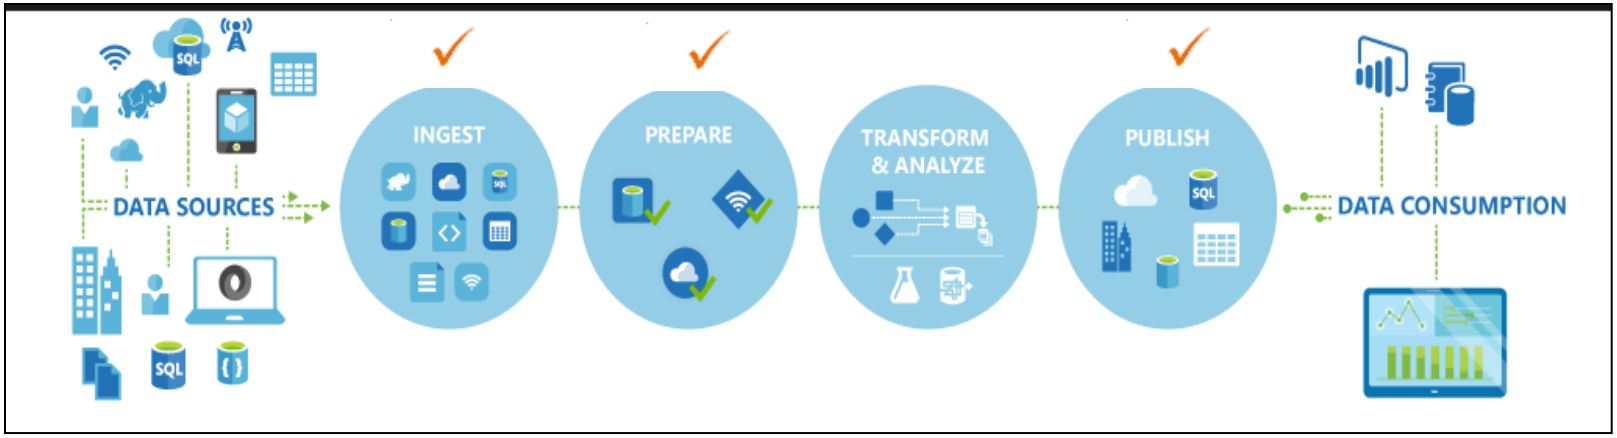

My Cost to Run the Data Factory Demo

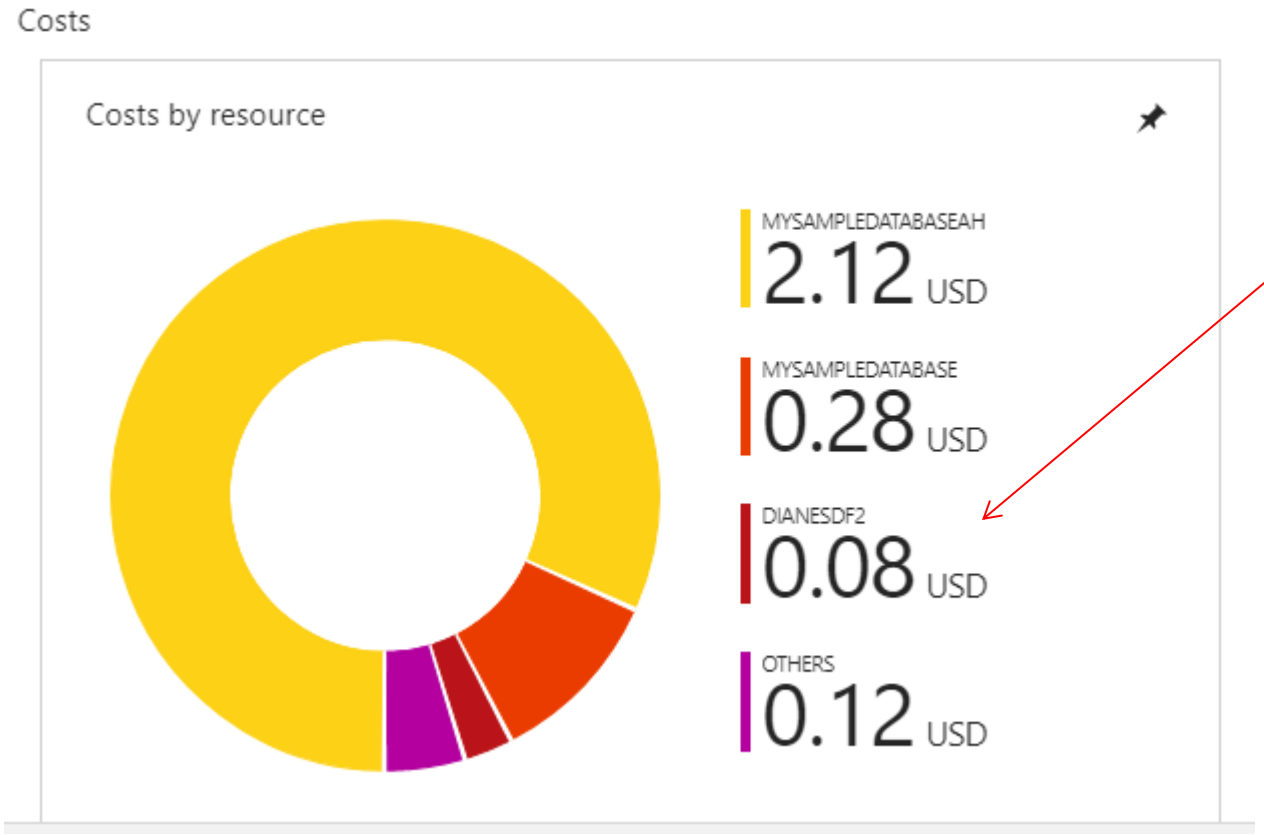

Errors from run

```
        Pipeline run status: Failed

        Activity run details

        Activity run status: Failed
        Errors: Failure happened on 'Source' side. ErrorCode=UserErrorSourceBlobNotExist,'Type=Microsoft.DataTransfer.Common.Sh
ared.HybridDeliveryException,Message=The required Blob is missing. ContainerName: https://dianesazurestorage.blob.core.windows.
net/adfv2tutorial&#44; ContainerExist: False&#44; BlobPrefix: input.txt&#44; BlobCount: 0.,Source=Microsoft.DataTransfer.Client
Library,'
```

Error in not defining the container name properly.

- Recommendations: Check how you defined your container vs. folder.

Issue with Credentials

```
---------------------------------------------------------------------------
CloudError                                Traceback (most recent call last)
<ipython-input-1-23807d1d8df9> in <module>()
    131
    132 # Start the main method
--> 133 main()

<ipython-input-1-23807d1d8df9> in main()
     67 # Create the resource group
     68 # Comment out if the resource group already exits
---> 69     resource_client.resource_groups.create_or_update(rg_name, rg_params)
     70
     71 # Create a data factory

~\Anaconda3_64bit\lib\site-packages\azure\mgmt\resource\resources\v2017_05_10\operations\resource_groups_operations.py in creat
e_or_update(self, resource_group_name, parameters, custom_headers, raw, **operation_config)
    145             exp = CloudError(response)
    146             exp.request_id = response.headers.get('x-ms-request-id')
--> 147             raise exp
    148
    149         deserialized = None

CloudError: Azure Error: AuthorizationFailed
Message: The client 'af54c570-7988-4719-9790-65681c0ebcc9' with object id 'af54c570-7988-4719-9790-65681c0ebcc9' does not have
authorization to perform action 'Microsoft.Resources/subscriptions/resourcegroups/write' over scope '/subscriptions/b92c0f6e-48
6f-4ae9-96af-218ba438580f/resourcegroups/DianesRG'.
```

- Check your container ID – was it copied correctly?

Can't run Pip

- Check your directory where python is installed

>where python
C:\Users\dhoward\anaconda3_64bit\python.exe

- Go to the directory where python is installed

>cd c:\users\dhoward\anaconda3_64bit

- Go to the Scripts directory where pip resides

>cd Scripts
c:\Users\dhoward\Anaconda3_64bit\Scripts>dir pip*

```
 Volume in drive C is OS
 Volume Serial Number is 2426-663B
 Directory of c:\Users\dhoward\Anaconda3_64bit\Scripts
09/25/2017  03:52 PM               197 pip-script.py
09/19/2017  08:10 AM            40,960 pip.exe
               2 File(s)         41,157 bytes
               0 Dir(s)  84,455,149,568 bytes free
```

- Run PIP from this directory (Scripts) to install your packages

>pip install azure-mgmt-resource
Requirement already satisfied: azure-mgmt-resource in c:\users\dhoward\anaconda3_64bit\lib\site-packages