

# HDInsight Hadoop Cluster, HDFS, Spark, Hive

## Lab 12

Deep.Azure@McKesson

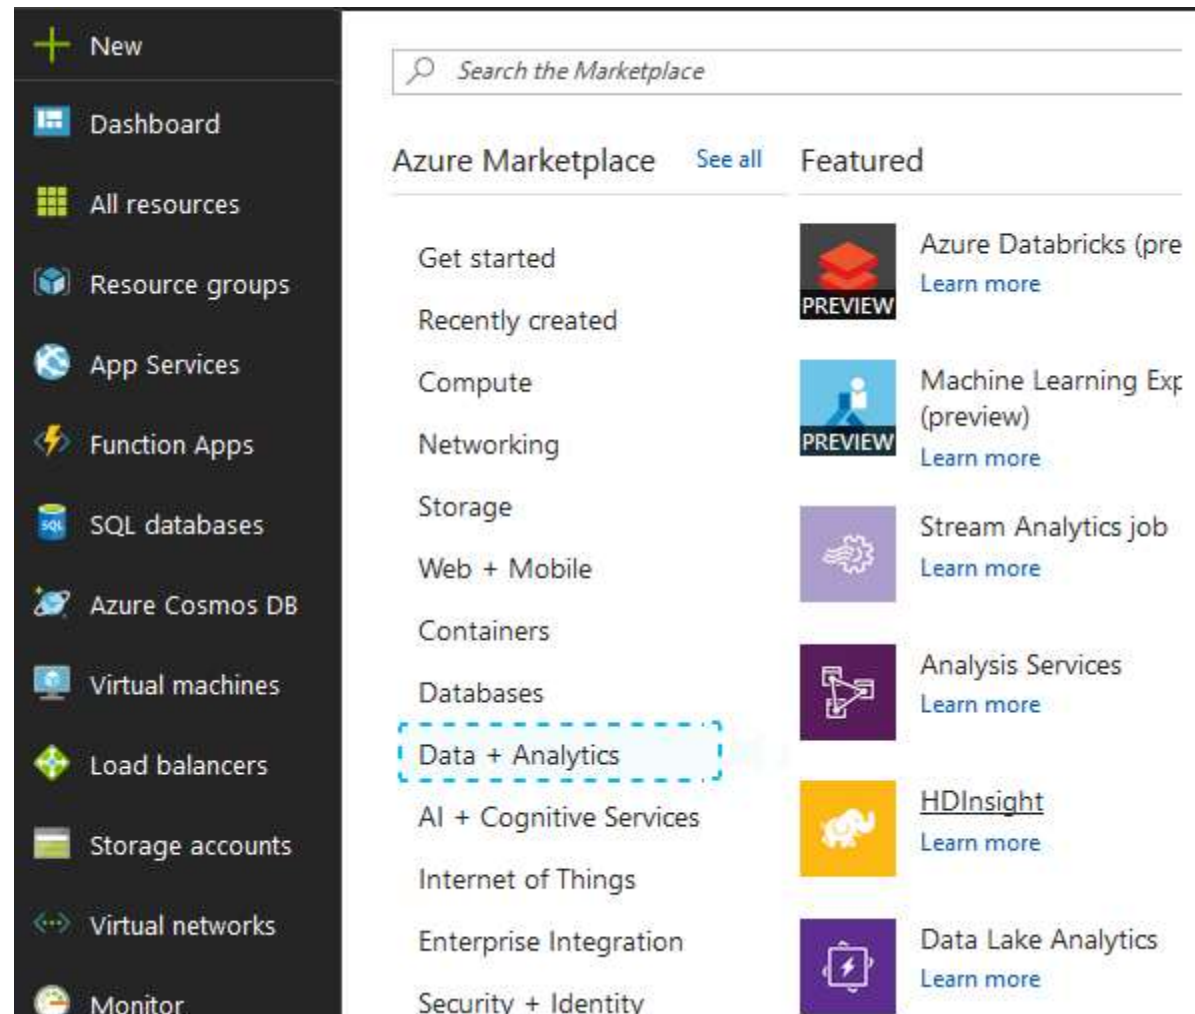
Zoran B. Djordjević

# HDInsight Apps run on Clusters

- Hadoop and most applications and frameworks in Hadoop ecosystem require clusters of machines to run and scale successfully.
- Typically, you can run Hadoop, Spark, Storm, Kafka and most other applications on a cluster with a single node. You do such "pseudo-cluster" runs for development and testing purposes.
- Let us create a HDInsight cluster that supports Hadoop.

# HDInsight Cluster

- In Azure Portal select +New, then Data + Analytics,
- Subsequently select HDInsight



# Basic and Cluster Configuration

- Enter cluster name, cluster username and password, secure shell username, subscription, resource group and most importantly cluster type. Select **Cluster type**
- Cluster type can be: Hadoop, Hbase, Storm, Spark, R Server, Kafka & Interactive Query
- Select Hadoop. Hit **Select**.

The screenshot displays the HDInsight configuration interface, divided into three main sections: a left sidebar, a 'Basics' tab, and a 'Cluster configuration' tab.

**Left Sidebar:**

- Buttons: 'Quick create' and 'Custom (size, settings, apps)'.
- Step 1: Basics (Configure basic settings) - Active.
- Step 2: Storage (Set storage settings).
- Step 3: Summary (Confirm configurations).
- Info: 'This cluster may take up to 20 minutes to create.'

**Basics Tab:**

- \* Cluster name: 'zjdjrdjehadoop'.
- \* Subscription: 'Pay-As-You-Go'.
- \* Cluster type: 'Hadoop' (with a link to 'Configure required settings').
- \* Cluster login username: 'admin'.
- \* Cluster login password: '\*\*\*\*\*' (with a green checkmark).
- Secure Shell (SSH) username: 'sshuser'.
- ☒ Use same password as cluster login.
- \* Resource group: 'zjdjrdjehadoopgrp' (with a green checkmark).
- \* Location: 'East US'.
- Info: 'Click here to view cores usage.'
- Next button.


**Cluster configuration Tab:**

- Learn about HDInsight and cluster versions. →
- Cluster configuration**
- \* Cluster type: 'Hadoop'.
- \* Operating system: 'Linux' (selected) and 'Windows'.
- \* Version: 'Hadoop 2.7.3 (HDI 3.6)'.
- ☐ Enterprise Security Package.
- Hadoop:** Petabyte-scale processing with Hadoop components like MapReduce, Hive (SQL on Hadoop), Pig, Sqoop and Oozie. If you are looking for Hive using LLAP, please create an Interactive Query cluster.
- Features**
- \* denotes preview feature
- + Secure shell (SSH) access
- + HDInsight applications
- + Custom virtual network
- + Custom Hive metastore
- + Custom Oozie metastore
- + Data Lake Store access
- + Data Lake Store as primary data storage
- Select button.

# View cores usage

- Before hitting Next, Click to view core usage. It appears that Azure allows you to build clusters with up to 60 cores without asking for more. We will not ask.
- We will try to reduce the number of cores.
- Hit Next in Basic setup

Each of your subscriptions has a per-location quota on the number of cores that HDInsight clusters can consume. If you'd like to increase the core quota in a location, please [request billing support](#).



LOCATION	CORES IN USE	AVAILABLE CORES	TOTAL CORES
Australia East	0	60	60
Australia Southeast	0	60	60
Brazil South	0	60	60
Canada Central	0	60	60
Canada East	0	60	60

# Storage

HDInsight by Microsoft

Quick create Custom (size, settings, apps)

- 1 Basics  
Configure basic settings ✓
- 2 Storage  
Set storage settings >
- 3 Summary  
Confirm configurations >

**Storage Account Settings**

\* Primary storage type  
☒ Azure Storage ☐ Data Lake Store

\* Selection method ⓘ  
☒ My subscriptions ☐ Access key

\* Create a new Storage account  
zdjordjestorageaccount ✓

Select existing

\* Default container ⓘ  
zdjordjehadoop-2018-01-10t23-15-30-9...

Additional storage accounts  
Optional >

Data Lake Store access ⓘ  
Optional >

**Metastore Settings (optional)**  
Filtered to location and subscription of cluster.

To preserve your metadata outside this cluster, link a SQL database to this account.

Select a SQL database for Hive  
No database in eastus for subscription. ▾

Next

- For storage, Azure gives you 2 options:
  - Azure storage or
  - Data Lake Store
- It turns out that Data Lake Store setup requires administrative privileges that most of us do not have and we will stay away from the data lake for now.
- If you do not have a storage account, create one.
- Hit [Next](#)

# Cluster Summary, Applications, Cluster size

- On Cluster Summary page select Applications ([Edit](#)).
- This, just to see that you have options.
- Do not choose anything but rather select [Next](#). Cluster size plate appears.
- Notice that Azure offers cluster with 2 head (master) and 4 worker nodes. Price per hour on my screen was \$3.11. Change the number of workers to 1, price of cluster should come down to \$1.34/hour. Number of cores used is reduced to 23.
- Select [Next](#)
- On Advanced settings page click [Next](#)

Cluster size

To learn more, visit our pricing page.  
[Learn more](#)

Number of Worker nodes

1

★ Worker node size

D4 v2 (1 node, 8 cores)


>

★ Head node size

D12 v2 (2 nodes, 8 cores)

>

WORKER NODES	$0.590 \times 1 = 0.590$
HEAD NODES	$0.374 \times 2 = 0.748$
TOTAL COST	<b>1.34</b>
USD/HOUR (ESTIMATED)	



This cluster will use 16 cores out of 60 available cores in East US. Your cores quota in East US is 60 cores for this subscription.

[Click here](#) to view cores usage.

This price estimate does not include storage costs, network egress costs, or subscription discounts.

Questions? [Contact billing support](#).

To learn more, visit our [pricing page](#).

Next

# Cluster summary

- On Cluster summary plate, you can decide to change parameters or hit [Create](#)

HDInsight  
by Microsoft

Quick create Custom (size, settings, apps)

1 Basics  
Configure basic settings ✓

2 Storage  
Set storage settings ✓

3 Applications (optional)  
Productivity through applicatio... ✓

4 Cluster size  
Choose node sizes ✓

5 Advanced settings  
Configure advanced features ✓

6 Summary  
Confirm configurations >

This cluster may take up to 20 minutes to create.

Cluster summary

Click 'Create' below to deploy your cluster. Once created you will be billed until your cluster is deleted.

Basics (Edit)

Cluster name	zjdjrdjehadoop
Subscription	Pay-As-You-Go
Cluster type	Hadoop 2.7 on Linux (HDI 3.6)
Cluster login username	admin
SSH username	sshuser
Resource group	zjdjrdjehadoopgrp
Location	East US

Storage (Edit)

Azure Storage account	zjdjrdjestorageaccount (new)
Additional Storage accounts	---
Metastores	---

Applications (optional) (Edit)

Applications (optional)	---
-------------------------	-----

Cluster size (Edit)

Nodes	Head (2 x D12 v2), Worker (1 x D4 v2)
-------	---------------------------------------

Advanced settings (Edit)

Script actions	---
Virtual network	---

1.34 USD  
USD/HOUR (ESTIMATED)

3 NODES (2 HEAD + 1 WORKER) 16 CORES - --

HADOOP 2.7 ON LINUX (HDI 3.6) ZDJORDJESTORAGEACCOUNT (EAST US)

Create Download template and parameters

@Zoran B. Djordjević, Nishava, Inc.

8



# Resource Group content

- My resource group, zdjordjehadoopgrp, no contains a HDInsight cluster named zdjordjehadoop

zdzordjehadoopgrp  
Resource group

Search (Ctrl+/)

Overview  
Activity log  
Access control (IAM)  
Tags

SETTINGS  
Quickstart  
Resource costs

+ Add Columns Delete resource group Refresh Move Assign Tags

Subscription (change)  
Pay-As-You-Go  
Subscription ID  
ab293589-2b01-4b26-97a4-24f36b894fd2

Deployments  
1 Succeeded


Filter by name... All types All locations No grouping

2 items

NAME ↑↓	TYPE ↑↓	LOCATION ↑↓
zdjordjehadoop	HDInsight cluster	East US
zdjordjehadooppaccount	Storage account	East US

- Click on cluster name (zdzordjehadoop)

# HDInsight Cluster View

 **zdjordjehadoop**  
HDInsight cluster

Search (Ctrl+ /)

Overview

Activity log

Access control (IAM)

Tags

Diagnose and solve problems

SETTINGS

Locks

Automation script

GETTING STARTED

Quick start

Tools for HDInsight

CONFIGURATION

Subscription cores usage

Scale cluster

SSH + Cluster login

Cluster Dashboard

Secure Shell (SSH)

Scale cluster

Move

Delete

Essentials ^

Resource group [\(change\)](#)  
**zdjordjehadoopgrp**

Status  
Running

Location  
East US

Subscription name [\(change\)](#)  
**Pay-As-You-Go**

Subscription ID  
ab293589-2b01-4b26-97a4-24f36b894fd2

Learn more

[Documentation](#)


Cluster type, HDI version  
Hadoop on Linux (HDI 3.6)


URL  
<https://zdjordjehadoop.azurehdinsight.net>


Getting started  
[Quickstart](#)

Head Nodes, Worker nodes  
D12 v2 (x2), D4 v2 (x1)

Quick links


 Cluster dashboard

 Ambari Views


 Scale cluster


Usage

Cluster nodes

**3 nodes** 

TYPE	NODE SIZE	CORES	NODES
Head	D12 v2	8	2
Worker	D4 v2	8	1

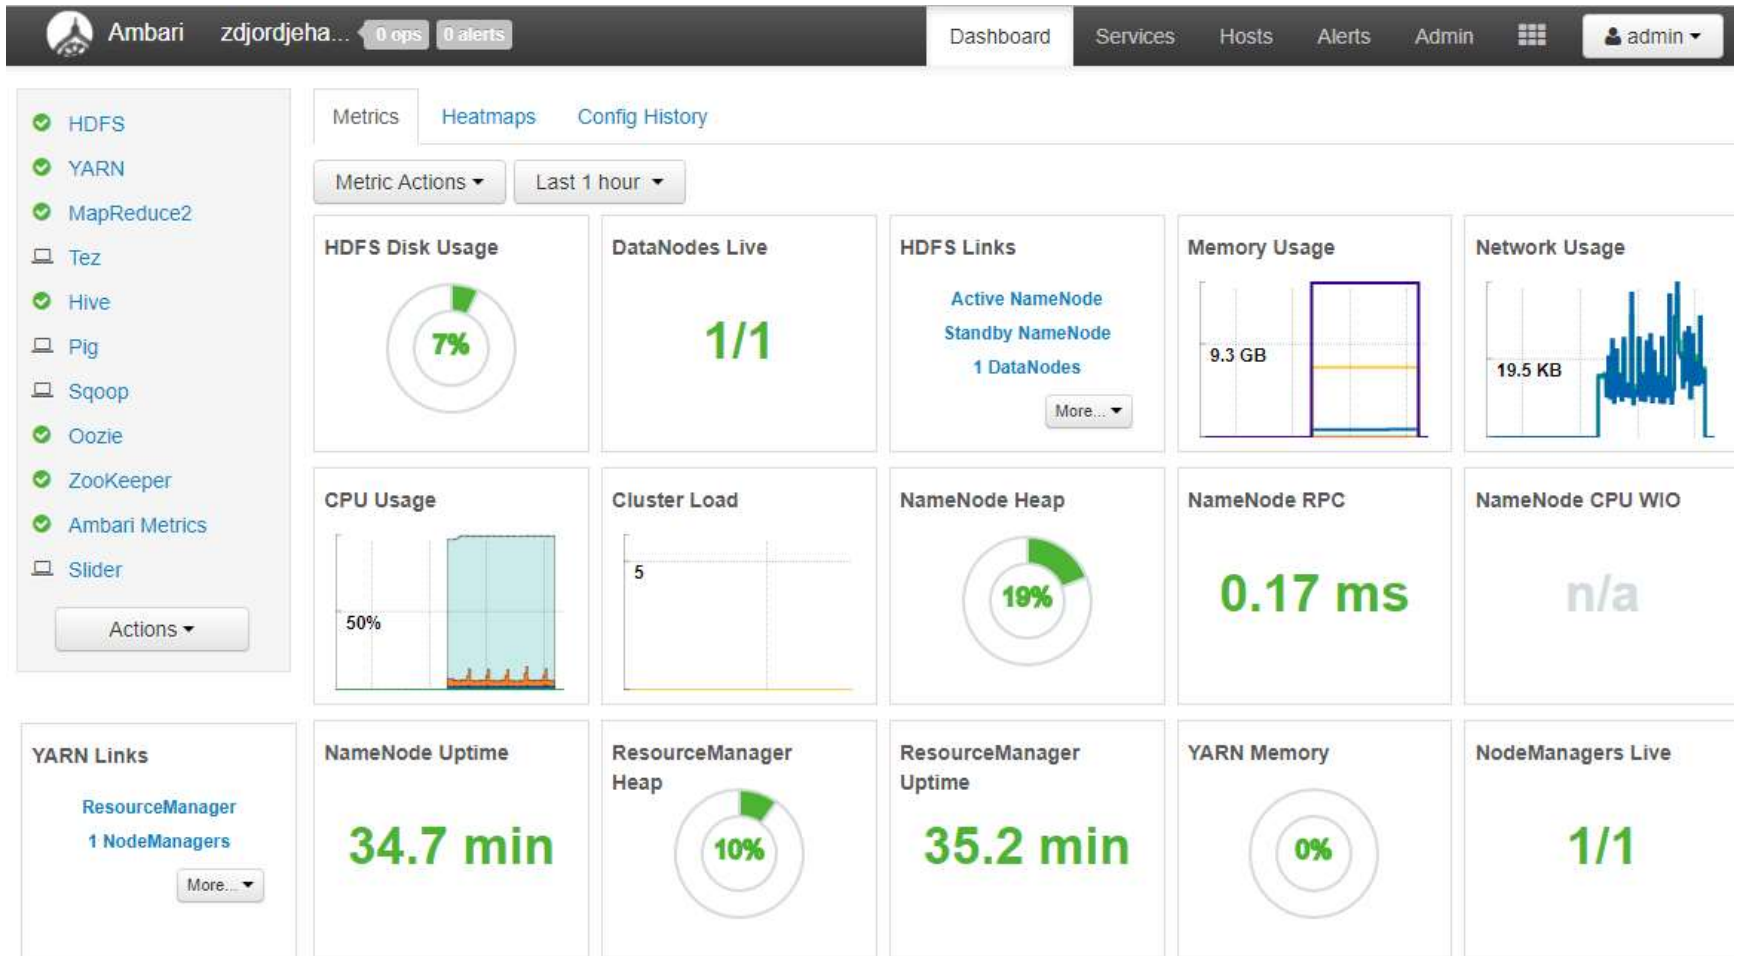
 Applications

 Script actions

@Zoran B. Djordjević, Nishava, Inc.


10


# Cluster Dashboard



# ssh Access


- HDInsight Dashboard gives you a convenient "Secure Shell (SSH)" link which tells you how to connect to the master node using ssh. We can do that using Cygwin.

 **SSH + Cluster login**  
zjdjrdjehadoop


 Cluster Dashboard


---

**Connect to cluster using secure shell (SSH)**


You can securely connect to the below endpoints in the HDInsight cluster with an SSH client.  
[Documentation](#) 


Hostname





**Connect to cluster using Cluster Login**

You can use the cluster dashboard to submit jobs, monitor resource usage and perform administrative actions. [Documentation](#) 

Cluster login username 

**Reset credential**

# Cygwin Login

```
$ ssh sshuser@zdjordjehadoop-ssh.azurehdinsight.net
The authenticity of host 'zdjordjehadoop-ssh.azurehdinsight.net (40.71.25.124)' can't
be established.
ECDSA key fingerprint is SHA256:PDav303yIBI+v6o/JVqDELV730NffnhtixwM/VA1Oz0.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'zdjordjehadoop-ssh.azurehdinsight.net,40.71.25.124' (ECDSA)
to the list of known hosts.
Authorized uses only. All activity may be monitored and reported.
sshuser@zdjordjehadoop-ssh.azurehdinsight.net's password:
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.4.0-101-generic x86_64)
 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

  Get cloud support with Ubuntu Advantage Cloud Guest:
    http://www.ubuntu.com/business/services/cloud

82 packages can be updated.
45 updates are security updates.

Welcome to HDInsight. Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted
by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

sshuser@hn0-zdjord:~$
```

# Hive View

- On HDInsight plate, hit [Cluster dashboard](#) link, then, again hit [HDInsight cluster dashboard](#). On the right navigation list select [Hive](#). If asked provide password for user admin. Under Summary, select [Hive View 2.0](#)

The screenshot shows the Ambari HDInsight cluster dashboard for a cluster named 'zdjorjdjehadoop'. The top navigation bar includes links for Dashboard, Services, Hosts, Alerts, and Admin. The 'Services' tab is active, and the 'Hive' service is selected in the left-hand navigation menu. The main content area displays the 'Summary' tab for the Hive service, which shows a list of components and their status. All components are 'Started' and have 'No alerts'. The components listed are: Hive Metastore (2 instances), HiveServer2 (2 instances), WebHCat Server (2 instances), HCat Client (1 installed), and Hive Clients (3 installed). At the bottom, there are links for 'Hive View 2.0' and 'Debug Hive Query', both with a 'Go To View' button. The URL in the browser address bar is 'https://zdjorjdjehadoop.azurehdinsight.net/#/main/services/HIVE/summary'.

Secure | <https://zdjorjdjehadoop.azurehdinsight.net/#/main/services/HIVE/summary>

Ambari | zdjorjdjehadoop | 0 ops | 0 alerts | Dashboard | Services | Hosts | Alerts | Admin | admin

Summary | Configs | Service Actions

**Summary** No alerts

- [Hive Metastore](#) Started No alerts
- [Hive Metastore](#) Started No alerts
- [HiveServer2](#) Started No alerts
- [HiveServer2](#) Started No alerts
- [WebHCat Server](#) Started No alerts
- [WebHCat Server](#) Started No alerts
- [HCat Client](#) 1 HCat Client Installed
- [Hive Clients](#) 3 Hive Clients Installed

HiveServer2 JDBC URL [jdbc:hive2://zk0-zdjord.13yy4wperviu5ab2cnpftsr2zf.bx.internal.cloudapp.net:2181,zk...](#)

Hive View 2.0 [Go To View](#)

Debug Hive Query [Go To View](#)

# Hive Query View

- In Hive Query view that appears enter:
  1. `show tables;` # note Hive queries are terminated with ";" (not Python)
- Hit Execute on the bottom. You will see a single table name: hivesampletable.
- Write another query:
  2. `describe hivesampletable;`
  3. `select * from hivesampletable;`

You will get table description and content:

Filter columns 

col_name	data_type	comment
clientid	string	""
querytime	string	""
market	string	""
deviceplatform	string	""
devicemake	string	""
devicemodel	string	""

hivesampletable.clientid	hivesampletable.querytime	hivesampletable.market
8	18:54:20	en-US
23	19:19:44	en-US
23	19:19:46	en-US
23	19:19:47	en-US
28	01:37:50	en-US

- To stop consuming resources, wipe out your resource group and included cluster.


# Hortonworks VMs

- Microsoft's Azure VMs are based on Hortonworks HDP framework. When developing software, you are better off working on Hortonworks VMs on your laptop, PC or Mac, where they do not cost you a penny.
- One problem with those VMs is that they require at least 8GB of RAM to run.
- Hortonworks has two family of products HDP (Hortonworks Data Platform) and HDF (Hortonworks Data Flow) and corresponding VMs. We will use HDP VMs and in particular so called sandbox VMs.
- The HDP Sandbox, a single node VM, lets you get started with Apache Hadoop, Falcon, Atlas, Tez, Sqoop, Flume, Kafka, Pig, Hive, HBase, Accumulo, Storm, Solr, Spark, Ranger, Knox, Ambari, ZooKeeper, Oozie, Phoenix, NiFi, HAWQ, Zeppelin, Atlas, Slider, Mahout, MapReduce, HDFS, YARN, Metron Apache projects
- The HDF Sandbox lets you get started with Apache NiFi, Apache Kafka, Apache Storm, Druid and Streaming Analytics Manager.
- Go to <https://hortonworks.com/products/sandbox> and download HDP Sandbox

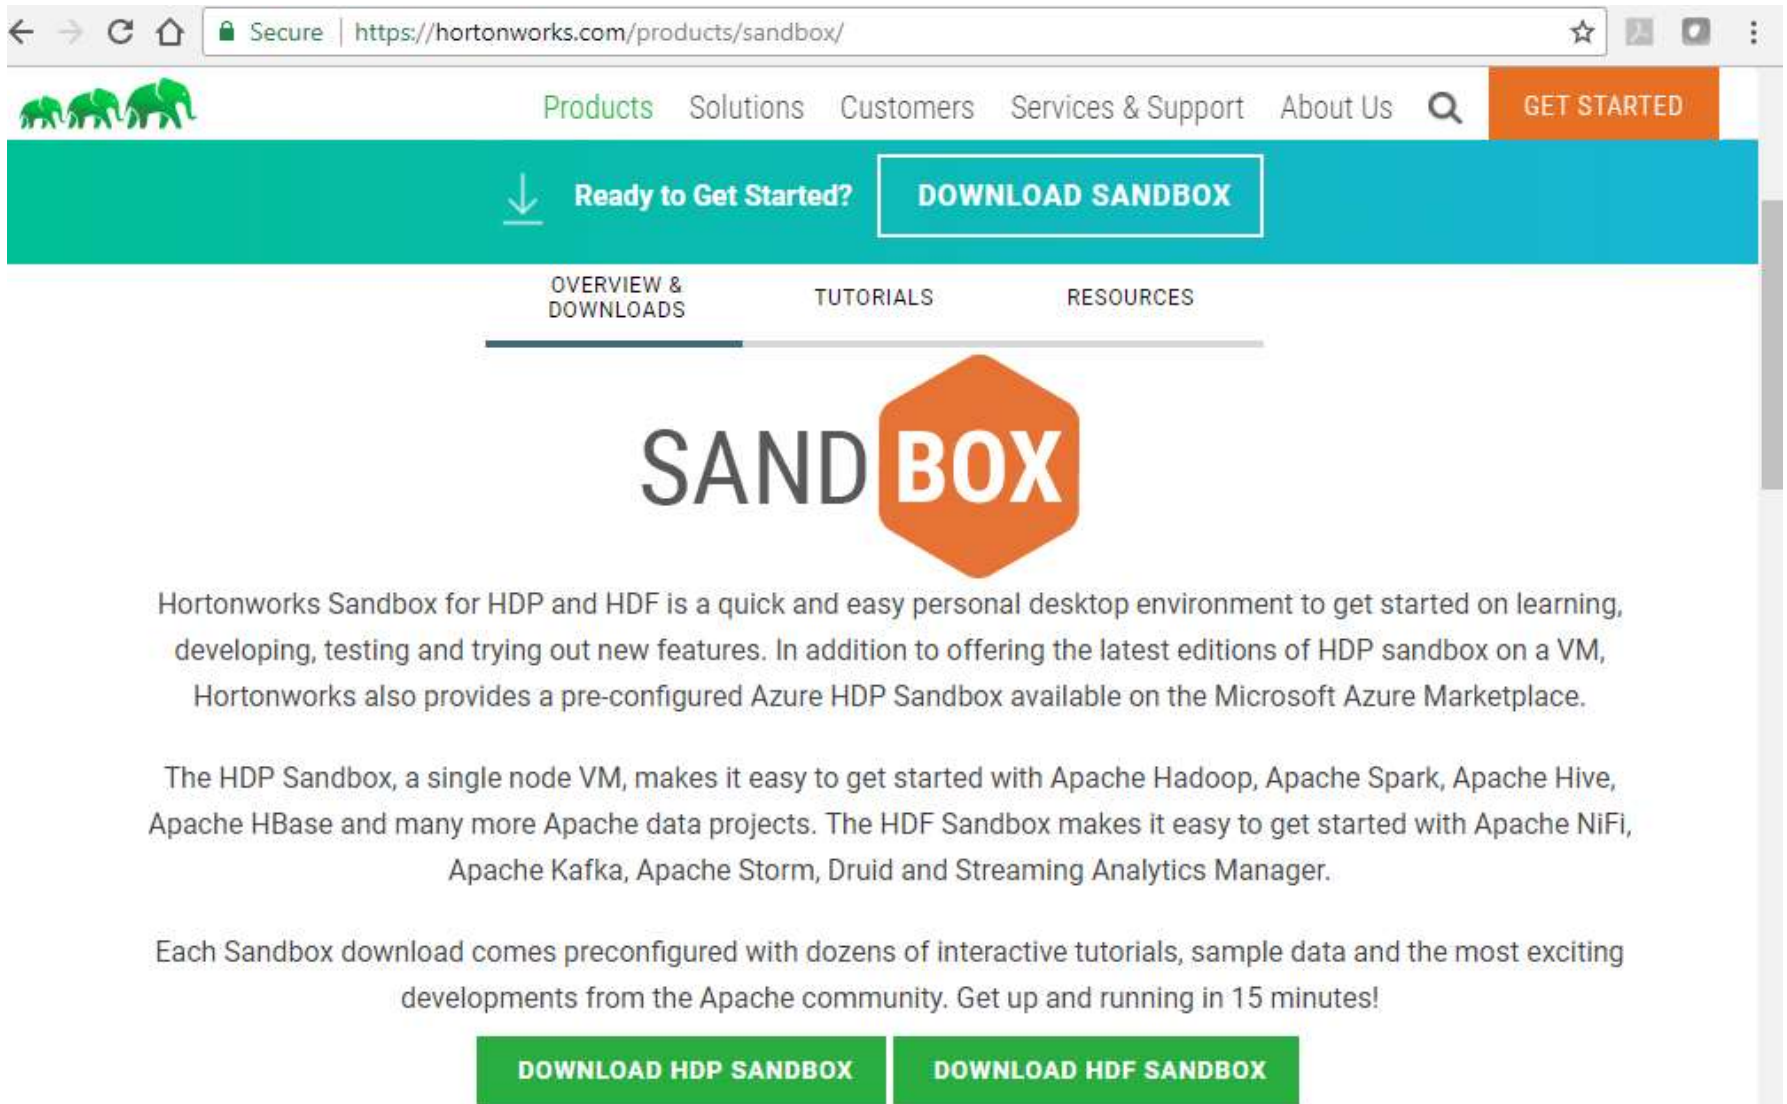


# Content of HDP

- On URL <https://hortonworks.com/product/data-platforms/hdp> you can see detailed content of HDP framework in terms of various Apache and other open source packages.



# Download Hortonworks HDP Sandbox



The screenshot shows the Hortonworks website's 'Sandbox' page. The browser address bar displays 'https://hortonworks.com/products/sandbox/'. The navigation bar includes links for 'Products', 'Solutions', 'Customers', 'Services & Support', and 'About Us', along with a search icon and a 'GET STARTED' button. A teal banner features a download icon, the text 'Ready to Get Started?', and a prominent 'DOWNLOAD SANDBOX' button. Below this, a horizontal menu highlights 'OVERVIEW & DOWNLOADS' with a blue underline, while 'TUTORIALS' and 'RESOURCES' are also visible. The main heading 'SAND BOX' is displayed, with 'SAND' in black and 'BOX' in white inside an orange hexagon. The text describes the sandbox as a quick and easy personal desktop environment for learning and testing HDP and HDF features. It mentions that Hortonworks provides pre-configured Azure HDP Sandboxes on the Microsoft Azure Marketplace. The HDP Sandbox is described as a single node VM for getting started with various Apache projects like Hadoop, Spark, Hive, HBase, Kafka, Storm, and Druid. The HDF Sandbox is similarly described for getting started with NiFi, Kafka, Storm, and Streaming Analytics Manager. It concludes by stating that each download comes preconfigured with tutorials and sample data, allowing users to get up and running in 15 minutes. At the bottom, there are two green buttons: 'DOWNLOAD HDP SANDBOX' and 'DOWNLOAD HDF SANDBOX'.

Secure | https://hortonworks.com/products/sandbox/

Products Solutions Customers Services & Support About Us GET STARTED

Ready to Get Started? DOWNLOAD SANDBOX

OVERVIEW & DOWNLOADS TUTORIALS RESOURCES

## SAND BOX

Hortonworks Sandbox for HDP and HDF is a quick and easy personal desktop environment to get started on learning, developing, testing and trying out new features. In addition to offering the latest editions of HDP sandbox on a VM, Hortonworks also provides a pre-configured Azure HDP Sandbox available on the Microsoft Azure Marketplace.


The HDP Sandbox, a single node VM, makes it easy to get started with Apache Hadoop, Apache Spark, Apache Hive, Apache HBase and many more Apache data projects. The HDF Sandbox makes it easy to get started with Apache NiFi, Apache Kafka, Apache Storm, Druid and Streaming Analytics Manager.

Each Sandbox download comes preconfigured with dozens of interactive tutorials, sample data and the most exciting developments from the Apache community. Get up and running in 15 minutes!

DOWNLOAD HDP SANDBOX DOWNLOAD HDF SANDBOX

# Handle Sandbox VM

- You can select a VMware VM, Virtual Box VM or a Docker container.
- I downloaded `HDP_2.6.3_vmware_16_11_2017.ova` file with a volume of 10GB.
- An OVA file is a virtual appliance used by virtualization applications such as VMware Workstation and Oracle VM Virtualbox. It is a package that contains files used to describe a virtual machine, which includes an .OVF descriptor file, optional manifest (.MF) and certificate files, and other related files.
- `HDP_263_vmware...ova` file can be opened by VMWare Workstation, which will transform it into a regular VMWare VM.
- When you open VMWare Workstation, select Open a Virtual Machine and then select downloaded \*.ova file. Select the storage path for new virtual machine.
- Select **Import**



# The first screen and IP Address

- Once the import is finished, run the VM. It will display a confirmation screen.

```
HDP 2.6.3
http://hortonworks.com

To initiate your Hortonworks Sandbox session,
please open a browser and enter this address
in the browser's address field:
http://192.168.135.131:8888/

Log in to this virtual machine: Linux/Windows <Alt+F5>, Mac OS X <Ctrl+Alt+F5>
```

- Note the IP address of your machine. My IP address is 192.168.135.131
- On your Windows machine **as an administrator open Notepad** and then edit file:  
C:\Windows\system32\Drivers\etc\hosts
- And add a line associating the above IP address with your favorite machine names.
- For example:  
192.168.135.131 localhost sandbox.hortonworks.com sandbox-hdp.hortonworks.
- On Mac you would add that line to /private/etc/host and on Linux to /etc/hosts file.
- If you follow messages displayed during the boot-up, you would notice that this VM runs a CentOS 7 operating system.

# Run ssh client

- To login into your VM, on Linux/Windows do `Alt+F5`. On Mac OS, do `Ctrl+Alt+F5`
- The default password for user `root` is `hadoop` (all lowercase)
- You can connect to your machine using Cygwin or any other ssh client. Note that this VM listens to ssh traffic on port 2222. Type:

```
$ ssh root@192.168.135.131 -p 2222
The authenticity of host '[192.168.135.131]:2222
([192.168.135.131]:2222)' can't be established.
RSA key fingerprint is
SHA256:0igBYD9FHqqbVcTryp+T3RxYa2l0JvTTbPNihSLUXyg.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '[192.168.135.131]:2222' (RSA) to the
list of known hosts.
root@192.168.135.131's password:hadoop
You are required to change your password immediately (root
enforced)
Last login: Thu Jan 11 21:08:19 2018 from 192.168.135.1
Changing password for root.
(current) UNIX password:hadoop
New password:*****
Retype new password:*****
[root@sandbox-hdp ~]#
```

# Copy files to and from VM using scp

- Using Cygwin or a terminal of your choice, you can transfer files to/from sandbox and local machine.
- Transfer file from local machine to sandbox directory /tmp by issuing the following command:

```
scp -P 2222 local_dir_or_file root@sandbox-hdp.hortonworks.com:/tmp
```

- Transfer file from sandbox to local machine:

```
scp -P 2222 root@sandbox-hdp.hortonworks.com:/tmp/sandbox_dir_or_file  
local_dir_or_file
```

- Notice the difference between the two commands?
- To send data from local machine to sandbox, the local machine directory path comes before sandbox directory. To transfer data from sandbox to local machine, the command arguments are reversed.

# Default Users

- VM comes with a set of default users: admin, maria\_dev, raj\_ops, holger\_gov, amy\_ds. Those users have password the same as the username.
- User admin requires change of password. To accomplish the change, run the following command on the Linux prompt:

```
[root@sandbox-hdp ~]# ambari-admin-password-reset
Please set the password for admin:
Please set the password for admin:
Please retype the password for admin:
```

The admin password has been set.

Restarting ambari-server to make the password change effective...

Using python /usr/bin/python

Restarting ambari-server

Waiting for server stop...

Ambari Server stopped

Ambari Server running with administrator privileges.

Organizing resource files at /var/lib/ambari-server/resources...

Ambari database consistency check started...

Server PID at: /var/run/ambari-server/ambari-server.pid

Server out at: /var/log/ambari-server/ambari-server.out

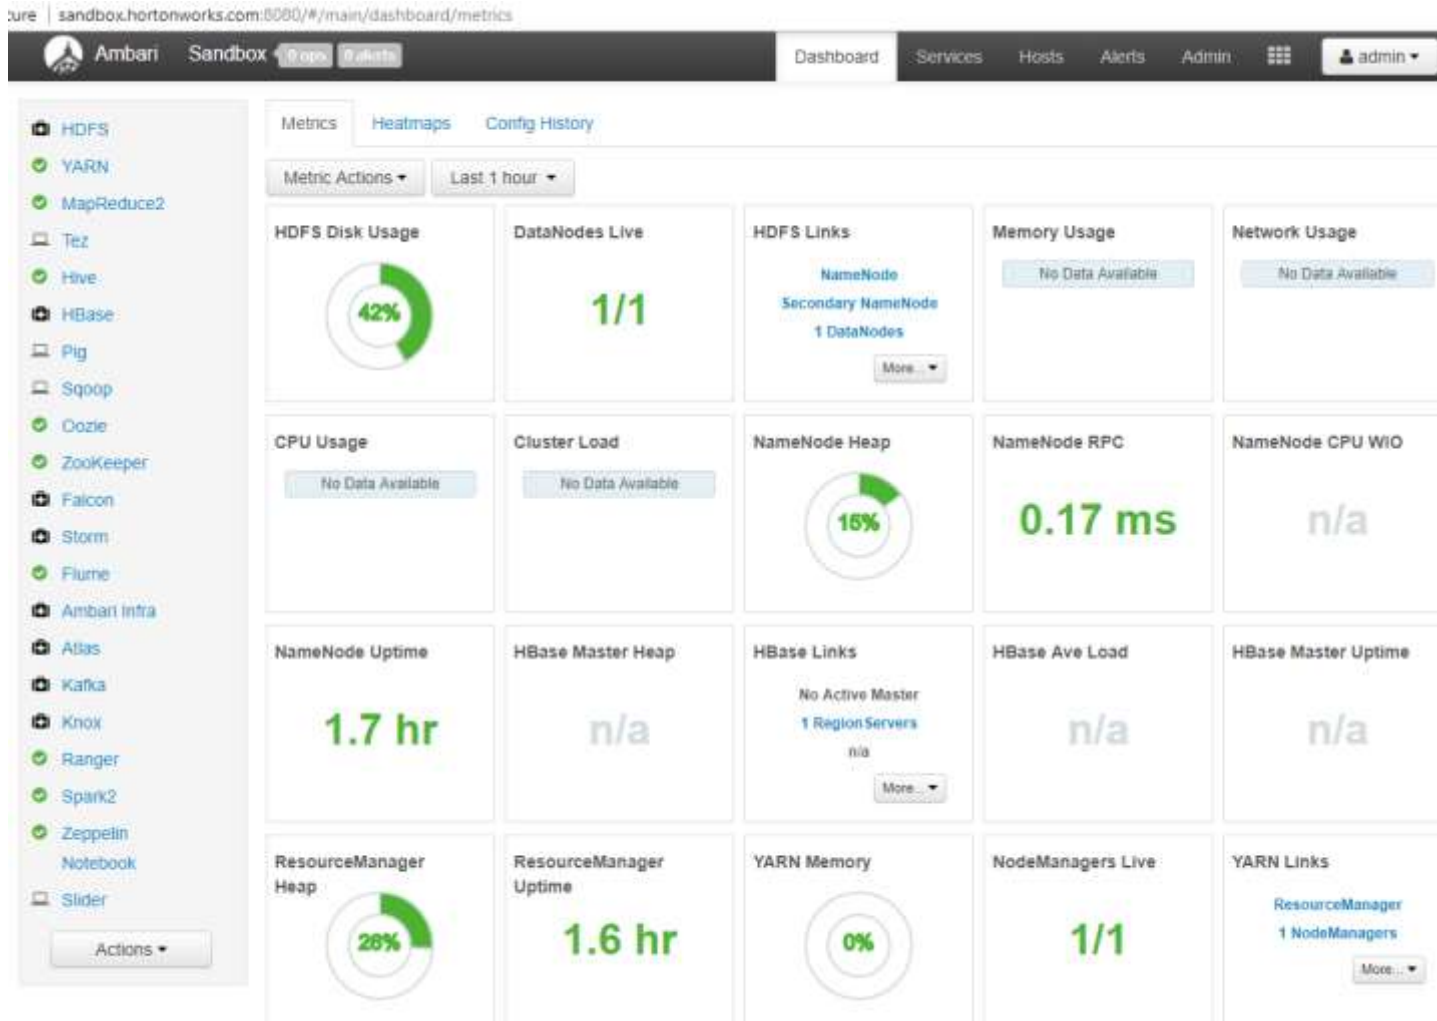
Server log at: /var/log/ambari-server/ambari-server.log

Waiting for server start.....



# Dashboard

- Go to <http://sandbox.hortonworks.com:8888>
- Then select **LAUNCH DASHBOARD**. Enter username admin and the password you have just assigned to that user. Singin! A dashboard appears:





# Please remember port 2222

- Hortonworks Sandbox runs on docker container.
- So some profile configs are loaded only when you do "ssh on port 2222"
- This is the way port forwarding works is that it exposes port 22 of the sandbox machine to the outside network of your PC via port 2222.
- If you try ssh-ing into your VM through port 22, you might get in but would see a different environment
- Once you create user `centos`, for example, you will `ssh` into the sandbox VM by typing:

```
$ ssh centos@192.168.1356.132 -p 2222
```

# Introduction to Hadoop Distribute File System (HDFS)

# Examine Directories in your HDFS

- User `hdfs` has complete insight into all directories in your HDFS.
- User `hdfs` can create user (home) directories and perform maintenance operation on HDFS.
- To see home directories in HDFS, type:

```
$ sudo -u hdfs hdfs dfs -ls /user/
```

```
drwxr-xr-x  - root root          0 2017-07-19 06:28 /user/root
drwxr-xr-x  - mapred  hadoop           0 2017-07-19 06:29 /user/history
drwxrwxrwx  - hive    supergroup      0 2017-07-19 06:31 /user/hive
drwxrwxrwx  - hue     supergroup      0 2017-07-19 06:30 /user/hue
drwxrwxrwx  - jenkins supergroup      0 2017-07-19 06:29 /user/jenkins
drwxrwxrwx  - oozie   supergroup      0 2017-07-19 06:30 /user/oozie
drwxrwxrwx  - root    supergroup      0 2017-07-19 06:29 /user/root
drwxr-xr-x  - hdfs    supergroup      0 2017-07-19 06:31 /user/spark
```

- You can see inside all of those directories, just type

```
$ sudo -u hdfs hdfs dfs -ls -R /user/hive
```

```
drwxrwxrwx  - hive supergroup      0 2017-07-19 06:31 /user/hive/warehouse
```

# Create Linux user `centos`

- If you need to create a new Linux user, e.g. `centos`, log in as `root`, or type:

```
$ su -
```

and enter `root`'s password.

- Now, as user `root`, type:

```
$ useradd -g mapred centos
```

- The above creates new user `centos`, as a member of group `mapred`.
- Please note that a user running Hadoop MapReduce programs must be a member of `mapred` group. To create password for new user, type:

```
$ passwd centos
```

- At the `New password: prompt`, enter a password for user `centos`, press [Enter].
- At the `Retype new password: prompt`, enter the same password to confirm.
- If user `root` wants to become user `centos`, `root` types:

```
[root@sandbox-hdp ~]$ su -- centos
```

```
Password: xxxxxxxxxxxx
```

```
[centos@sandbox-hdp root]$ exit # turns you back into root
```

- Should you want to remove Linux user `centos`, type:

```
$ sudo userdel centos
```

# Make centos a sudo user

- Allow users in group wheel to run `sudo` command without password.
- As user root, type

```
$ chmod +rw /etc/sudoers
```

```
$ visudo -f /etc/sudoers
```

- Uncomment the last (second) line with `%wheel` :

```
## Allows people in group wheel to run all commands
```

```
# %wheel          ALL=(ALL)          ALL
```

```
## Same thing without a password
```

```
%wheel  ALL=(ALL)          NOPASSWD: ALL
```

- Add user centos to group wheel:

```
[root]$ usermod -aG wheel centos
```

# Create HDFS Directories for user `centos`

- Next, create HDFS home directory for a new MapReduce user `centos`.
- By the way, on a truly distributed cluster you will do all of this on the Name Node.
- Type:

```
$ sudo -u hdfs hdfs dfs -mkdir /user/centos
```

- `hdfs dfs -mkdir` command automatically creates parent directories if they don't already exist. This is similar to the Unix `mkdir` command with the `-p` option.
- Once we have the directory we want to grant the ownership to user `centos`.

```
$ sudo -u hdfs hdfs dfs -chown centos:mapred /user/centos
```

- Finally we want to give full read-write-execute right on that directory.

```
$ sudo -u hdfs hdfs dfs -chmod 1777 /user/centos
```

- Alternatively, if the Linux user already exist, you can login as that user and create the home directory as follows:

```
$ sudo -u hdfs hdfs dfs -mkdir /user/$USER
```

```
$ sudo -u hdfs hdfs dfs -chown $USER /user/$USER
```

- If your system does not have HDFS directories for user `root`, you could use the above procedure to create them

# hadoop script

- You used user `hdfs` and command `hdfs` only for high level system maintenance of HDFS.
- To manipulate files in your (HDFS) home directory, transfer files from local file system to HDFS and back, ordinary users use command `hadoop`.
- Let us find out what `hadoop` is. If you type:

```
$ which hadoop      # you will get:
```

```
/usr/bin/hadoop
```

- If you open `/usr/bin/hadoop` file, you will see that it invokes another file `/usr/hdp/2.6.3.0-235//hadoop/bin/hadoop.distro`
- That other `hadoop` file is also a script which runs various Java programs depending on the options you pass to it. In the original Hadoop script we see that Hortonworks VM uses the following environmental variables:

```
HADOOP_HOME=/usr/hdp/2.6.3.0-235/hadoop
```

```
HADOOP_MAPRED_HOME=/usr/hdp/2.6.3.0-235/hadoop-mapreduce
```

```
HADOOP_YARN_HOME=/usr/hdp/2.6.3.0-235/hadoop-yarn}
```

```
HADOOP_LIBEXEC_DIR=${HADOOP_HOME}/libexec
```

- To see the options of `hadoop` command, invoke `hadoop` by itself:

```
$ hadoop
```

# Options of `hadoop` script

```
[root]# hadoop
```

```
Usage: hadoop [--config confdir] COMMAND
```

```
where COMMAND is one of:
```

<code>fs</code>	run a generic filesystem user client
<code>version</code>	print the version
<code>jar &lt;jar&gt;</code>	run a jar file
<code>checknative [-a -h]</code>	check native hadoop and compression libraries availability
<code>distcp &lt;srcurl&gt; &lt;desturl&gt;</code>	copy file or directories recursively
<code>archive -archiveName NAME -p &lt;parent path&gt; &lt;src&gt;* &lt;dest&gt;</code>	create a hadoop archive
<code>classpath</code>	prints the class path needed to get the
<code>credential</code>	interact with credential providers
	Hadoop jar and the required libraries
<code>daemonlog</code>	get/set the log level for each daemon
<code>trace</code>	view and modify Hadoop tracing settings or
<code>CLASSNAME</code>	run the class named CLASSNAME

Most commands print help when invoked w/o parameters.



# Hadoop Distributed File System Shell, `fs` command

- Hadoop has access to both local, Linux, file system, and its own distributed file system (HDFS, Hadoop Distributed File System)
- We access HDFS through Hadoop distributed file system shell, `fs`;
- By invoking command itself:

```
$ hadoop fs
```

- we get a long list of options. Some of those resemble Unix (Linux) commands. Some are different.
- We use those commands to create directories in the HDFS, copy files between HDFS and the local file system, Internet and AWS S3 buckets.
- When you use `fs`, you always prefix it with `hadoop`.

# Options of file system shell, `hadoop fs`

```
[root]$ hadoop fs
```

```
Usage: hadoop fs [generic options]
```

```
[-appendToFile <localsrc> ... <dst>]  
[-cat [-ignoreCrc] <src> ...]  
[-checksum <src> ...]  
[-chgrp [-R] GROUP PATH...]  
[-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]  
[-chown [-R] [OWNER][:[GROUP]] PATH...]  
[-copyFromLocal [-f] [-p] [-l] <localsrc> ... <dst>]  
[-copyToLocal [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]  
[-count [-q] [-h] [-v] [-x] <path> ...]  
[-cp [-f] [-p | -p[topax]] <src> ... <dst>]  
[-createSnapshot <snapshotDir> [<snapshotName>]]  
[-deleteSnapshot <snapshotDir> <snapshotName>]  
[-df [-h] [<path> ...]]  
[-du [-s] [-h] [-x] <path> ...]  
[-expunge]  
[-find <path> ... <expression> ...]  
[-get [-p] [-ignoreCrc] [-crc] <src> ... <localdst>]  
[-getfacl [-R] <path>]  
[-getfattr [-R] {-n name | -d} [-e en] <path>]  
[-getmerge [-nl] <src> <localdst>]
```

# File system shell fs

```
[-help [cmd ...]]
  [-ls [-C] [-d] [-h] [-q] [-R] [-t] [-S] [-r] [-u] [<path> ...]]
  [-mkdir [-p] <path> ...]
  [-moveFromLocal <localsrc> ... <dst>]
  [-moveToLocal [-crc] <src> <localdst>]
  [-mkdir <path>]
  [-setrep [-R] [-w] <rep> <path/file>]
  [-touchz <path>]
  [-test -[ezd] <path>]
  [-stat [format] <path>]
  [-tail [-f] <file>]
  [-chmod [-R] <MODE[,MODE]... | OCTALMODE> PATH...]
  [-chown [-R] [OWNER][:[GROUP]] PATH...]
  [-chgrp [-R] GROUP PATH...]
  [-help [cmd]]
```

Generic options supported are

-conf <configuration file>	specify an application configuration file
-D <property=value>	use value for given property
-fs <local namenode:port>	specify a namenode
-jt <local jobtracker:port>	specify a job tracker
-files <comma separated list of files>	specify comma separated files to be copied to the map reduce cluster
-libjars <comma separated list of jars>	specify comma separated jar files to include in the classpath.

# Working with HDFS and Files

- If you run `hadoop fs -ls` command as an ordinary user, e.g. `centos` or `root`, and do not specify the root of HDFS directory tree, `/user/`, you only see the directories and files that belong to you and reside above your home directory.

- Let's check

```
[root]$ hadoop fs -ls
```

- There might be nothing there if user `root` has no subdirectories.
- If you want to see all the subdirectories, in a way similar to Unix's `ls` with the `-R` option, you can use `hadoop fs -ls -R` command.

```
[root]$ hadoop fs -ls -R    You'll see all the files and directories recursively.
```

```
[root]$ hadoop fs -ls /    # / is the root of the HDFS directory tree
```

- Command finds 12 items, some of which are:

```
drwxr-xr-x    - mapred hdfs      0 2017-11-10 14:38 /mapred
drwxrwxrwx    - mapred hadoop    0 2017-11-10 14:38 /mr-history
drwxr-xr-x    - hdfs  hdfs      0 2017-11-10 14:37 /ranger
drwxrwxrwx    - spark hadoop    0 2018-01-12 23:36 /spark2-history
drwxrwxrwx    - hdfs  hdfs      0 2017-11-10 15:00 /tmp
drwxr-xr-x    - hdfs  hdfs      0 2018-01-12 23:18 /user
drwxr-xr-x    - hdfs  hdfs      0 2017-11-10 14:38 /webhdfs
```

# Copying a file to new HDFS directory

- We are ready to add files to HDFS.
- We could fetch the .txt version of James Joyce's Ulysses by issuing the following command on the command prompt:

```
$ wget http://www.gutenberg.org/files/4300/4300-0.txt
```

- **We can place file 4300-0.txt into ulysses directory of user root**

```
$ hadoop fs -mkdir ulysses
```

```
$ hadoop fs -put 4300-0.txt ulysses
```

```
$ hadoop fs -ls ulysses
```

```
[root]$ hadoop fs -ls ulysses
```

```
Found 1 items
```

```
-rw-r--r-- 1 root root 1580890 2018-01-12 19:54 ulysses/4300-0.txt
```

- The number 1 in the above listing tells us how many times is a particular file replicated. Since we have a single machine, 1 is appropriate.
- The replication factor is 3 by default, but could be set to any number.

# Fetching and examining Files from HDFS

- The Hadoop command `get` does the reverse of `put`. It copies files from HDFS to the local file system.

- Let us first move file `/home/centos/4300-0.txt` to `/tmp`

```
$ mv 4300-0.txt /tmp
```

```
$ ls
```

- Nothing there.
- To retrieve file `4300-0.txt` from HDFS and copy it into the current local working directory, we run the command

```
hadoop fs -get ulysses/4300-0.txt .
```

- File `4300-0.txt` reappeared in `/home/centos` directory.
- A way to examine files in HDFS is to display data. For small HDFS files, Hadoop `cat` command is convenient.

```
hadoop fs -cat ulysses/4300-0.txt | wc -l  
32710
```

- We can follow any Hadoop file command by Linux pipe to forward its output for further processing by a Linux commands. For example, if the file is huge (as typical Hadoop files are) and you're interested in a quick check of its content, you can pipe the output of Hadoop's `cat` into a Unix `head`.

```
hadoop fs -cat ulysses/4300-0.txt | head
```

- Hadoop natively supports `tail` command for looking at the last kilobyte of a file.

```
hadoop fs -tail ulysses/4300-0.txt
```

# Deleting Files and Directories

- Hadoop command for removing files is `rm`.

```
hadoop fs -rm ulysses/4300-0.txt
```

- To delete files and directories recursively use

```
hadoop fs -rm -R directory/*
```

- To delete empty directories use

```
hadoop fs -rmdir directory
```

# Running a MapReduce Example on YARN

- As user `centos`, we made a directory in HDFS called `input` and copied file `4300-0.txt` into that directory by running the following commands:

```
hadoop fs -mkdir input
```

```
hadoop fs -put 4300-0.txt input
```

```
hadoop fs -ls input
```

```
Found 1 items
```

```
-rw-r--r--    1 centos mapred    1580890 2018-01-12 23:57 input/4300-0.txt
```

- To run MapReduce examples we need to set `HADOOP_MAPRED_HOME` for user `centos`. Best, enter that environmental variable into `.bash_profile` file and then source that file. You can do it for a single use only, as well.

```
$ export HADOOP_MAPRED_HOME=/usr/hdp/2.6.3.0-235/hadoop-mapreduce
```

- Hadoop's MapReduce jobs always read and write from and to HDFS files and/or directories. The above is a standard procedure with regular Hadoop's MapReduce jobs.



# Running a MapReduce grep Job

- On our VM example MapReduce scripts (jobs) are contained in `$HADOOP_MAPRED_HOME/hadoop-mapreduce-examples.jar` file.
- One of the scripts is referred to as `grep` job which counts how many times every word appears in the analyzed corpus.
- In our case, Hadoop `grep` would scan the file (with James Joyce novel) placed in the specified (HDFS) directory `"input"` and create a tab delimited report named `ulysses_freq`.
- Hadoop `grep` uses simple regular expression `'\w+'` to select all multi-character words. In this regular expression `'\w'`, with a lowercase `w`, represents word characters.
- This `grep` is different from Unix (Linux) `grep`. Unix `grep` returns lines where a pattern appears. Hadoop example `grep` counts word frequencies.
- The command to run Hadoop `grep` reads:  

```
$ hadoop jar /usr/hdp/2.6.3.0-235/hadoop-mapreduce/hadoop-mapreduce-examples.jar grep input/4300-0.txt ulysses_freq '\w+'
```
- Job takes a few minutes. You could monitor progress of all map jobs and
- reduce jobs. The output is placed in HDFS directory `ulysses_freq`

# WordCount in Hadoop MapReduce API

```
package org.apache.hadoop.examples; import java.io.IOException;
import java.util.StringTokenizer; import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path; import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text; import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;
public class WordCount {
    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken()); context.write(word, one);
            }
        }
    }
}
```

# WordCount in Hadoop MapReduce API

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {    sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    String[] otherArgs = new
        GenericOptionsParser(conf,args).getRemainingArgs();
    if (otherArgs.length < 2) {
        System.err.println("Usage: wordcount <in> [<in>...] <out>");
        System.exit(2);
    }
}
```

# WordCount in Hadoop MapReduce API

```
Job job = Job.getInstance(conf, "wordcount");
job.setJarByClass(WordCount.class);
job.setMapperClass(TokenizerMapper.class);
job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
for (int i = 0; i < otherArgs.length - 1; ++i) {
    FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
}
FileOutputFormat.setOutputPath(job,
    new Path(otherArgs[otherArgs.length - 1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

# Running a MapReduce example

- After a lot of console output, the job completes.
- We can find the output in the HDFS directory named `ulysses_freq`, because we specified that particular output directory to Hadoop.

- **Type:**

```
$ hadoop fs -ls
```

- You will see directories `input` and `ulysses_freq`. We list the content of `ulysses_freq`

```
$ hadoop fs -ls ulysses_freq
```

```
Found 2 items
```

```
-rw-r--r--    1 centos mapred          0 2018-01-13 00:17 ulysses_freq/_SUCCESS
-rw-r--r--    1 centos mapred 351794 2018-01-13 00:17 ulysses_freq/part-r-00000
```

- The content of the output file `part-r-hadoop00000` can be seen using `fs -cat` command:

```
[centos]$ hadoop fs -cat ulysses_freq/part-r-00000 | head
```

```
13686    the
8169     of
6697     and
5892     a
4872     to
4743     in
3063     his
2982     I
2970     he
2701     s
```

# Introduction to Spark

# Select Version of Spark

- Hortonworks VM comes with two versions of Spark, Spark 1.6.3, referred to as `Spark1` and Spark 2.2.0.2 referred to as `Spark2`.
- To choose one or the other version you set the environmental variable `SPARK_MAJOR_VERSION` to value 1 or value 2, like:

```
$ export SPARK_MAJOR_VERSION=1 or
```

```
$ export SPARK_MAJOR_VERSION=2
```

- The following slide presents the output of `pyspark`, the command shell used for Python Spark interactive programming, when the version is set to `Spark1`.

pyspark

- If you type `pyspark` on the Linux command prompt, in `centos` home directory, you will see the following:

```
[centos@sandbox-hdp ~]$ pyspark
SPARK_MAJOR_VERSION is set to 1, using Spark
Python 2.6.6 (r266:84292, Aug 18 2016, 15:13:37)
[GCC 4.4.7 20120313 (Red Hat 4.4.7-17)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
18/01/13 00:31:18 INFO spark.SparkContext: Running Spark version 1.6.3
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).
18/01/13 00:42:04 INFO netty.NettyBlockTransferService: Server created on 44301
18/01/13 00:42:04 INFO storage.BlockManagerMaster: Trying to register BlockManager
18/01/13 00:42:04 INFO storage.BlockManagerMasterEndpoint: Registering block manager
localhost:44301 with 511.1 MB RAM, BlockManagerId(driver, localhost, 44301)
18/01/13 00:42:04 INFO storage.BlockManagerMaster: Registered BlockManager
Welcome to
```

version 1.6.3

```
Using Python version 2.6.6 (r266:84292, Aug 18 2016 15:13:37)
SparkContext available as sc, HiveContext available as sqlContext. >>>
>>> quit()
$
```



# Load Data (RDD), Spark 1.6

- In Spark, we express our computation through operations on distributed collections that are automatically parallelized across the cluster.
- These collections are called *resilient distributed datasets*, or RDDs.
- When we load some data, i.e. a file into a shell variable, we are creating an RDD, like

```
>>> lines = sc.textFile("file:///home/centos/4300-0.txt")
```

```
>>> lines.count()
```

```
32742
```

- `4300-0.txt` is a text file residing in the home directory of user `centos`.
- What we've just done is create and populate an RDD named `lines` using a mysterious object `"sc"` and its method `textFile()`. We populated RDD `lines` with data in file `4300-0.txt`.
- `"sc"` stands for an implicit `SparkContext`.
- `SparkContext` communicates with the execution environment.
- It appears that RDD-s are also (Object Oriented) objects and have methods, such as `count()` which gave use the exact number of lines in file `4300-0.txt`.
- We could ran the above load process using command:

```
lines = sc.textFile("file:///home/centos/4300-0.txt")
```

- Emphasizing that the file resides in a regular file system

# Spark could read HDFS files

- We could have ran previous commands like:

```
lines2 = sc.textFile("/user/centos/ulysses/4300-0.txt")  
lines2.count()  
32710
```

- Spark, as this command shows, by default, reads from HDFS files. Owe could have emphasized use of HDFS by formulating this command as:

```
lines3 = sc.textFile("hdfs:///user/centos/ulysses/4300-0.txt")  
lines3.count()  
32710
```

# Shell verbosity and `log4j.properties` file

- The output we have seen is long and annoying. We could shorten it by creating `log4j.properties` files in the directories:  
`/usr/hdp/2.6.3.0-235/etc/spark2/conf/log4j.properties.template`  
`/usr/hdp/2.6.3.0-235/etc/spark/conf/log4j.properties.template`
- You create those files by copying provided file `log4j.properties.template` and by changing line  
`log4j.rootCategory=INFO, console`
- to read:  
`log4j.rootCategory=ERROR, console`
- That lowers (raises?) the logging level so that we see only the ERROR messages, and above.
- You might want to replace all INFO levels with WARN, which is more verbose than ERROR but less than INFO.
- You might also want to replace all WARN levels with ERROR. Adjust those levels as you find convenient.
- If, after changing those `log4j.properties` files, you open `pyspark` you will see that the console output is much less verbose.

# Spark 2

- If we run:

```
$ export SPARK_MAJOR_VERSION=2
$ pyspark
SPARK_MAJOR_VERSION is set to 2, using Spark2
Python 2.6.6 (r266:84292, Aug 18 2016, 15:13:37)
[GCC 4.4.7 20120313 (Red Hat 4.4.7-17)] on linux2
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use
setLogLevel(newLevel).
18/01/13 01:23:44 WARN Utils: Service 'SparkUI' could not bind on port 4040.
Attempting port 4041.
/usr/hdp/current/spark2-client/python/pyspark/context.py:205: UserWarning:
Support for Python 2.6 is deprecated as of Spark 2.0.0
  warnings.warn("Support for Python 2.6 is deprecated as of Spark 2.0.0")
Welcome to
```

[illegible]

Using Python version 2.6.6 (r266:84292, Aug 18 2016 15:13:37)  
SparkSession available as 'spark'.

&gt;&gt;&gt;

# Load Data into DataFrame, Spark 2.2

- In Spark 2.2 RDDs are still around but you are advised to use a different type of data objects called DataFrame-s. Also, in Spark 2.2 SparkContext ( object "sc") is implicit and you rather use spark session object, denoted by "spark".

- In Spark 2, previous commands would read:

```
>>> lines = spark.read.text("file:///home/centos/4300-0.txt")
```

```
>>> lines.count()
```

```
32742
```

```
>>> lines
```

```
DataFrame[value: string]
```

- In spark-shell which we use to work with Scala, you would use a very similar command:

```
scala> val textFile = spark.read.textFile("file:///home/centos/4300-0.txt ")
```

```
textFile: org.apache.spark.sql.Dataset[String] = [value: string]
```

- Spark 2.x DataFrame is strongly-typed like an RDD, but with richer optimizations under the hood.
- DataFrame objects have better performance than RDDs.

# Load Data (DataFrame) from HDFS

- We could load the same data from the Hadoops Distributed File System (HDFS)
- We will learn how to use HDFS in the next lecture.
- If we happen to have the `4300-0.txt` file in `centos` home directory in HDFS, we could do the following:

```
>>> blines = spark.read.text ("hdfs:///user/centos/4300-0.txt")
>>> blines.count()
32710
>>> blines.first()
Row(value=u'')
```

- What we did above was create an RDD named `blines` and populate that RDD with data from HDFS resident file `4300-0.txt`.
- We also see in action another method of DataFrames-s, `first()`, which tells us that the first line in `blines` is some uninterested collection of characters (`u''`).
- Please note that we are not terminating our commands with a semi-colon (`;`) or anything else aside from the carriage return. Savings on typing all those semi-colons is one of the greatest contributions of Python to the computer science.
- By the way, our commands are in Python

# Examine a DataFrame

- DataFrame method `show()` list first 20 rows

```
>>> lines.show()
```

```
+-----+
|          value|
+-----+
|The Project Guten...|
|                  |
|This eBook is for...|
|no restrictions w...|
|it under the term...|
|eBook or online a...|
|    Title: Ulysses|
|                  |
| Author: James Joyce|
|                  |
|Release Date: Aug...|
|Last Updated: Aug...|
|                  |
|  Language: English|
|                  |
|Character set enc...|
+-----+
```

```
only showing top 20 rows
```

```
>>>
```

## RDD `filter()` Method

- RDD method `filter()` takes a function returning `True` or `False` and applies it to a sequence (list) and returns only those members of the sequence for which the function returned `True`.

- So, in the Python code :

```
heavens =lines.filter(lambda line: "Heaven" in line)
```

- Method `filter()` acts on the collection `lines`, and passes every element of that collection as the variable `line` as the argument to the anonymous function created using lambda construct. That anonymous function uses a simple regular expression to test whether string "Heaven" exists in variable `line`.
- If the regular expression returns `True` for a particular `line`, an element of collection `lines`, the anonymous function will return `True` and for that particular `line`, `filter()` will return/add variable `line` building up a new collection called `heavens`.
- You can accomplish the same in Java 1.7 and older with several lines of code. In Java 1.8 you have similarly efficient lambda constructs.



## RDD `filter()` method

```
>>> lines = sc.textFile("/home/centos/4300-0.txt")  
heaven = lines.filter(lambda line: "heaven" in line)  
>>> heaven.count()
```

50

- We see in action a method of RDD-s, `filter()`, which apparently let us inquire how many times is `heaven` mentioned in the Ulysses.
- Heaven is mentioned 52 time, i.e. some 0.15% of the time (> Bible)

# DataFrame filtering

- If we load data into a DataFrame, like

```
>>> dset = spark.read.text("file:///home/centos/4300-0.txt")
```

- and want to extract all lines with word `heaven` we would use slightly different syntax:

```
>>> dset.count()
```

```
32710
```

```
>>> heavens = dset.filter(dset.value.contains('heaven'))
```

```
>>> heavens.count()
```

```
47
```

# Python lambda Syntax

- Python supports the creation of anonymous inline functions (i.e. functions that are not bound to a name) at runtime, using a construct called "lambda".
- The following code shows the difference between a normal function definition ("f") and a lambda function ("g"):

```
>>> def f(x): return x**2
```

```
>>> print f(8)
```

```
64
```

```
>>> g = lambda x: x**2
```

```
>>> print g(8)
```

```
64
```

- As you can see, `f()` and `g()` do exactly the same thing. The lambda definition does not include a "return" statement. The last expression is returned.
- You can put a lambda definition anywhere a function is expected, and you don't have to assign it to a variable at all.

# Standalone Application in Python 1.6

- To create an application (Python script) we need to import some Python classes and create `SparkContext` object.
- The rest of the application is coded as if you are writing code in PySpark shell

```
from pyspark import SparkConf, SparkContext
```

```
conf = SparkConf().setMaster("local").setAppName("MyApp")
sc = SparkContext(conf = conf)
lines = sc.textFile("/user/centos/Ulysses/4300-0.txt")
lifeLines = lines.filter(lambda line: "life" in line)
print lifeLines.first()
```

- If we invoke the above with `spark-submit` script:

```
$ spark-submit my_script.py
```

- We get:

```
Arius, warring his life long upon the consubstantiality of
the Son with
```

# Standalone App in Python, Spark 2.2

- Applications can be submitted to a cluster of any type using the `spark-submit` script.

```
from pyspark.sql import SparkSession

textfile = "/user/centos/Ulysses/4300-0.txt" # some file on your HDFS
spark = SparkSession.builder.appName(appName).getOrCreate()
textfile = spark.read.text(textfile).cache()

numAs = textfile.filter(textfile.value.contains('a')).count()
numBs = textfile.filter(textfile.value.contains('b')).count()

print("Lines with a: %i, lines with b: %i" % (numAs, numBs))

spark.stop()
```

- This program just counts the number of lines containing 'a' and the number containing 'b' in a text file.

```
$ spark-submit your_script.py
```

# Introduction to Hive

# Access Hive, using `hive` command line client

- On the master node of your HDInsight cluster or on the command line of our Hortonworks HDP VM, type `hive` and Hive shell opens:

```
[centos@sandbox-hdp ~]$ hive
log4j:WARN No such property [maxFileSize] in
org.apache.log4j.DailyRollingFileAppender.
```

```
Logging initialized using configuration in file:/etc/hive/2.6.3.0-235/0/hive-
log4j.properties
```

```
hive> show tables;
```

```
OK
```

```
sample_07
```

```
sample_08
```

```
Time taken: 3.782 seconds, Fetched: 2 row(s)
```

```
hive> select * from sample_07 limit 10;
```

```
OK
```

00-0000	All Occupations	134354250	40690	
11-0000	Management occupations	6003930	96150	
11-1011	Chief executives	299160	151370	
11-1021	General and operations managers	1655410	103780	
11-1031	Legislators	61110	33880	
11-2011	Advertising and promotions managers		36300	91100
11-2021	Marketing managers	165240	113400	
11-2022	Sales managers	322170	106790	
11-2031	Public relations managers		47210	97170
11-3011	Administrative services managers		239360	76370

```
Time taken: 1.775 seconds, Fetched: 10 row(s)
```

```
hive>
```

# Examine Table Definitions

- To connect to hive, a new client for Hive you do this:

```
hive> desc sample_08;
```

```
OK
```

```
code                string
```

```
description         string
```

```
total_emp           int
```

```
salary              int
```

```
Time taken: 0.69 seconds, Fetched: 4 row(s)
```

```
hive> desc sample_07;
```

```
OK
```

```
code                string
```

```
description         string
```

```
total_emp           int
```

```
salary              int
```

```
Time taken: 0.638 seconds, Fetched: 4 row(s)
```

```
hive>
```



# Hive Visual Client

- Go to Ambari dashboard and select Hive in the left column.
- On the next screen select Hive View 2.0 [Go to View](#)

The screenshot shows the Ambari dashboard interface. The top navigation bar includes the Ambari logo, the environment name 'Sandbox', and buttons for '0 ops' and '0 alerts'. The main content area is divided into a left sidebar and a right pane. The sidebar lists various services: HDFS, YARN, MapReduce2, Tez, Hive, HBase, Pig, and Sqoop. The 'Hive' service is highlighted. The right pane shows the 'Hive' service configuration page. It includes tabs for 'Metrics', 'Heatmaps', and 'Config History'. A 'Metric Actions' dropdown menu is open, showing a list of services including YARN, MapReduce2, Tez, Hive (selected), HBase, Pig, Sqoop, Oozie, ZooKeeper, and Falcon. The 'Summary' section on the right lists the status of various Hive components: Hive Metastore, HiveServer2, MySQL Server, WebHCat Server, and Hive Client. It also displays the HiveServer2 JDBC URL and provides links to 'Go To View' for Hive View 2.0 and Debug Hive Query.

Ambari Sandbox 0 ops 0 alerts Dashboard Services

Metrics Heatmaps Config History

Metric Actions Last 1 hour

HDFS D

YARN

MapReduce2

Tez

Hive

HBase

Pig

Sqoop

Oozie

ZooKeeper

Falcon

Summary

[Hive Metastore](#) Started No alerts

[HiveServer2](#) Started No alerts

[MySQL Server](#) Started No alerts

[WebHCat Server](#) Started No alerts

[Hive Client](#) 1 Hive Client Installed

HiveServer2 JDBC URL jdbc:hive2://sandbox-hd

Hive View 2.0 [Go To View](#)

Debug Hive Query [Go To View](#)

# Hive Query Editor

**HIVE**

[+ NEW JOB](#)[+ NEW TABLE](#)

[QUERY](#)[JOBS](#)[TABLES](#)[SAVED QUERIES](#)[UDFs](#)[SETTINGS](#)[NOTIFICATIONS](#)

Worksheet1 \*

+

DATABASE

Select or search database/schema

x default

Browse ▾

1 describe sample\_07;

2 select \* from sample\_07 limit 5;

✓ Execute

Save As

Insert UDF ▾

Visual Explain

RESULTS

LOG

VISUAL EXPLAIN

TEZ UI

Filter columns

x

≡

←

→

↗

sample_07.code	sample_07.description	sample_07.total_emp	sample_07.salary
00-0000	All Occupations	134354250	40690
11-0000	Management occupations	6003930	96150
11-1011	Chief executives	299160	151370
11-1021	General and operations managers	1655410	103780

# Sample Data, Shakespeare and King James

- On my machine there are two files:  
all-shakespeare and all-bible
- One contains complete works of Shakespeare and the other King James Bible. Both works use somewhat archaic form of English.
- You could examine the files by perhaps doing:

```
$ cat all-shakespeare | wc  or
$ cat all-shakespeare | tail -n 100
```

```
$ cat all-bible | wc  or
$ cat all-bible | tail -n 100
```

- We could transfer both files to our VM using `scp` command:

```
$ scp -P 2222 all-bible all-shakespeare centos@192.168.135.132:/home/centos
centos@192.168.135.132's password:
all-bible          100% 5135KB  27.4MB/s   00:00
all-shakespeare    100% 5160KB  47.4MB/s   00:00
```

# Copy all-shakespeare, all-bible into HDFS

- We will copy these file into HDFS directories: `input_bible` and `input_shake` respectively:

```
$ hadoop fs -mkdir input_bible
$ hadoop fs -mkdir input_shake
$ hadoop fs -put all-bible input_bible
$ hadoop fs -put all-shakespear input_shake
```

- We could convince ourselves that the data inside HDFS is still the same old Shakespeare by typing something like:

```
$ hadoop fs -cat input_shake/all-shakespeare | head -n 20
```

ALL'S WELL THAT ENDS WELL

DRAMATIS PERSONAE

KING OF FRANCE (KING:)

DUKE OF FLORENCE (DUKE:)

BERTRAM Count of Rousillon.

LAFEU an old lord.

PAROLLES a follower of Bertram.

# Running a MapReduce `grep` Job

- Next we will use MapReduce `grep` job to count the number of word occurrences in both texts.
- We will place respective outputs in HDFS directories `bible_freq` and `shake_freq`:

```
$ hadoop jar /usr/hdp/2.6.3.0-235/hadoop-mapreduce/hadoop-mapreduce-examples.jar grep input_bible bible_freq '\w+'
```

```
$ hadoop jar /usr/hdp/2.6.3.0-235/hadoop-mapreduce/hadoop-mapreduce-examples.jar grep input_shake shake_freq '\w+'
```

# Examine Result of `grep`

- Examine the HDFS directories by typing:

```
drwx----- - centos mapred      0 2018-01-13 14:40 .Trash
drwxr-xr-x - centos mapred      0 2018-01-13 01:38 .hiveJars
drwx----- - centos mapred      0 2018-01-13 15:39 .staging
drwxr-xr-x - centos mapred      0 2018-01-13 15:37 bible_freq
drwxr-xr-x - centos mapred      0 2018-01-12 23:57 input
drwxr-xr-x - centos mapred      0 2018-01-13 15:26 input_bible
drwxr-xr-x - centos mapred      0 2018-01-13 15:27 input_shake
drwxr-xr-x - centos mapred      0 2018-01-13 15:39 shake_freq
drwxr-xr-x - centos mapred      0 2018-01-12 23:42 ulysses
```

# Examine Content of the output files

- We could also see partial content of the output files:

```
[centos@sandbox-hdp ~]$ hadoop fs -cat bible_freq/part-r-00000 | head -n 10
62394    the
38985    and
34654    of
13526    to
12846    And
12603    that
12445    in
9764     shall
9672     he
8940     unto
cat: Unable to write to output stream.
[centos@sandbox-hdp ~]$ hadoop fs -cat shake_freq/part-r-00000 | head -n 10
25578    the
23027    I
19654    and
17462    to
16444    of
13524    a
12697    you
11296    my
10699    in
8857     is
cat: Unable to write to output stream.
[centos@sandbox-hdp ~]$
```

- These are frequency - word pairs, as expected.

# Create tables to accept `grep` data

- Before opening `hive` client, and in particular before importing any data from HDFS, add your current Linux user, e.g. `centos`, to group Hadoop:

```
$ su -
```

```
Passwd:
```

```
$ usermod -aG hadoop centos
```

- We want to import of Shakespeare frequency data. We first create tables `shakespeare` and `bible` by issuing the following commands:

```
hive> create table shakespeare (freq INT, word STRING) ROW FORMAT  
      DELIMITED FIELDS TERMINATED BY '\t' stored as textfile;
```

- This created table `shakespeare` with out any data

```
hive> show tables;
```

```
shakespeare
```

```
Time taken: 8.268 seconds
```

```
hive> describe shakespeare;
```

```
OK
```

```
freq      int
```

```
word      string
```

```
Time taken: 1.253 seconds
```

- Similarly, for `bible`:

```
hive> create table bible (freq INT, word STRING) ROW FORMAT  
      DELIMITED FIELDS TERMINATED BY '\t' stored as textfile;
```



# Load grep Data into `shakespeare` Table

- To load data we go back to the `hive` and type:

```
hive> LOAD DATA INPATH "/user/centos/shake_freq" OVERWRITE INTO
      TABLE shakespeare;      # From HDFS file system or
hive> LOAD DATA LOCAL INPATH "/home/centos/part-r-00000" INTO TABLE
      shakespeare;      # From the local file system
Loading data to table shakespeare
OK
Time taken: 0.213 seconds
```

- `OVERWRITE` clause is optional. You need it to avoid duplicates or simply overwrite an existing table content.
- On the load command, Hive moved HDFS content of `shake_freq` into its own HDFS directory. That directory on Hortonworks HDP VM is located in HDFS directory:  
`/apps/hive/warehouse/shakespeare`
- Please examine that directory using `hadoop fs` commands
- Note again, the directory `/apps/hive/warehouse` is in HDFS, not on Linux OS.

# Verify that shakespeare has grep Data

```
hive> select * from shakespeare limit 10;
```

```
OK
```

```
25848    the
```

```
23031    I
```

```
19671    and
```

```
18038    to
```

```
16700    of
```

```
14170    a
```

```
12702    you
```

```
11297    my
```

```
10797    in
```

```
8882     is
```

```
Time taken: 0.095 seconds
```

```
hive>
```

- This statement read from the table (actually as part of optimization, it read directly from the HDFS file) and presented us with the first 10 lines.
- This is the same data we saw previously.

# More Advanced Query

- Slightly more advanced query would perhaps be this one:

```
hive> SELECT * FROM shakespeare  
WHERE freq > 100 SORT BY freq ASC  
LIMIT 10;
```

- Notice that for a large data set this is not an entirely trivial job.
- Data has to be sorted before we could see 10 rows of words that have frequency just above 100.
- Notice how hive reports on map-reduce job it is starting.
- If the job takes too long you are given the job id and the command that you could execute to tell Hadoop to kill the job:

```
Starting Job = job_201404021324_0005, Tracking URL =  
    http://quickstart:50030/jobdetails.jsp?jobid=job_201404021324_0005  
Kill Command = /usr/lib/hadoop/bin/hadoop job -  
    Dmapred.job.tracker=quickstart:8021 -kill job_201404021324_0005
```

## Even More Complex Query

- The "users", linguists perhaps, would like to know the number of words which appear with the most common frequencies.

```
hive> SELECT freq, COUNT(1) AS f2
      FROM shakespeare GROUP BY freq SORT BY f2 DESC LIMIT 10;
```

- OK
- 1      13426
- 2      4274
- 3      2342
- 4      1502
- 5      1111
- 6      873
- 7      656
- 8      598
- 9      474
- 10     381
- This tells us that there are 13426 words that appears only once.
- 4274 words appear twice. 2342 words appear three times, etc.
- SQL command with minor deviation: ORDER BY is replaced by SORT BY.

# Joining Tables

- One of the most powerful feature of Hive is the ability to create queries that joins tables together using regular SQL syntax.
- We have (`freq`, `word`) data for Shakespeare works and for King James Bible.
- We want to examine which words show up in both volumes of text.

# Import data into bible

```
hive> LOAD DATA INPATH "/user/cloudera/bible_freq" INTO TABLE  
Bible;
```

```
OK
```

```
Time taken: 0.781 seconds
```

```
hive> select * from bible limit 20;
```

```
OK
```

```
62394    the  
38985    and  
34654    of  
13526    to  
12846    And  
12603    that  
12445    in  
6913     be  
6884     is  
6649     him  
6647     LORD
```

```
. . .
```

```
Time taken: 0.111 seconds
```

```
hive>
```

# Create an Intermediate Table

- We need a table that will list most common words in both volumes with corresponding frequencies

```
hive> CREATE TABLE merged  
      (word STRING, shake_f INT, kjb_f INT);
```

- For this table we do not need to specify how will data be stored.
- Hive will determine that by itself.
- Next, we will run a query that will select data from tables: `shakespeare` and `bible`, create a join and insert, i.e. overwrite the content of new table.
- In our case the table happens to be empty. If it were not empty and we insist on overwriting, table data would be lost. If we only perform an insert, new data would be appended to the old.

# Populate merged table using join

```
hive> INSERT OVERWRITE TABLE merged
SELECT s.word, s.freq, k.freq FROM
shakespeare s JOIN bible k ON
(s.word = k.word)
WHERE s.freq >= 1 AND k.freq >= 1;
```

```
Query ID = centos_20180113172124_cc3bc7a0-928f-4c28-a36d-f6200dfa2d0d
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1515797541938_0017)
```

```
-----
      VERTICES      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 .....  SUCCEEDED      1          1          0          0          0          0
Map 2 .....  SUCCEEDED      1          1          0          0          0          0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 9.64 s
-----
```

```
Loading data to table default.merged
Table default.merged stats: [numFiles=1, numRows=12432, totalSize=150351,
  rawDataSize=137919]
OK
Time taken: 19.161 seconds
```



# Most common common words

- What the words that appear most frequently in both corpuses?

```
hive> SELECT word, shake_f, kjb_f, (shake_f + kjb_f) AS ss  
FROM merged SORT BY ss DESC LIMIT 20;
```

the	25848	62394	88242
and	19671	38985	58656
of	16700	34654	51354
I	23031	8854	31885
to	18038	13526	31564
in	10797	12445	23242
a	14170	8057	22227
that	8869	12603	21472
And	7800	12846	20646
is	8882	6884	15766
my	11297	4135	15432
you	12702	2720	15422
he	5720	9672	15392
his	6817	8385	15202
not	8409	6591	15000
be	6773	6913	13686
for	6309	7270	13579
with	7284	6057	13341
it	7178	5917	13095
shall	3293	9764	13057

# To examine common non-Stop Word, go deeper

```
SELECT word, shake_f, kjb_f, (shake_f + kjb_f) AS ss  
FROM merged SORT BY ss DESC LIMIT 200;
```

```
. . . . .  
heaven  626      578      1204  
When    847      349      1196  
Of       1006     63       1191  
most    1017     135      1152  
where   813      335      1148  
tell    960      188      1148  
blood   699      447      1146  
doth    961      63       1146  
set     451      694      1145  
It       890     241      1131  
ever    634      475      1109  
Which   977      130      1107  
whom    375      732      1107  
Time taken: 46.988 seconds
```